

CS2200

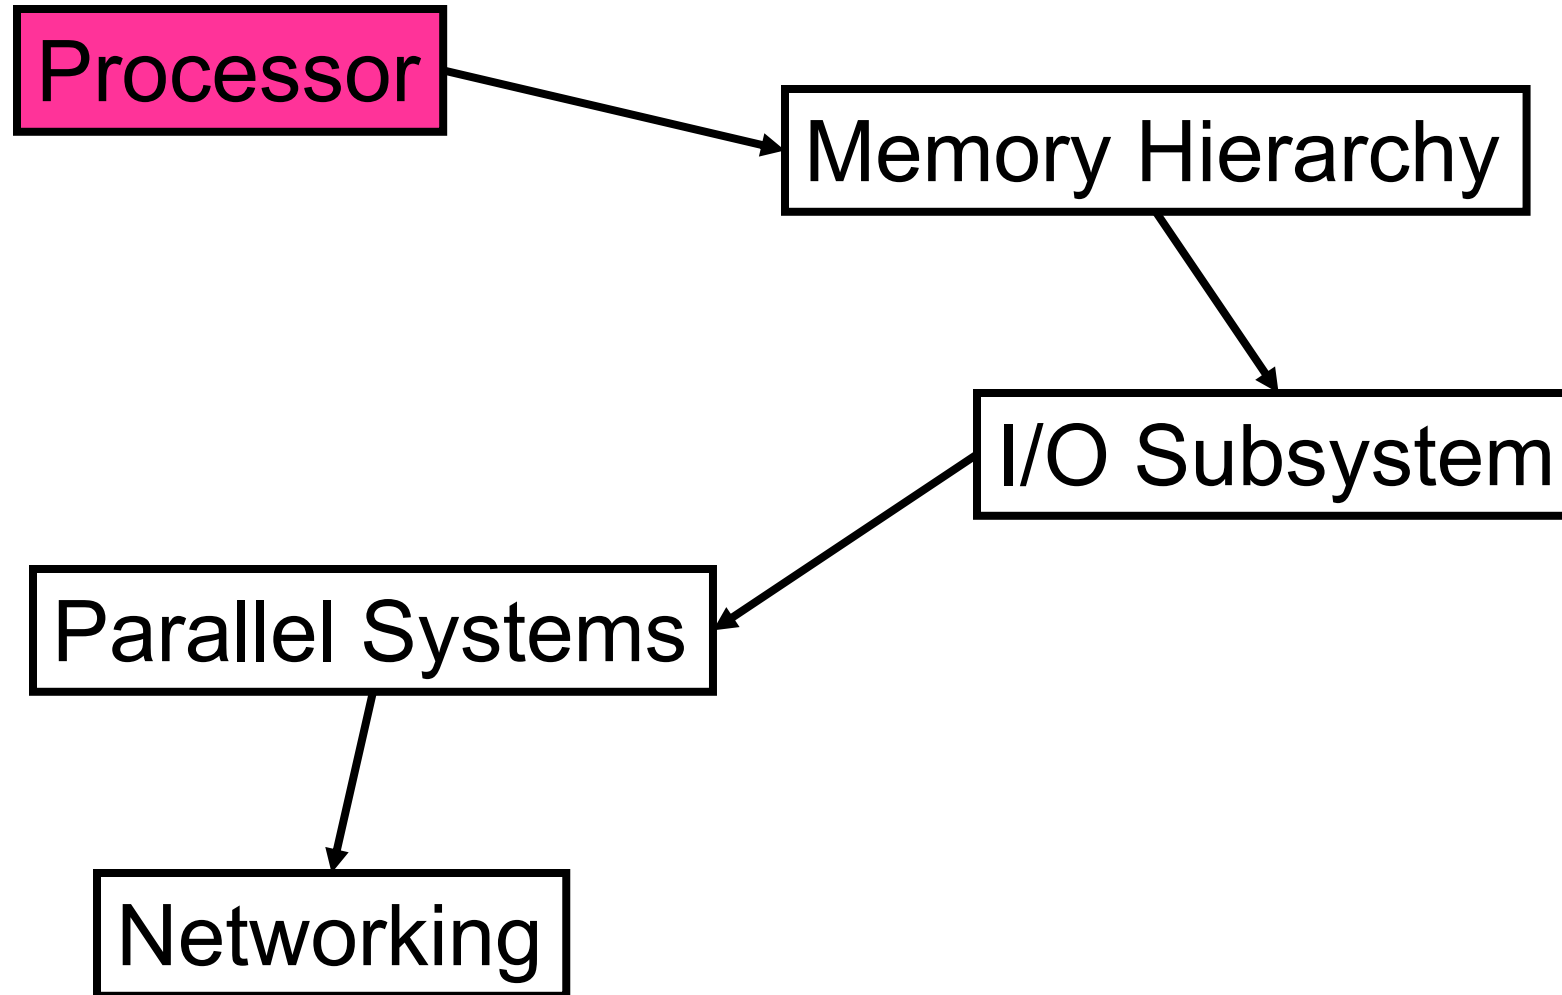
Systems and Networks

Spring 2024


Lecture 2: Processors

Alexandros (Alex) Daglis
School of Computer Science
Georgia Institute of Technology
adaglis@gatech.edu

Our Road Map




A Computer's Key Components



What does the processor do?

- Knows where it is in program
- Can get and put data into memory
- Can do some arithmetic
- Can make tests and take different paths depending on the results
- Do you need a language to make a computer run?

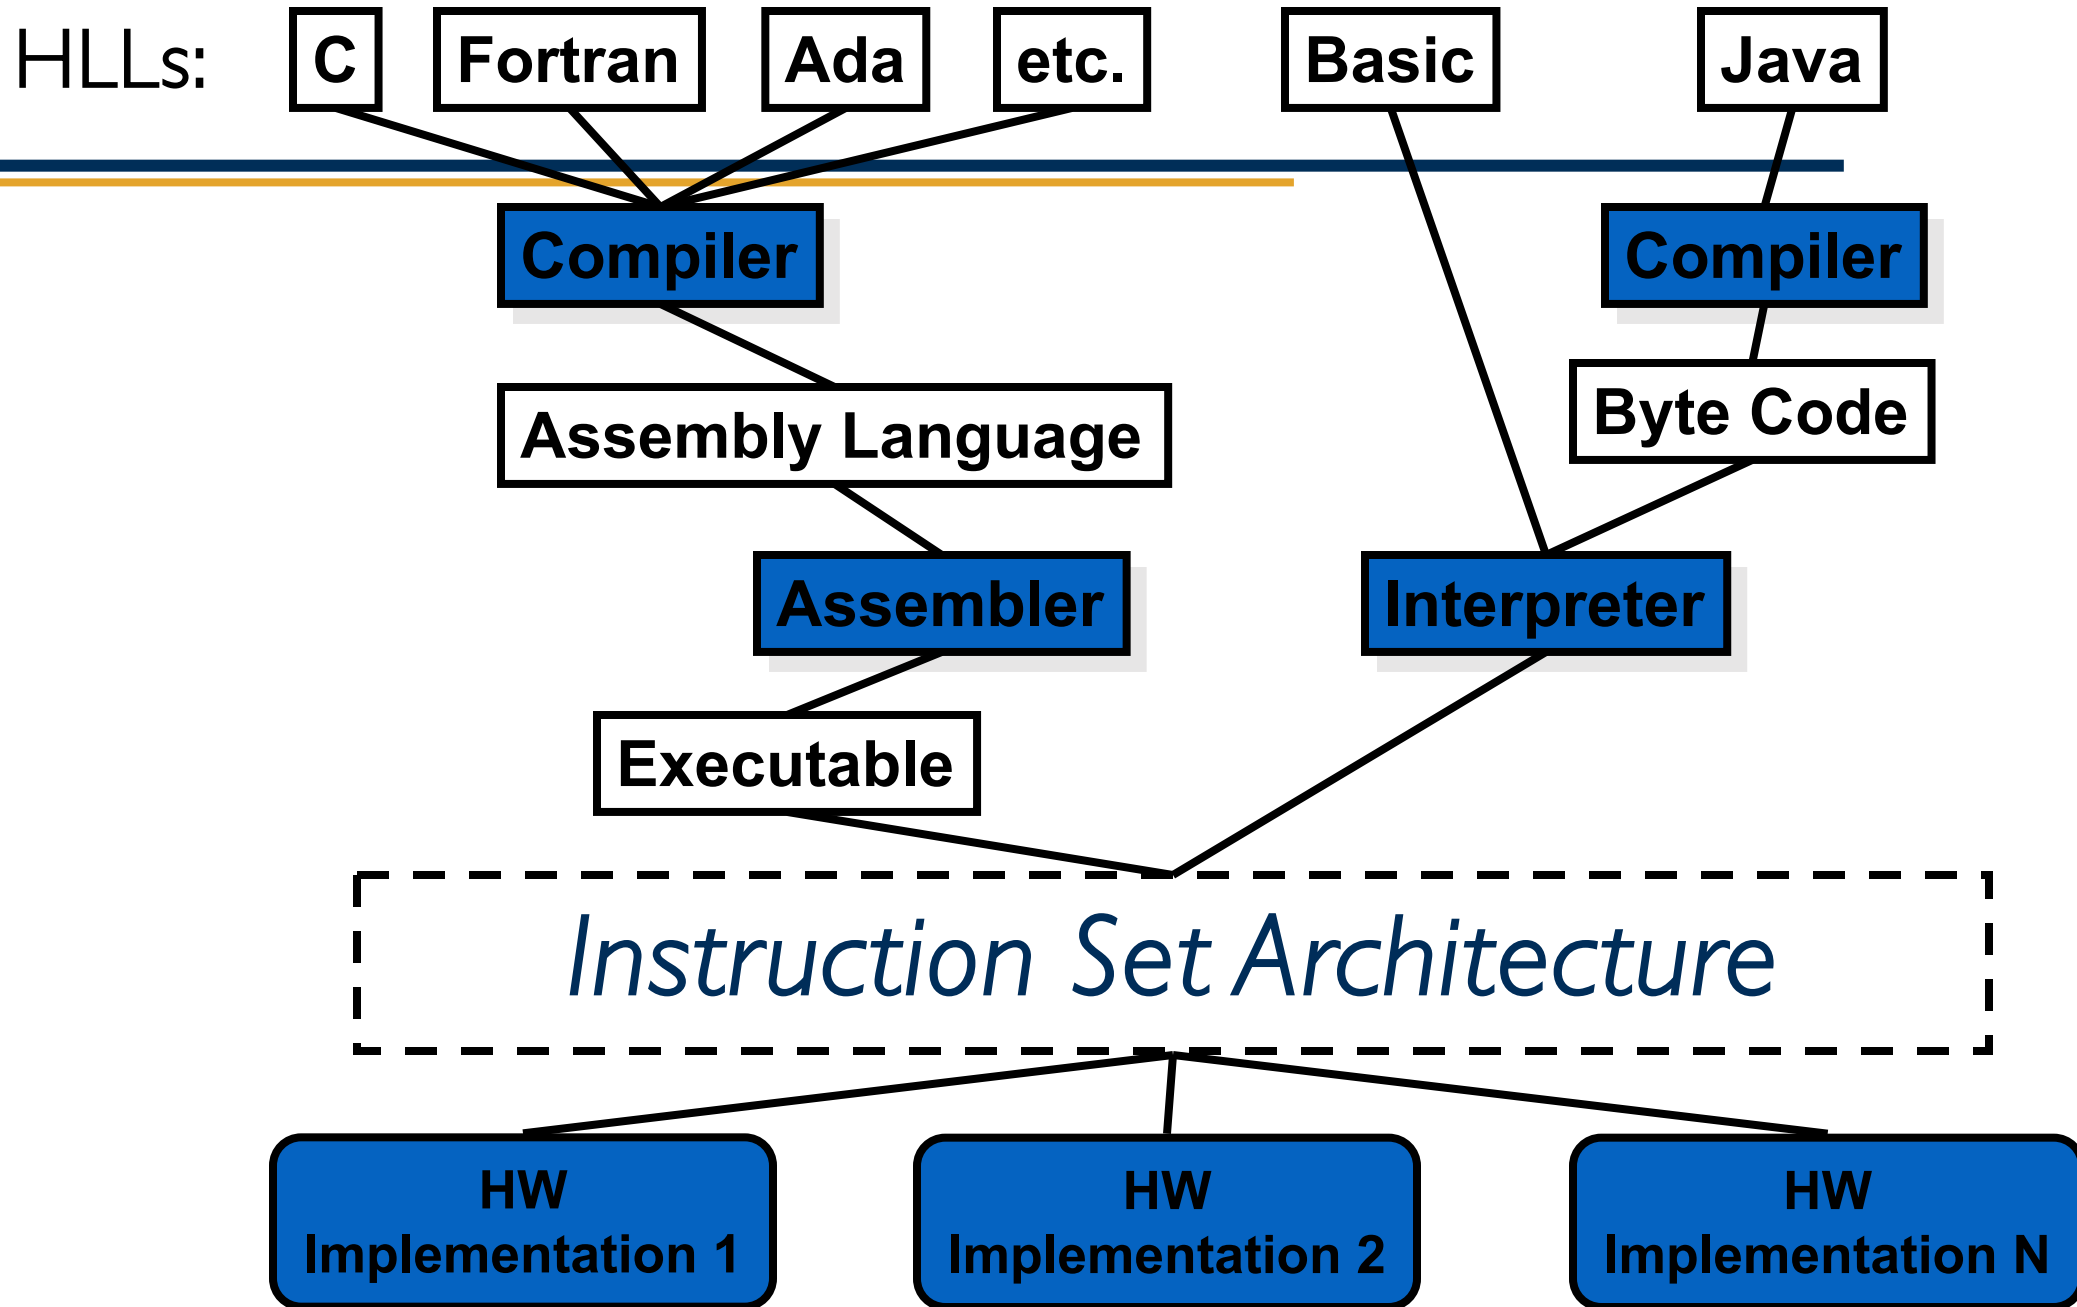
A Little History



Replacing a bad tube meant checking among ENIAC's 19,000 possibilities.

A Little History



- First computers programmed by hand
1000110010100000
- Somewhat tedious, so invented:
- Assembler
add A,B
- If we can convert from Assembly Language to machine code why not from some higher-level language to Assembler?
 $A + B$



Instructions

- Language of the machine
- Vocabulary is the instruction set (ISA)
- Two levels
 - Human readable (assembly)
 - Machine readable (machine code)

Computing Evolution in ~70 Years



17T ops/s in 16sq in




3 billion ENIACs in your palm

55 trillion times higher compute density

Moving forward: ISA Construction


- Instruction set design from HLL constructs
 - Expressions, assignments => ALU instructions
 - Data abstraction => Addressing modes
 - Conditional & loop statements => Branch instructions
 - Procedure calls/returns => stack management
- Please note the reading assignments in the schedule: Start reading chapter 2

Simple Machine Model



- Remember the LC-3? It used a greatly simplified ARM instruction set
- We'll be introducing the LC-2200, a greatly simplified MIPS instruction set
 - Architecture that's similar, but not the same as the LC-3.


Simple Machine Model



- Let's consider the execution of a HLL
- ```
a = a + 1
c = a + b
if (c == d) {
 ...
}
```

# How to Design an Instruction Set?

---




Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?

# Arithmetic/Logical Expressions

---



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?


$c = a + b$

becomes

add c, a, b

# Arithmetic/Logical Expressions

---



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?


$c = a + b$  ➔ add  $c, a, b$

Keep adding to repertoire

$c = a - b$  ➔ sub  $c, a, b$

$c = !(a \& b)$  ➔ nand  $c, a, b$

# Arithmetic/Logical Expressions



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?

$c = a + b$  ➔ add c, a, b

➔ memory operands


➔ memory addressing mode

What do you call these?

Where are they?



# Arithmetic/Logical Expressions



Start thinking like a compiler writer

➔ What instructions are needed for each HLL construct?


$c = a + b$  ➔ add c, a, b

Is there a downside  
to operands in  
memory?

How can we address  
that?

A trip to memory is  
**EXPENSIVE!**

# Operands?




$c = a + b \rightarrow \text{add } c, a, b$

How about introducing Load/Store instructions?

ld  $r_1, a$

st  $c, r_2$

# Load/Store Instructions



ld r<sub>1</sub>, a


st c, r<sub>2</sub>

➔ We've got operands in registers

➔ Register addressing mode!

So how do we compile  $c = a + b$  now?

# Register Operands



Old way:

add c, a, b

New way:


ld r<sub>1</sub>, a

ld r<sub>2</sub>, b

add r<sub>3</sub>, r<sub>1</sub>, r<sub>2</sub>

st c, r<sub>3</sub>

# Compiling with Register Operands



Old way:

add c, a, b

1 instruction

**THIS IS THE WAY**

New way:

ld r<sub>1</sub>, a

ld r<sub>2</sub>, b

add r<sub>3</sub>, r<sub>1</sub>, r<sub>2</sub>

st c, r<sub>3</sub>

4 instructions

This looks dumb!

Not really.  
Why?

We can re-use the  
values in registers!


# Keep Frequently Used Tools Nearby!

---



...or, the principle of **locality**

# Reusing Values



```
c = a + b
d = a * b + c
if (c == d) {
 ...
}
```

With operands left in registers, we can save three memory accesses here!



# Why would we consider loading values into registers before computing with them?

Doesn't that use more instructions?

- 0% A. Yes it does, but it doesn't matter how many instructions it takes.
- 0% B. No it doesn't. You counted wrong.
- 0% C. Yes it does, but it saves memory accesses because we can re-use the values.
- 0% D. Yes it does and it's a terrible design choice.






# Structs in HLL

```
struct {
 int b;
 int c;
} a;
```

Elements of a struct are contiguous in memory

How do we load b and c into registers?



# Accessing Struct Members

```
struct {
 int b;
 int c;
} a;
```


Let's say `&a` is already in register  $R_1$

To load `b`:

$R_2 \leftarrow \text{memory}[R_1 + 0]$

To load `c`:

$R_3 \leftarrow \text{memory}[R_1 + 4]$



# Base + Offset Addressing Mode

```
struct {
 int b;
 int c;
} a;
```


Let's say &a is already in register  $R_1$

$\text{ld } R_i, \text{offset}(R_{\text{base}})$

To load b and c:

$\text{ld } R_2, 0(R_1)$

$\text{ld } R_3, 4(R_1)$



# Operand Granularity

---

char → 8 bits → byte  
short → 16 bits → half word  
int → 32 bits\* → word  
long → 64 bits\*

\*depends on the word size of the architecture:  
int=16, long=32 and others can happen

We need some instruction variants:

ldb, ldh, ldl, ...

similar for store instructions


# Operand Alignment

---

```
struct {
 char a;
 char b[3];
}
```

# Dense Packing


---



- We're packing operands to save space.

# A Different Struct


```
struct {
 char a;
 int b;
}
```



If we have a 32-bit-aligned path to memory, how are we going to load b?

# Why Alignment Rules Matter

```
struct {
 char a;
 int b;
}
```



Now, how  
to load b  
is obvious

This is one of many space/time tradeoffs that ISA designers must address.



# Accessing Array Operands

```
int a[100];
```

|            |   |             |
|------------|---|-------------|
| <b>100</b> |   | <b>a[0]</b> |
| <b>104</b> |   | <b>a[1]</b> |
| <b>108</b> |   | <b>a[2]</b> |
| <b>112</b> |   | <b>a[3]</b> |
|            | . |             |
|            | . |             |
|            | . |             |
|            | . |             |
| <b>128</b> |   | <b>a[7]</b> |
| <b>132</b> |   | <b>a[8]</b> |
|            | . |             |
|            | . |             |
|            | . |             |
|            | . |             |

One approach, use base+offset

$$r_1 = 100$$

To load a[8]

$$\text{ld } r_2, 32(r_1)$$

Is this the best we can do?

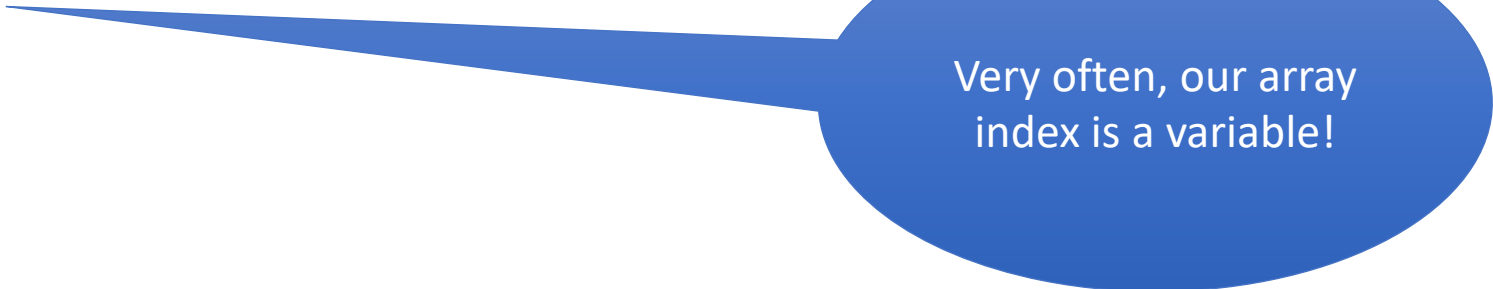
What does code that accesses an array usually look like?

# Typical Array Use

---

- How do we typically use arrays?

```
loop
 c[i] = a[i]
 ...
 i = i + 1
end loop
```



Very often, our array index is a variable!

- Looks like it's time for a new addressing mode
- Let's implement base+index addressing

# Accessing Array Operands

```
int a[100];
```

|     |   |      |
|-----|---|------|
| 100 |   | a[0] |
| 104 |   | a[1] |
| 108 |   | a[2] |
| 112 |   | a[3] |
|     | . |      |
|     | . |      |
|     | . |      |
|     | . |      |
| 128 |   | a[7] |
| 132 |   | a[8] |
|     | . |      |
|     | . |      |
|     | . |      |
|     | . |      |

With base+index mode

$$r_1 = 100$$

$$r_2 = i * 4$$

Why?

To load a[i]

ld r<sub>8</sub>, r<sub>2</sub>(r<sub>1</sub>)

Using base + index

Now we can have this ld in a loop, only need to update r<sub>2</sub> value in each iteration


# Endianness

Say we have a 32-bit register, `RI`, that contains the value `1` in two's-complement. We store that register into an (aligned) memory location

```
ST RI, Mem[100]
```

The value is stored in addresses `100-103` (of course).

Which byte contains the `1`? (The other 3 will be zero, right?)



# Endianness

Tip: Endianness often shows up in byte-addressable memories because we can access individual bytes out of longer data types

Little Endian → addresses the Least Significant Byte (LSB) of the word

`int b = 0x11223344; //placed in address 104`

| +3 |    |    | +2 | +1 | +0 |     |
|----|----|----|----|----|----|-----|
|    |    |    |    |    | a  | 100 |
| 44 | 33 | 22 | 11 |    |    | 104 |
| 11 | 22 | 33 | 44 |    |    |     |

Big Endian → addresses the Most Significant Byte (MSB) of the word

# So What's the Difference?

The difference only shows up when taking a “word” apart  
(in this case, loading 8 bits from a 32-bit integer)

ld**b** r<sub>1</sub>, Mem[104]

Big Endian → loads **0x11** into r<sub>1</sub>

Little Endian → loads **0x44** into r<sub>1</sub>

| +3            |  | +2            |  | +1            |  | +0            |  |     |
|---------------|--|---------------|--|---------------|--|---------------|--|-----|
|               |  |               |  |               |  | a             |  | 100 |
| b... 44<br>11 |  | b... 33<br>22 |  | b... 22<br>33 |  | b... 11<br>44 |  | 104 |

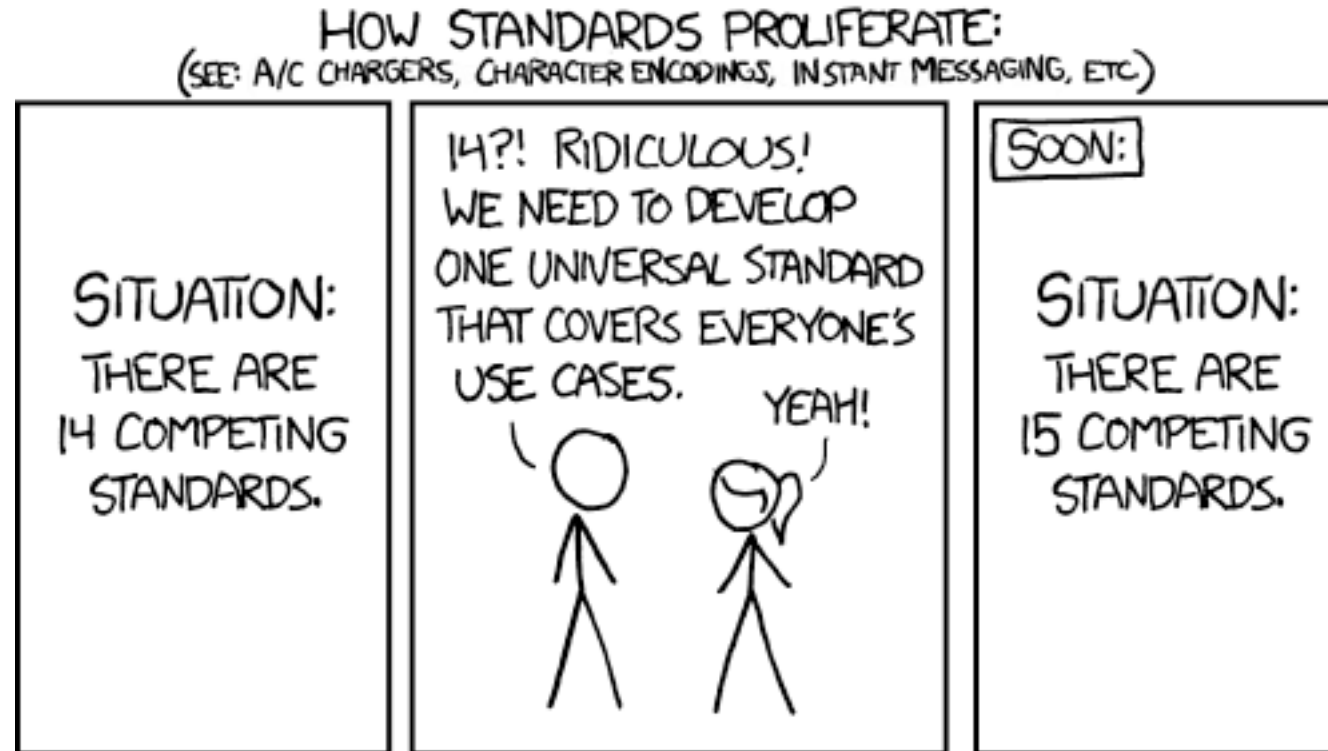
# So What's the Difference

- The difference shows up when taking a word apart
- Let's store 0x11223344 (a 32-bit integer) at location 104

|                               | Little Endian Result | Big Endian Result |
|-------------------------------|----------------------|-------------------|
| ldb r <sub>1</sub> , Mem[104] | 0x44                 | 0x11              |
| ldh r <sub>1</sub> , Mem[104] | 0x3344               | 0x1122            |
| ldw r <sub>1</sub> , Mem[104] | 0x11223344           | 0x11223344        |

|               | 107  | 106  | 105  | 104  |
|---------------|------|------|------|------|
|               | b... | b... | b... | b... |
| Big Endian    | 44   | 33   | 22   | 11   |
| Little Endian | 11   | 22   | 33   | 44   |

# ... why do two different Endiannesses even exist?





# So, about the LC-3

---

- Was it Big Endian or Little Endian?
- Think carefully...
- You can't tell!
- There are no instructions that manipulate more or fewer than 16 bits
- So there isn't any way to show this implementation detail!
- Was this accidental or on purpose?

# Recap

| Software                                                                                  | Hardware                                                  |
|-------------------------------------------------------------------------------------------|-----------------------------------------------------------|
| Expressions & assignments                                                                 | ALU instructions                                          |
| Variable reuse                                                                            | register addressing mode<br>ld/st instructions            |
| Data abstraction <ul style="list-style-type: none"><li>• struct</li><li>• array</li></ul> | base + offset addr mode<br>base + index addr mode         |
| Granularity of operands                                                                   | ldb/ldh/ldw instructions<br>addressability (byte, word)   |
| Packing operands                                                                          | Memory alignment<br>(space/time tradeoff)                 |
| Endianness 0x11223344                                                                     | Little (first byte is 0x44)<br>/ Big (first byte is 0x11) |



# Review Question 1

---

An instruction set...

- 0% A. Serves as a level of abstraction between software and hardware.
- 0% B. Provides the low-level details of the machine implementation.
- 0% C. Deals with the datapath and control implementation of the processor.
- 0% D. None of the above.





# Review Question 2

---

Addressing mode...

- 0% A. Refers to the kinds of opcodes supported in an architecture.
- 0% B. Refers to the way the operands are specified in an instruction that accesses memory.
- 0% C. Refers to the granularity of the memory element that can be addressed in an instruction.
- 0% D. Is a critical tool for USPS operations.
- 0% E. None of the above.





# Review Question 3

Endianness of an architecture...

- 0% A. Is a key determinant of processor performance.
- 0% B. Is a key determinant of how the compiler lays out data structures in memory.
- 0% C. Matters if one declares a datatype of a particular granularity and accesses it at a different granularity.
- 0% D. Is the official name of the party held after completing cs2200.
- 0% E. None of the above.

