



CS2200  
Systems and Networks  
Spring 2024

## Lecture 26: Networking

Alexandros (Alex) Daglis  
School of Computer Science  
Georgia Institute of Technology  
[adaglis@gatech.edu](mailto:adaglis@gatech.edu)

# Networking

---

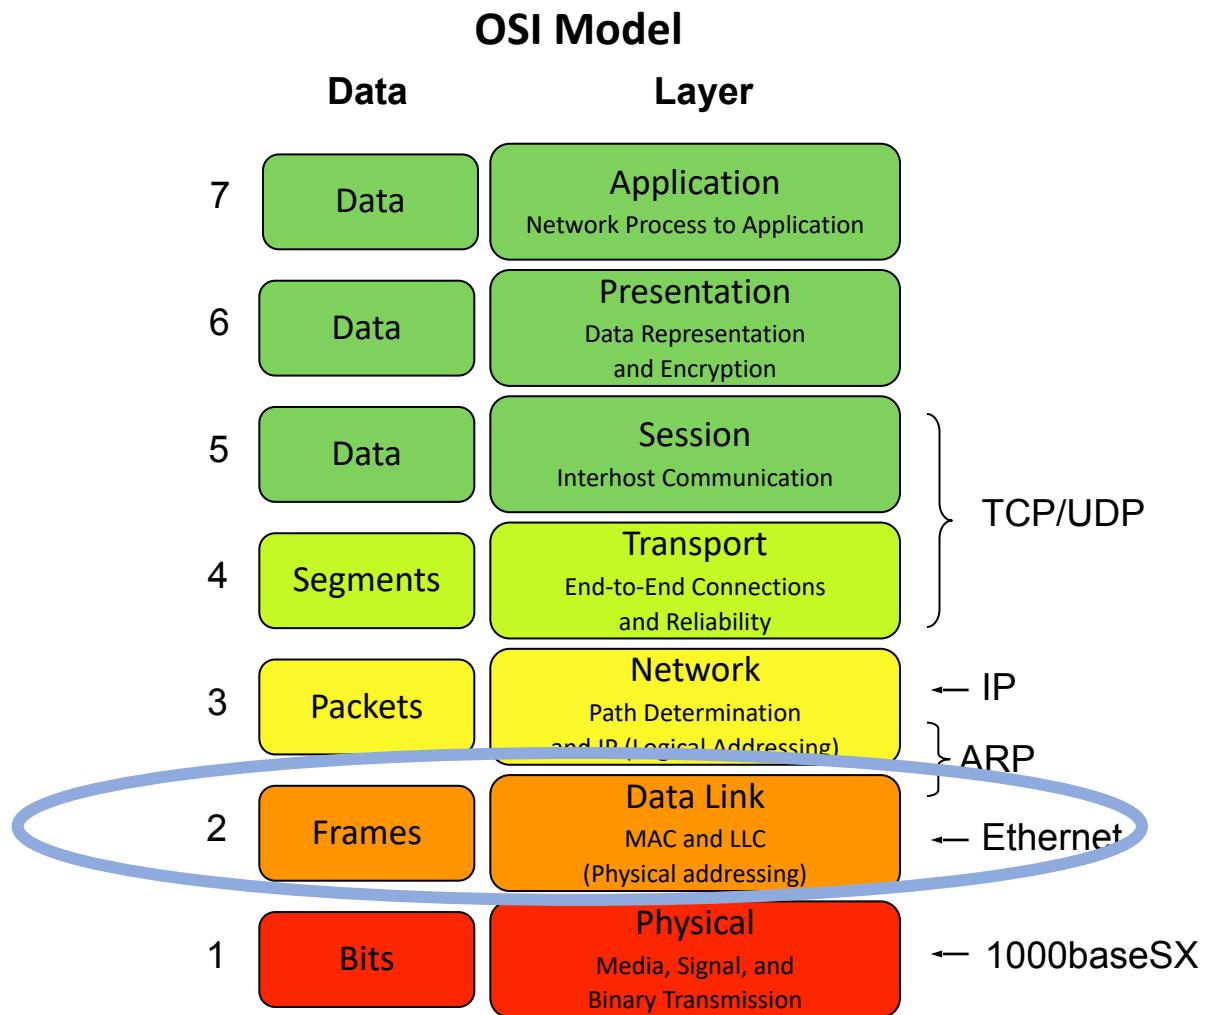
- Network stack overview – layer encapsulation
- Three transport protocol types
- Basic distinctions between TCP and UDP
- Network layer
- Ports and connections

Today:

- Data Link layer (Ethernet)
- Performance metrics
- Routing and IP table construction

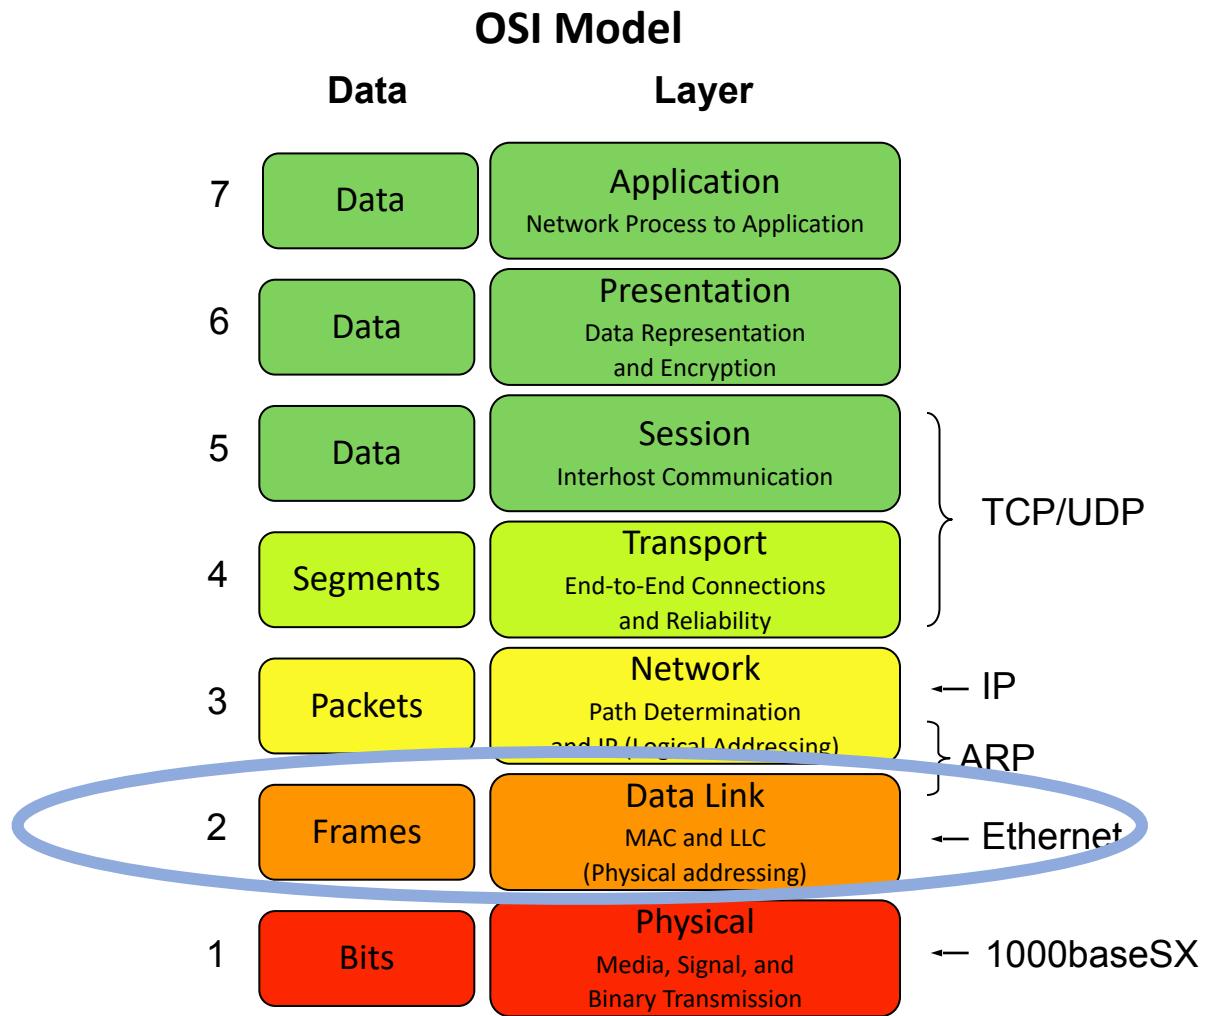
Last 3 lectures: Disk scheduling & filesystem

# Now let's look at the data link layer



# Now let's look at the data link layer

- Ethernet is a very common data link network



# Let's talk about Ethernet

---

- Ethernet is a particular type of data link network
- It was designed for relatively small areas, like a building, hence the term Local Area Network (LAN)
- It has its very own standard (IEEE 802.3)
- It's evolved a great deal since the 1980s

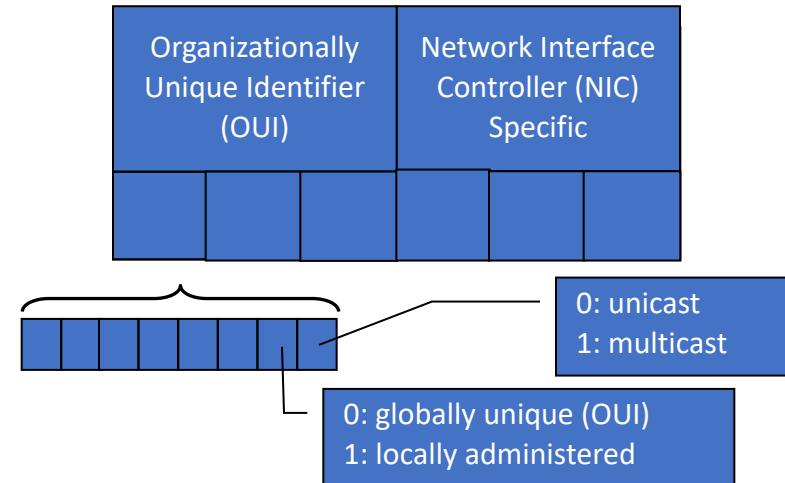
# Terminology

---

- Collision domain - All the NICs that share a physical ethernet; only one can transmit at a time
  - Broadcast domain - All the NICs that can hear an ethernet broadcast frame
- 
- Repeater/hub - layer 1 replication/amplification of bits (both sides stay in one collision domain)
  - Bridge/switch - layer 2 re-transmission of ethernet frames (both sides stay in one broadcast domain)
  - Router - layer 3 forwarding of packets based on a layer-3 (IP) routing table
  - Host – a computer with one or more NICs; if it forwards traffic using multiple NICs, we're going to call it a router, not a host.

# Ethernet Addressing

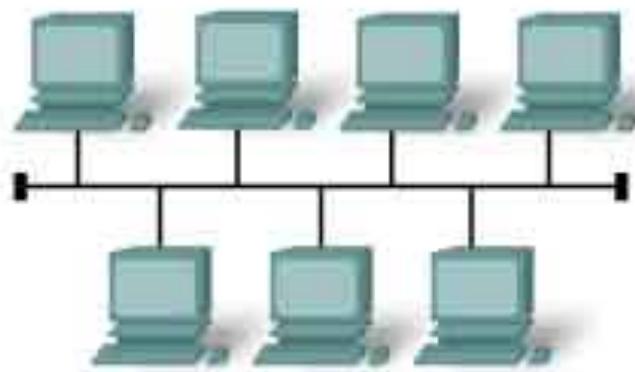
Every ethernet device gets a unique  
burned-in 48-bit MAC address



- 6 bytes/48 bits
  - 3 bytes OUI
  - 3 bytes NIC specific
- Examples
  - Apple: 00:25:00:f4:6d:c2
  - Oracle: 00-03-ba-a1-02-df
- To identify the OUI
  - <https://www.wireshark.org/tools/oui-lookup.html>

# Ethernet Frames

Preamble	Destination MAC	Source MAC	802.1Q tag	Ether type/Length	Payload	CRC	Interframe Gap
8	6	6	0-4	2	46-1500	4	12



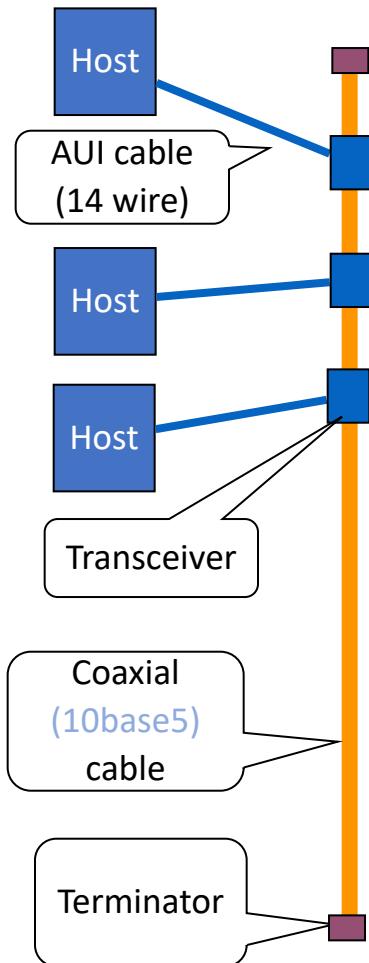
- To send a frame
  - Wait until no one using the cable (carrier sense)
  - Assert carrier signal on the cable
  - Broadcast the frame to all stations on the cable
  - If you detect a **collision** while transmitting, stop and requeue the frame for retransmission (known as CSMA/CD)
- If you hear a frame on the cable
  - Accept it if the destination MAC belongs to us (or if it has a broadcast or multicast destination MAC)
  - Drop it if you see a collision during its receipt or its CRC doesn't match
- Ethernet works like a logical bus!

# Something is Missing...

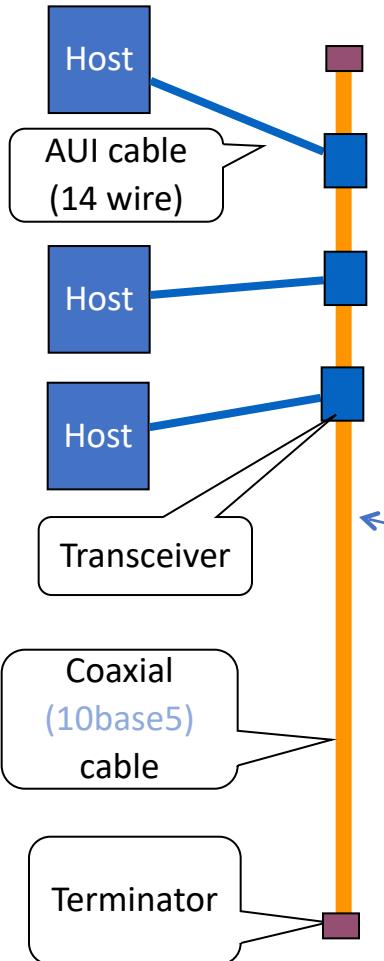
---

- Only one station can send at a time?
- Where are the IP addresses?
- We'll get back to this...

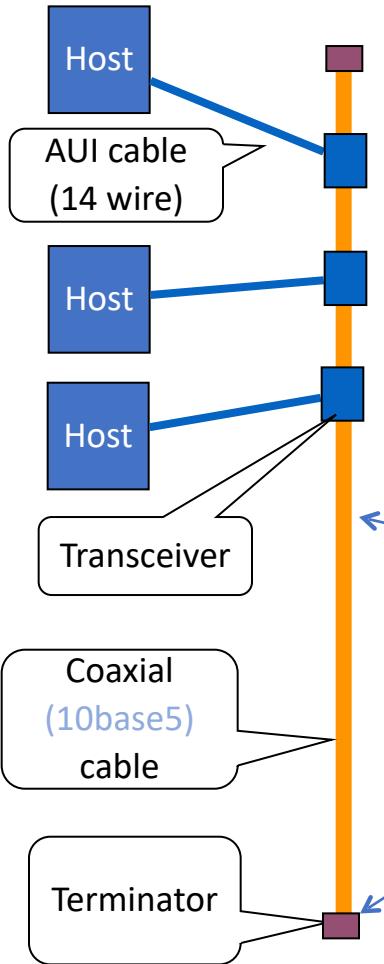
# Ethernet Topology (some history)



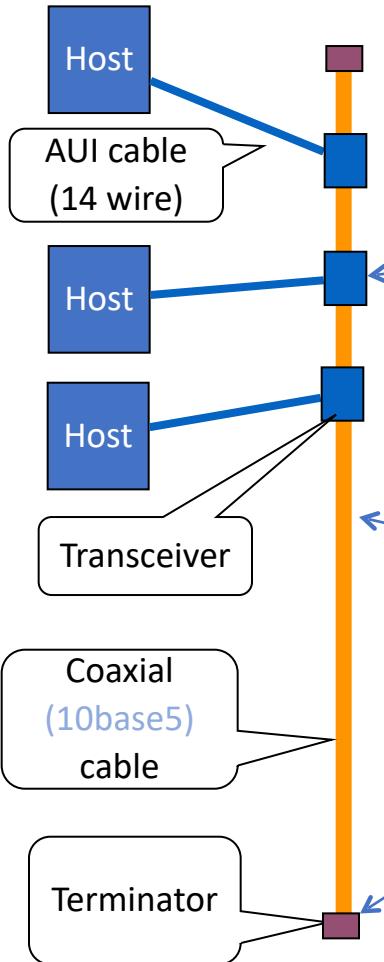
# Ethernet Topology (some history)



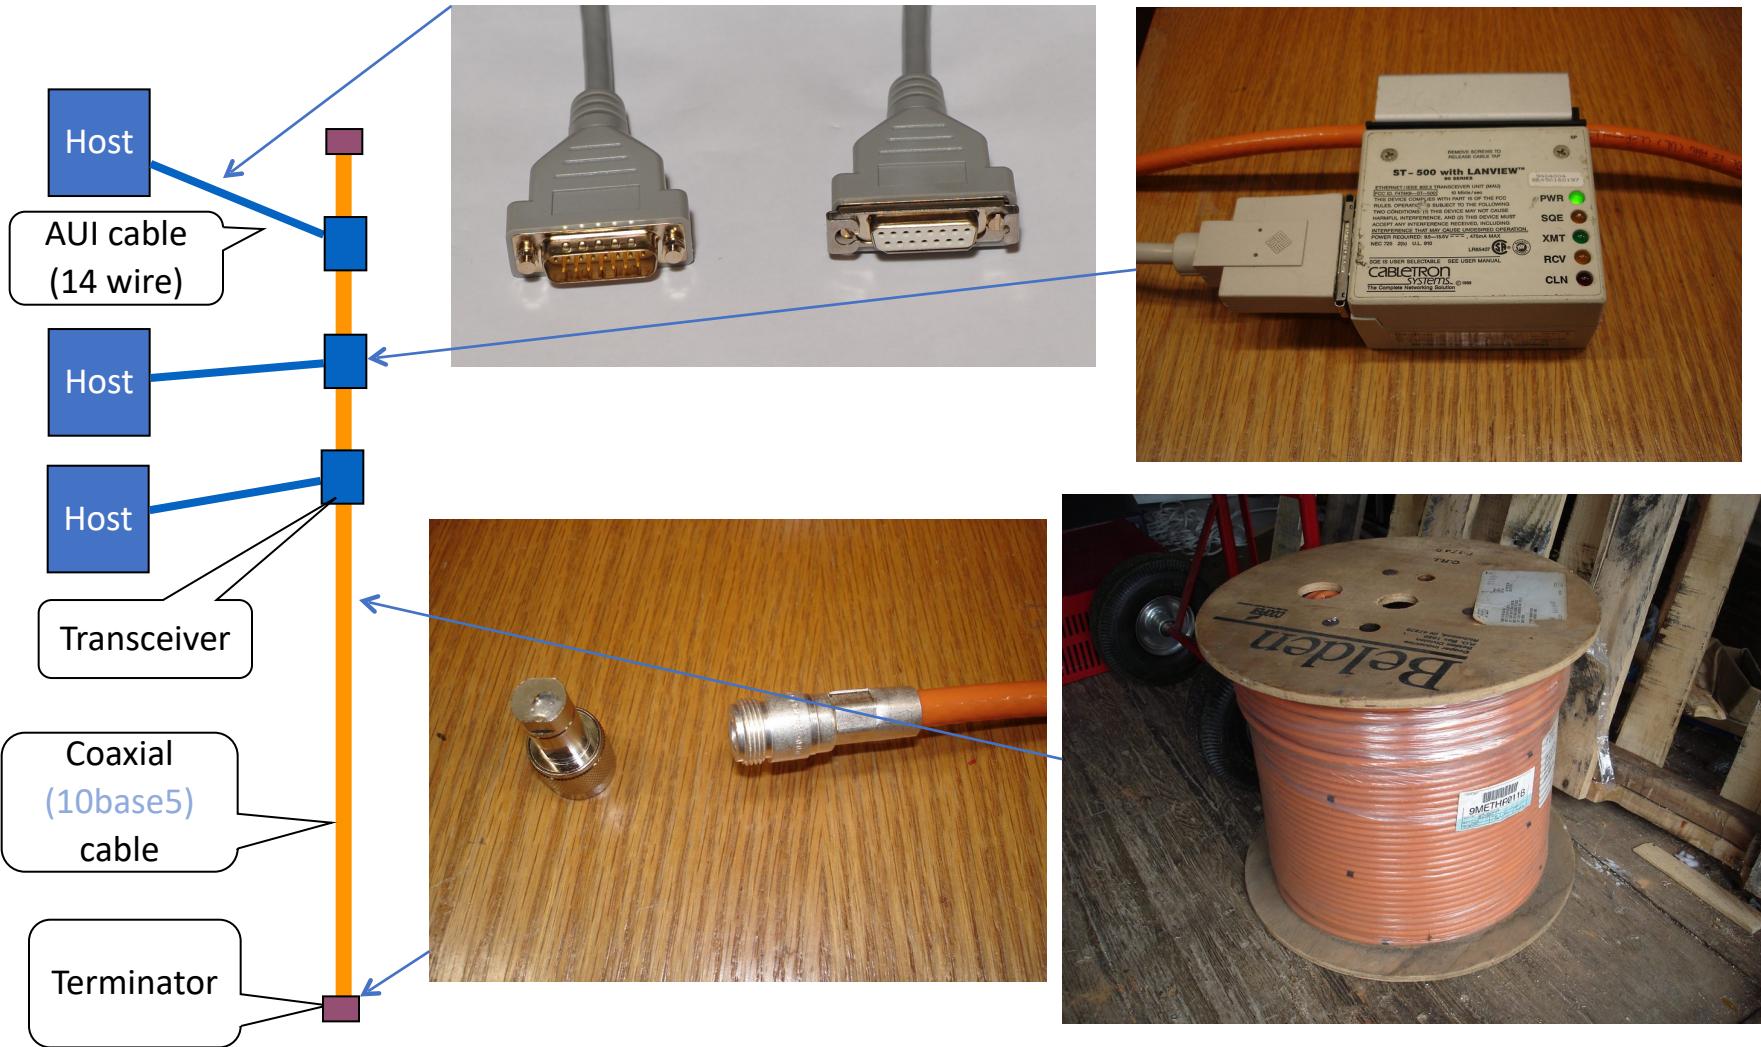
# Ethernet Topology (some history)



# Ethernet Topology (some history)



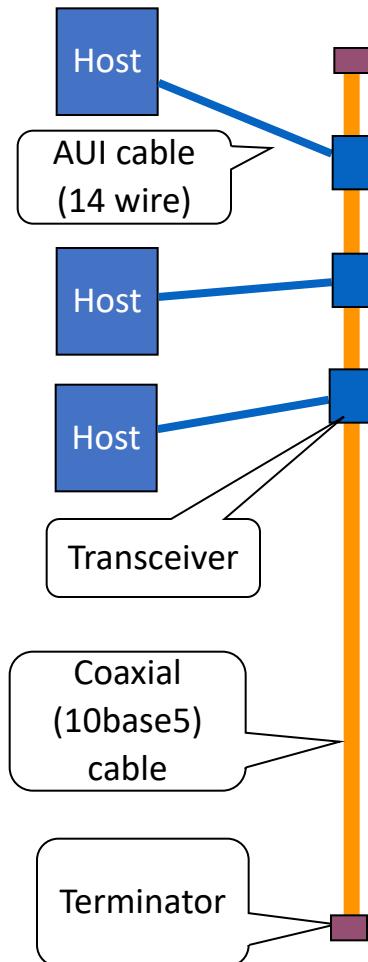
# Ethernet Topology (some history)



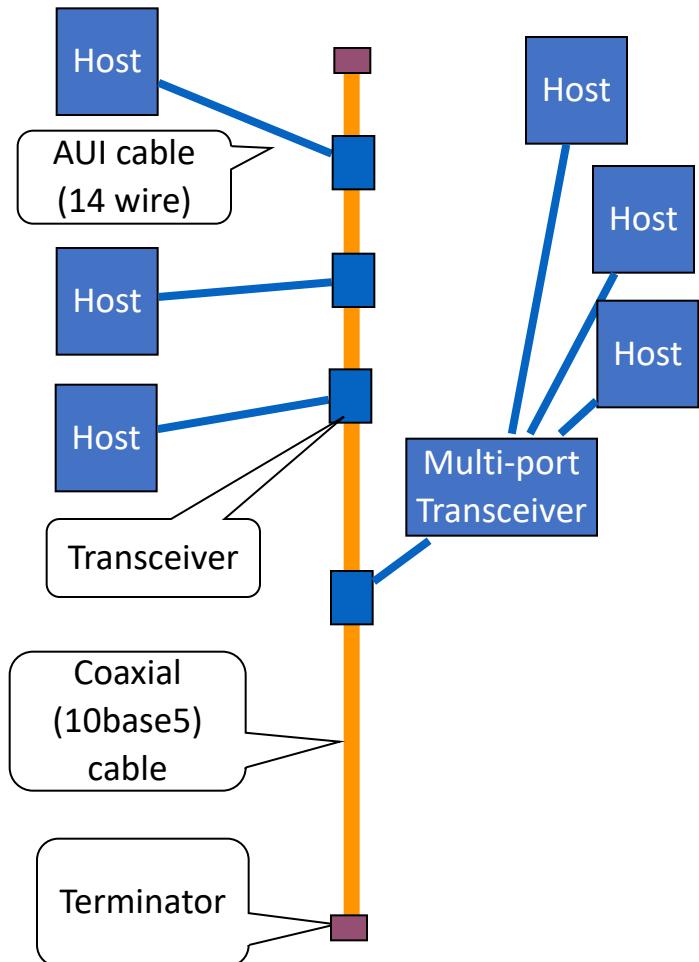
# Ethernet Topology (some history)

---

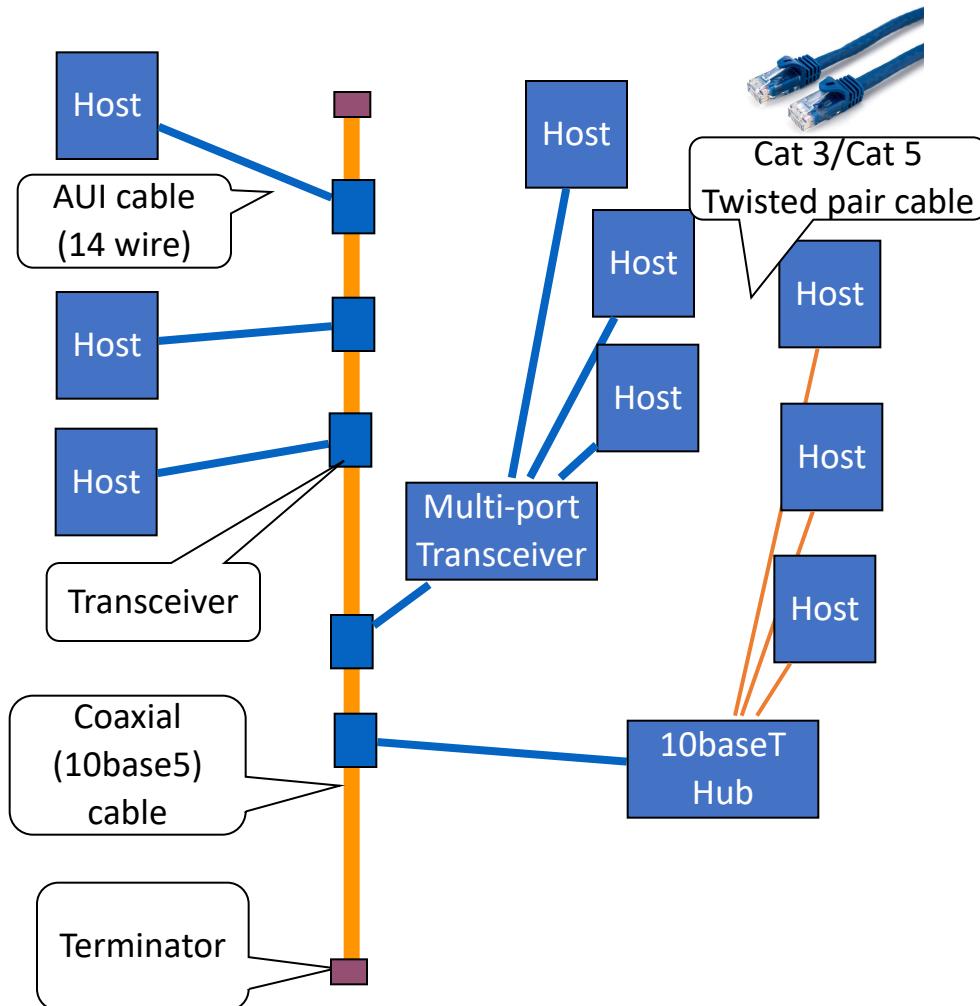
# Ethernet Topology (some history)



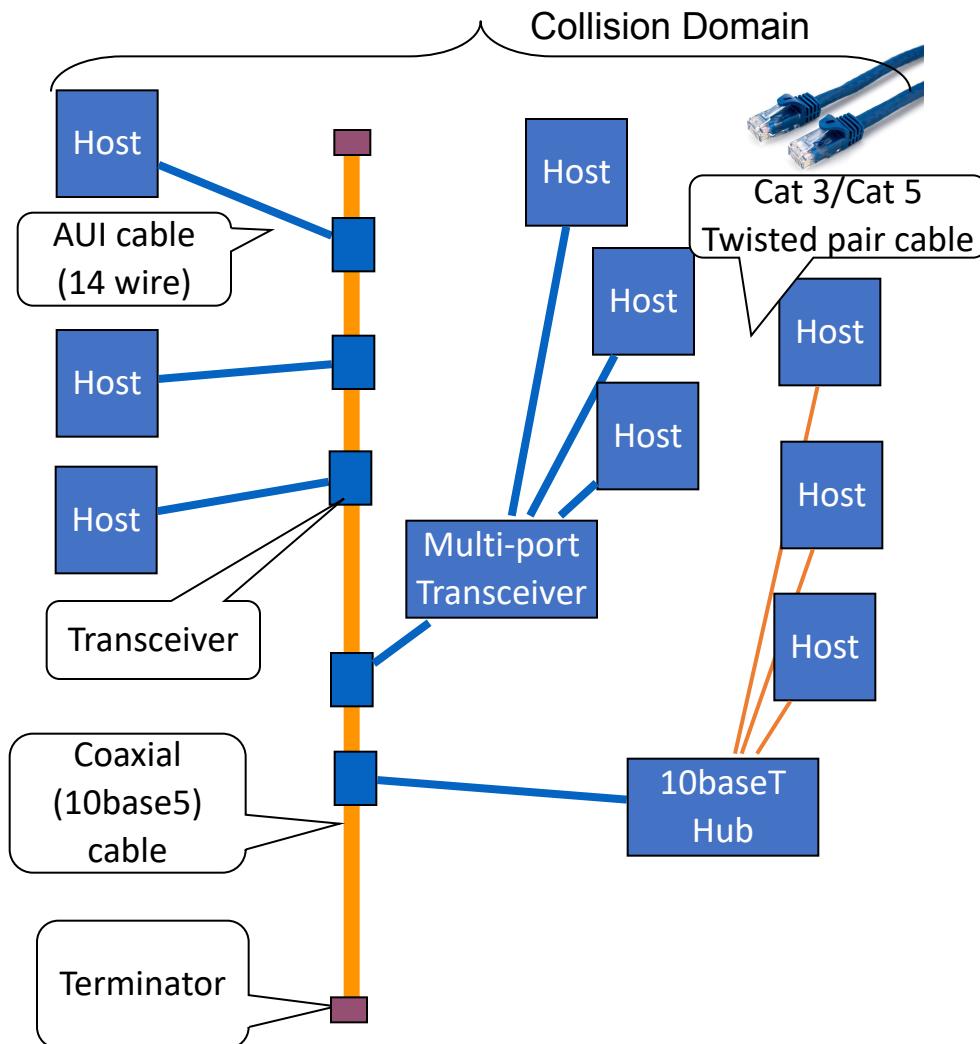
# Ethernet Topology (some history)



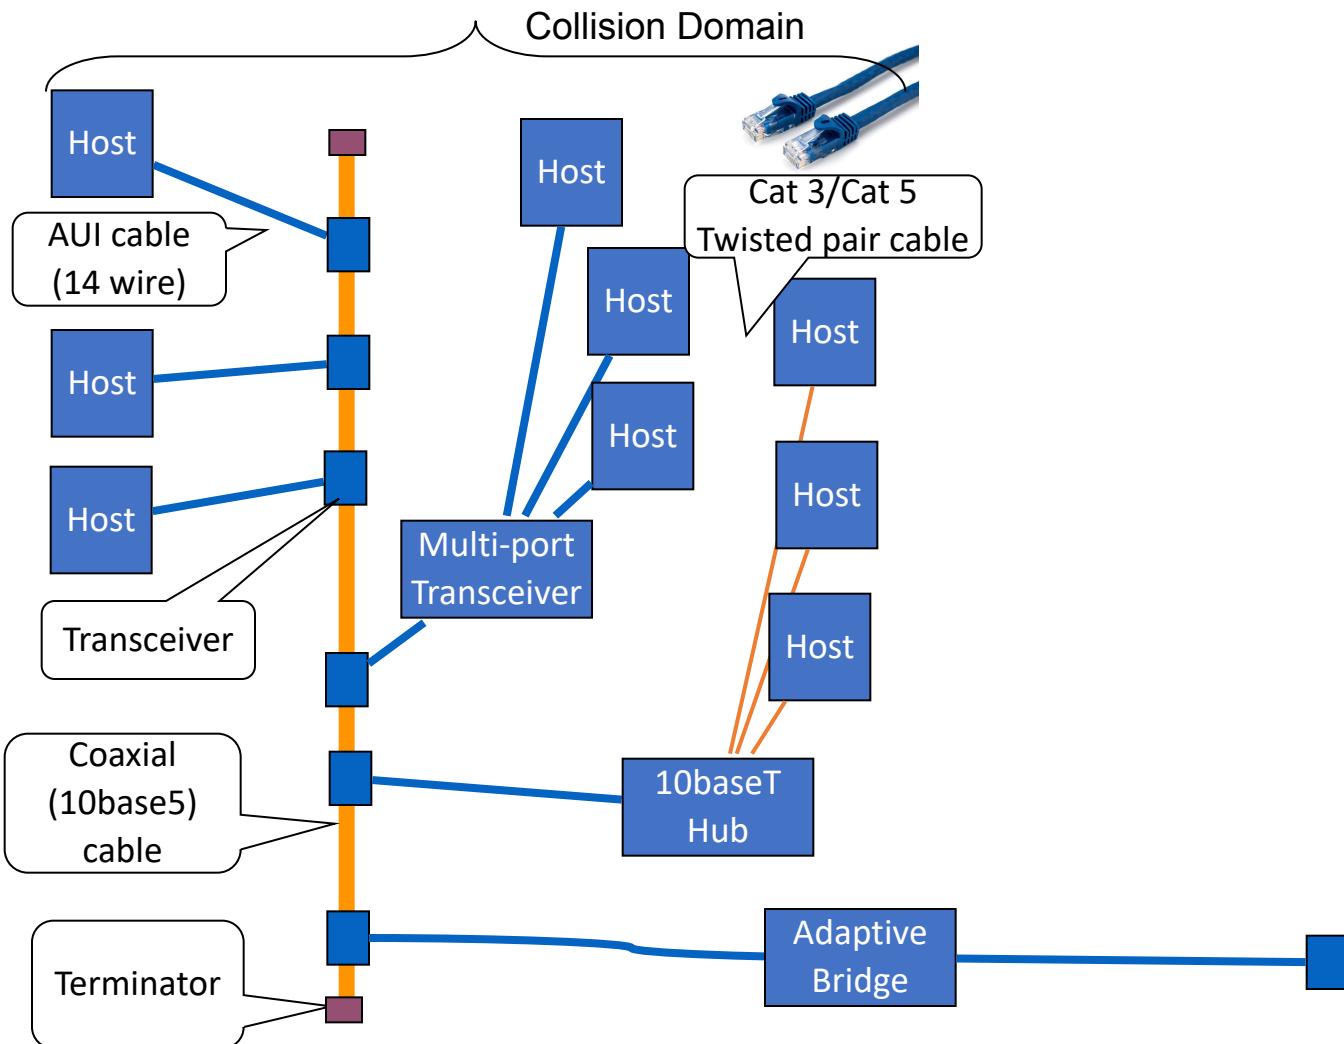
# Ethernet Topology (some history)



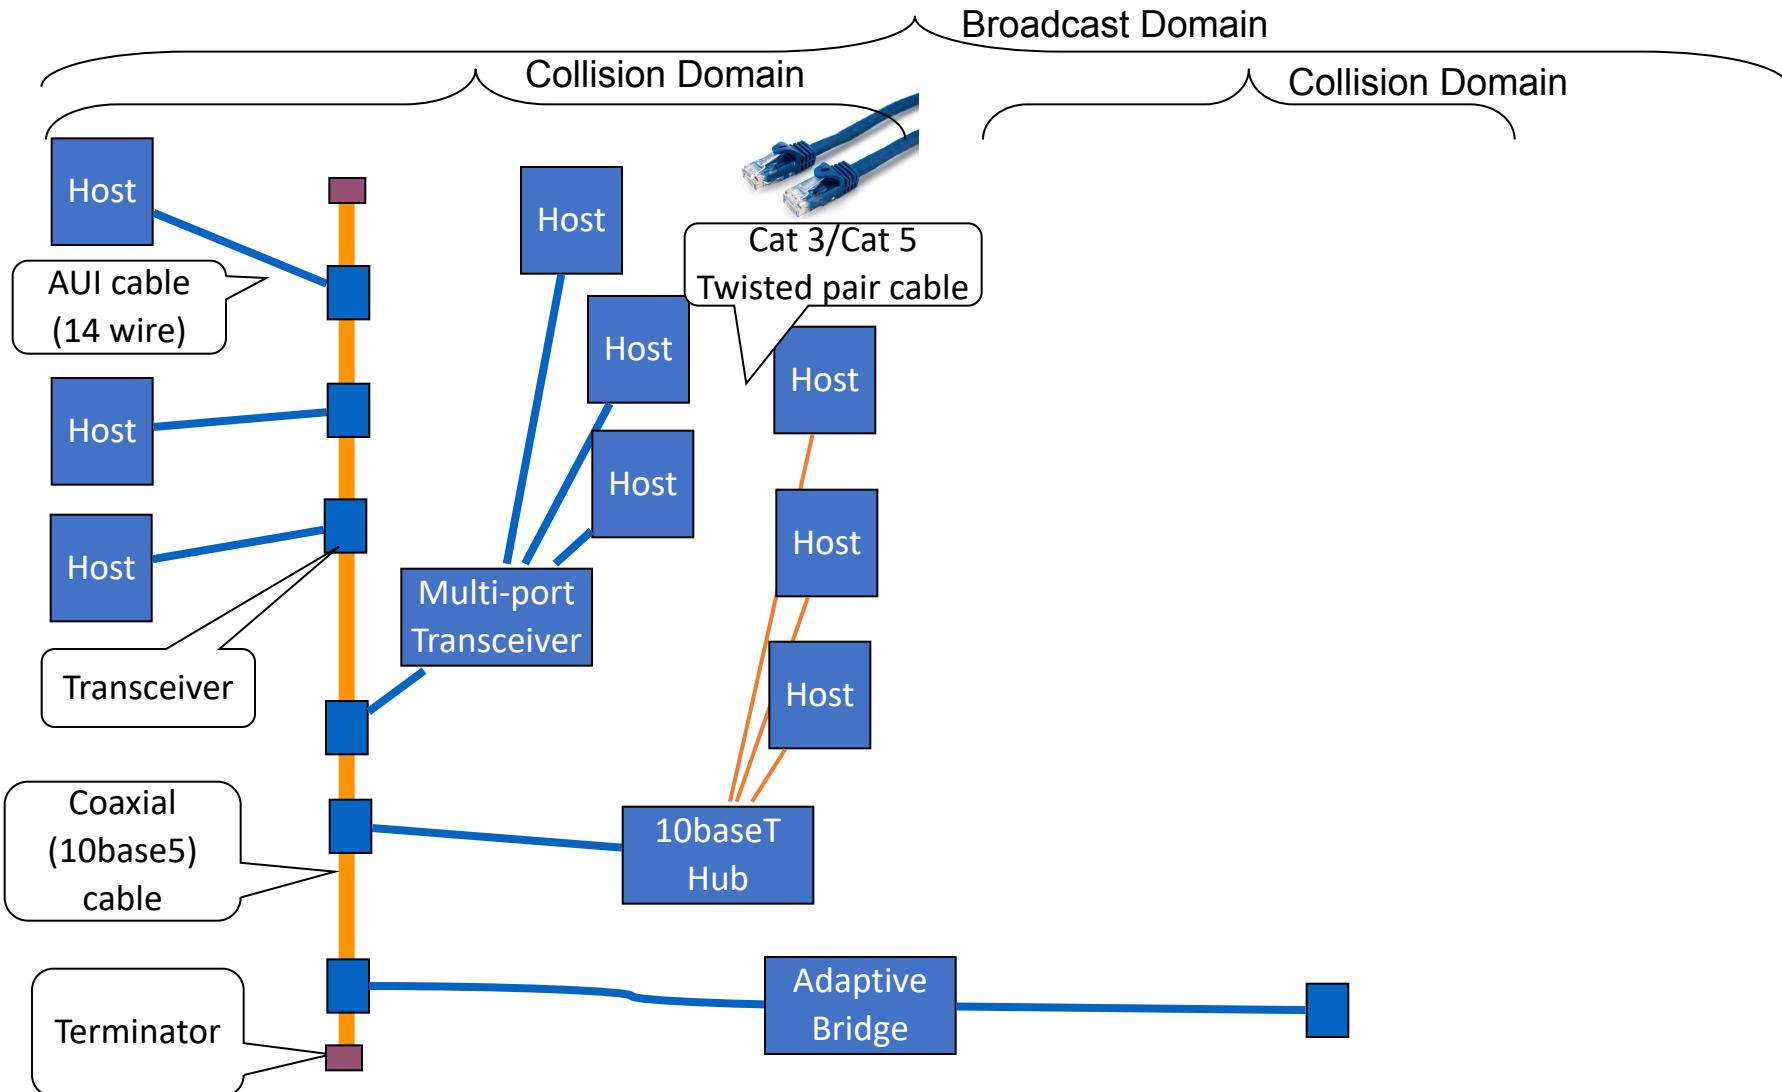
# Ethernet Topology (some history)



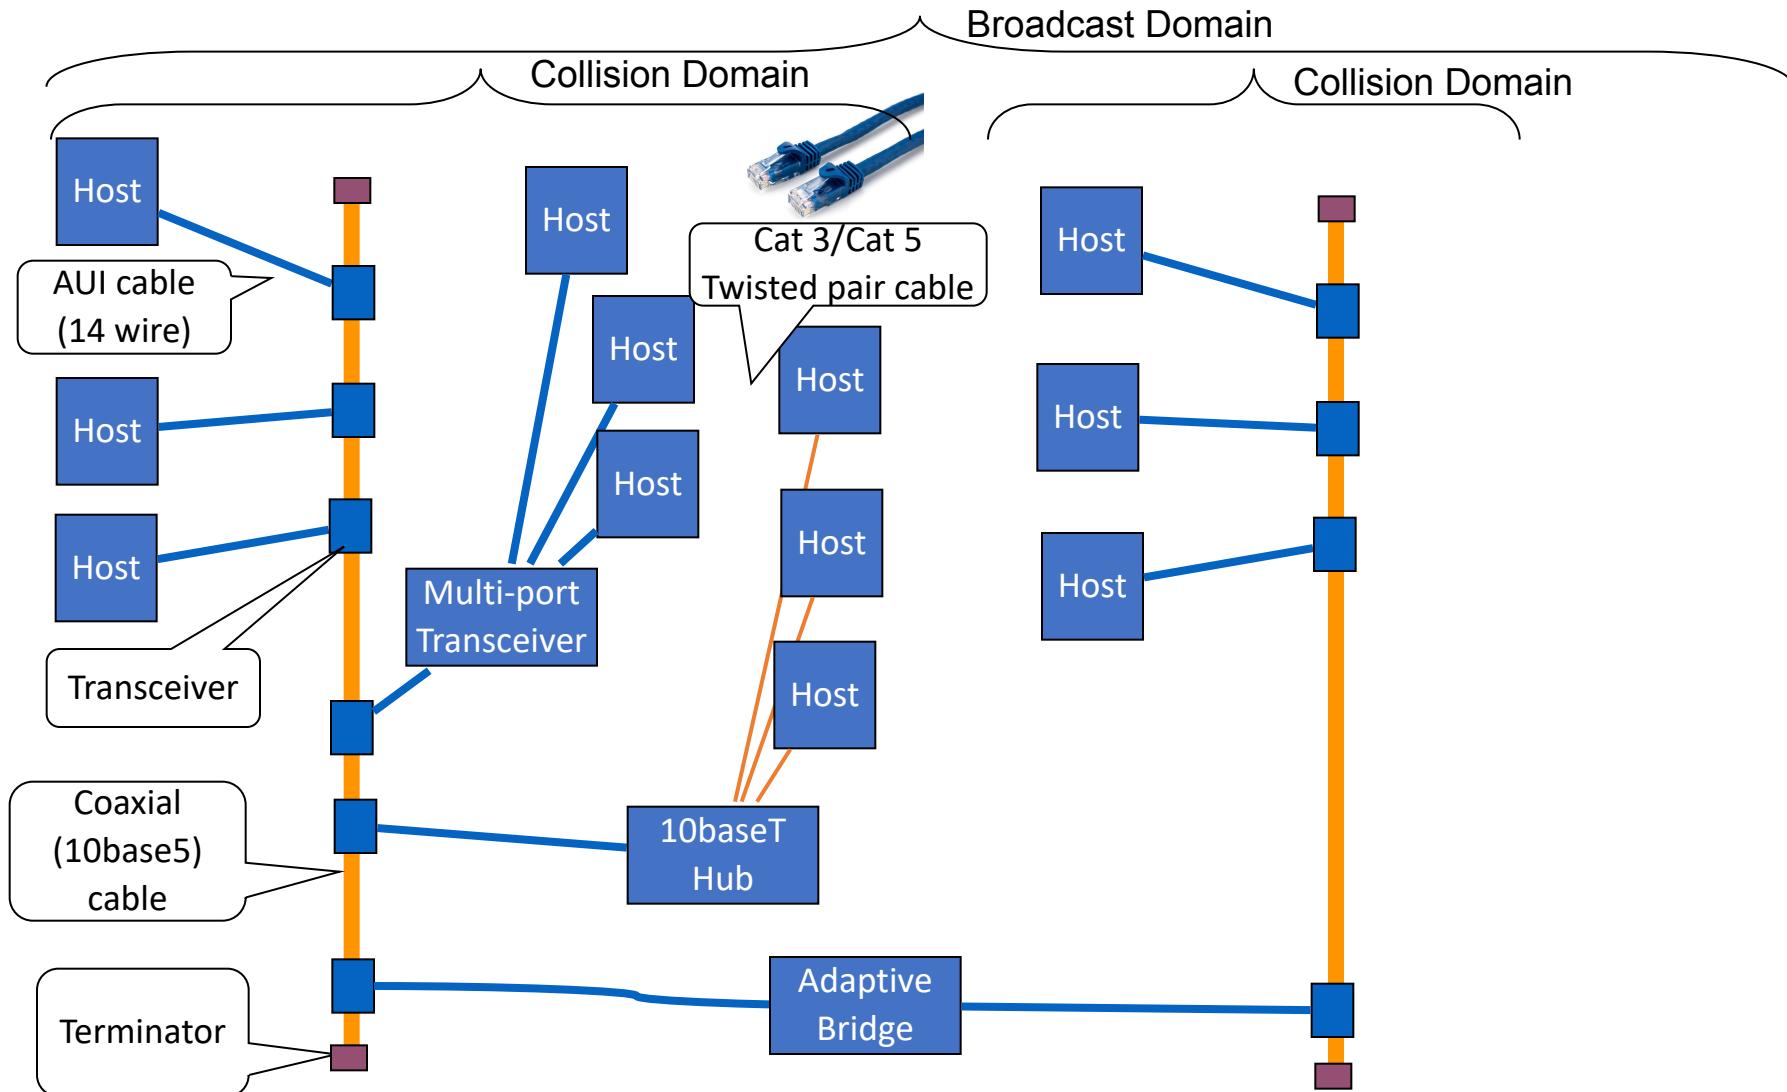
# Ethernet Topology (some history)



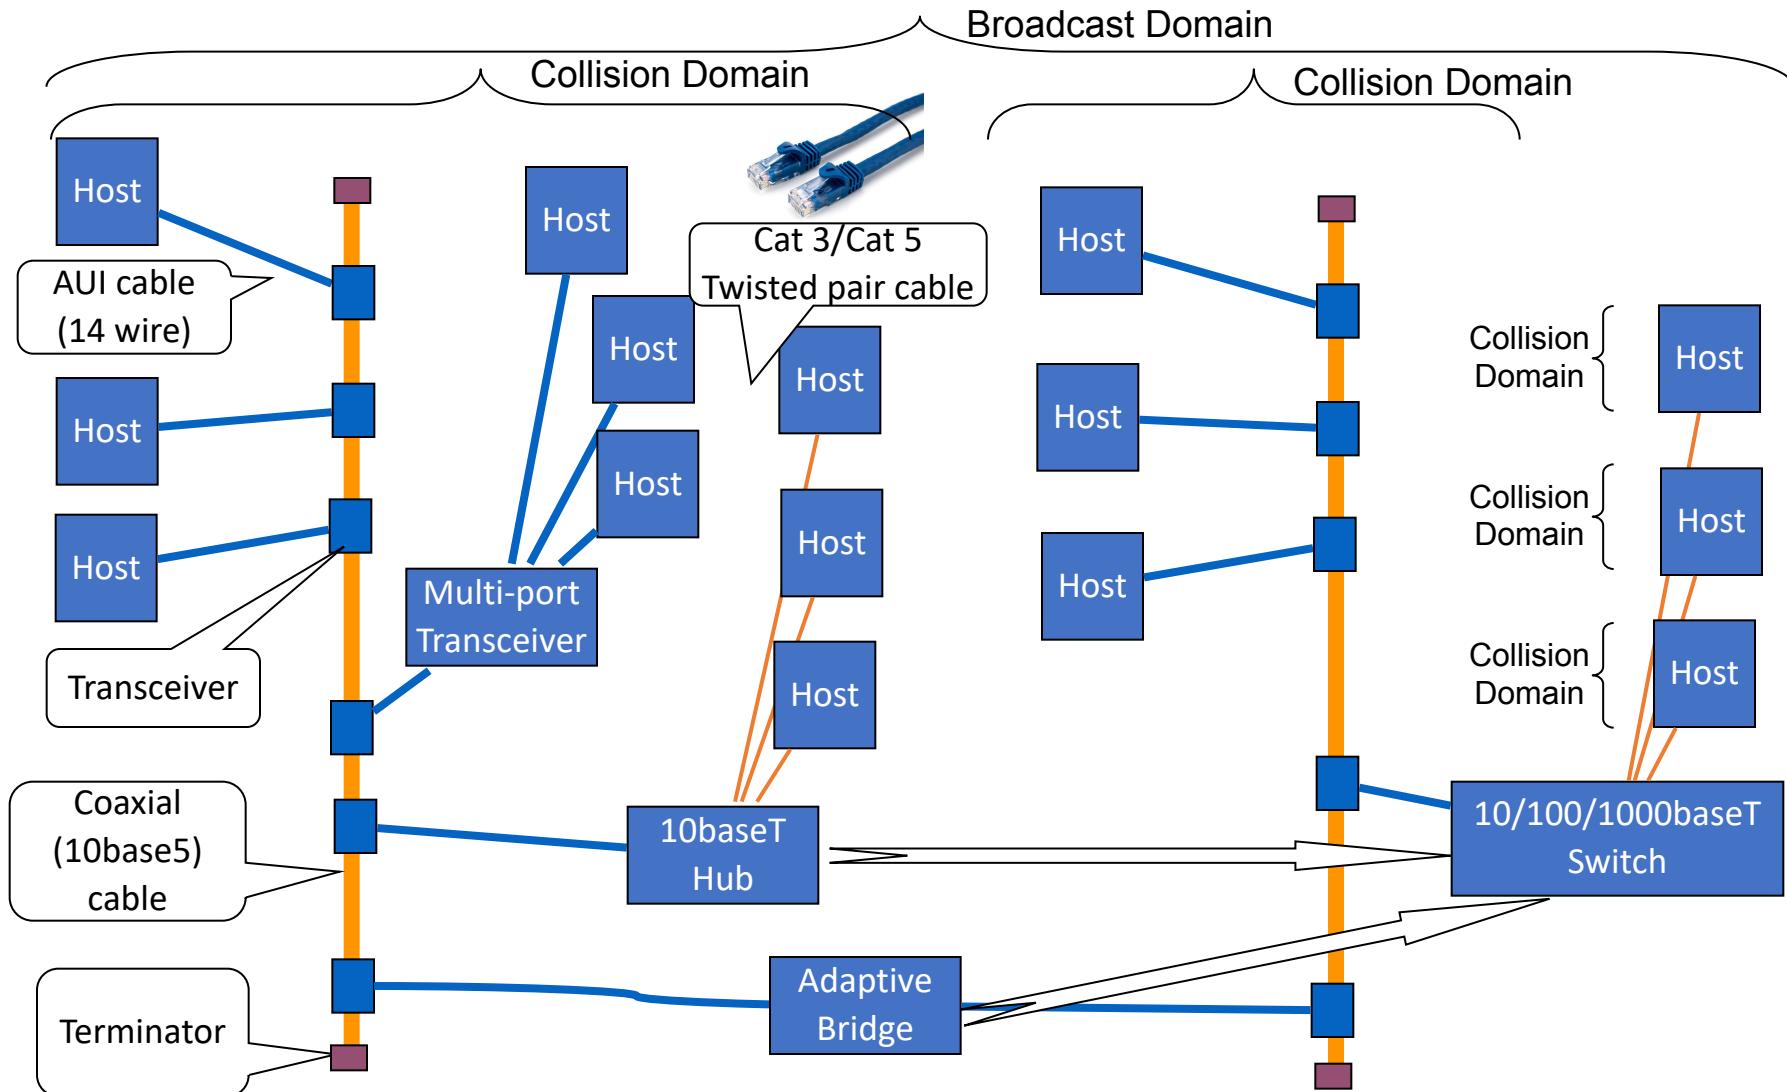
# Ethernet Topology (some history)



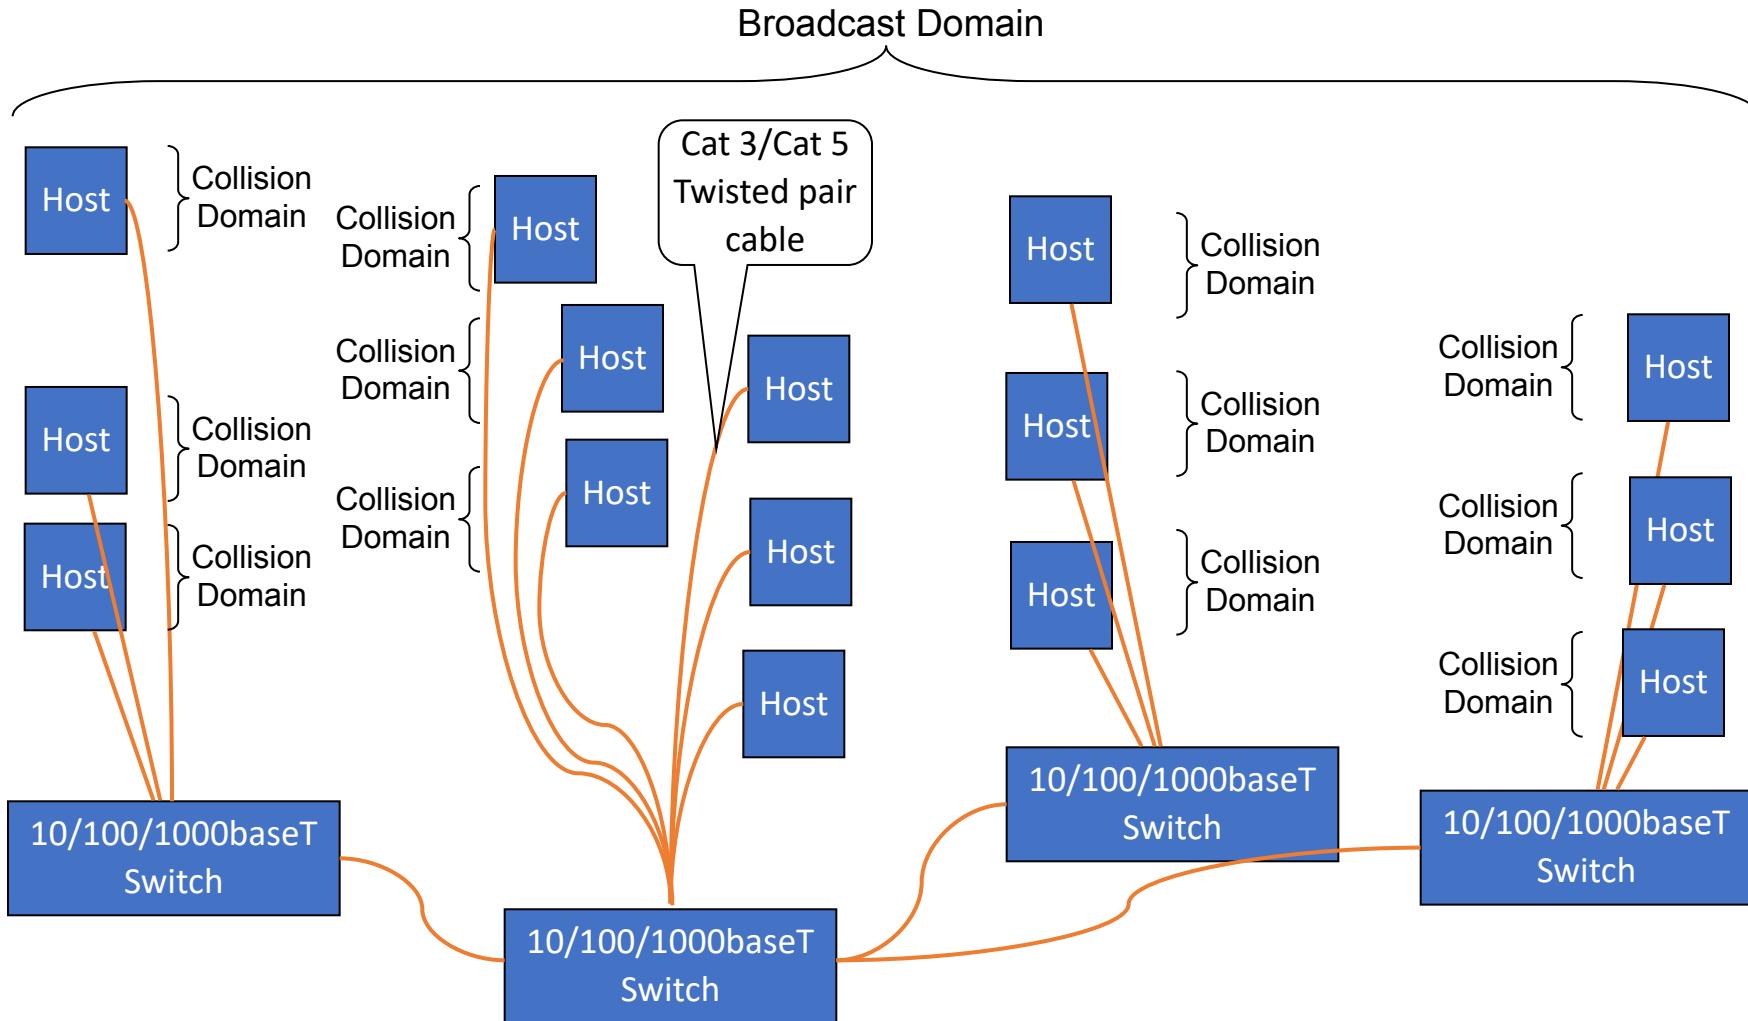
# Ethernet Topology (some history)



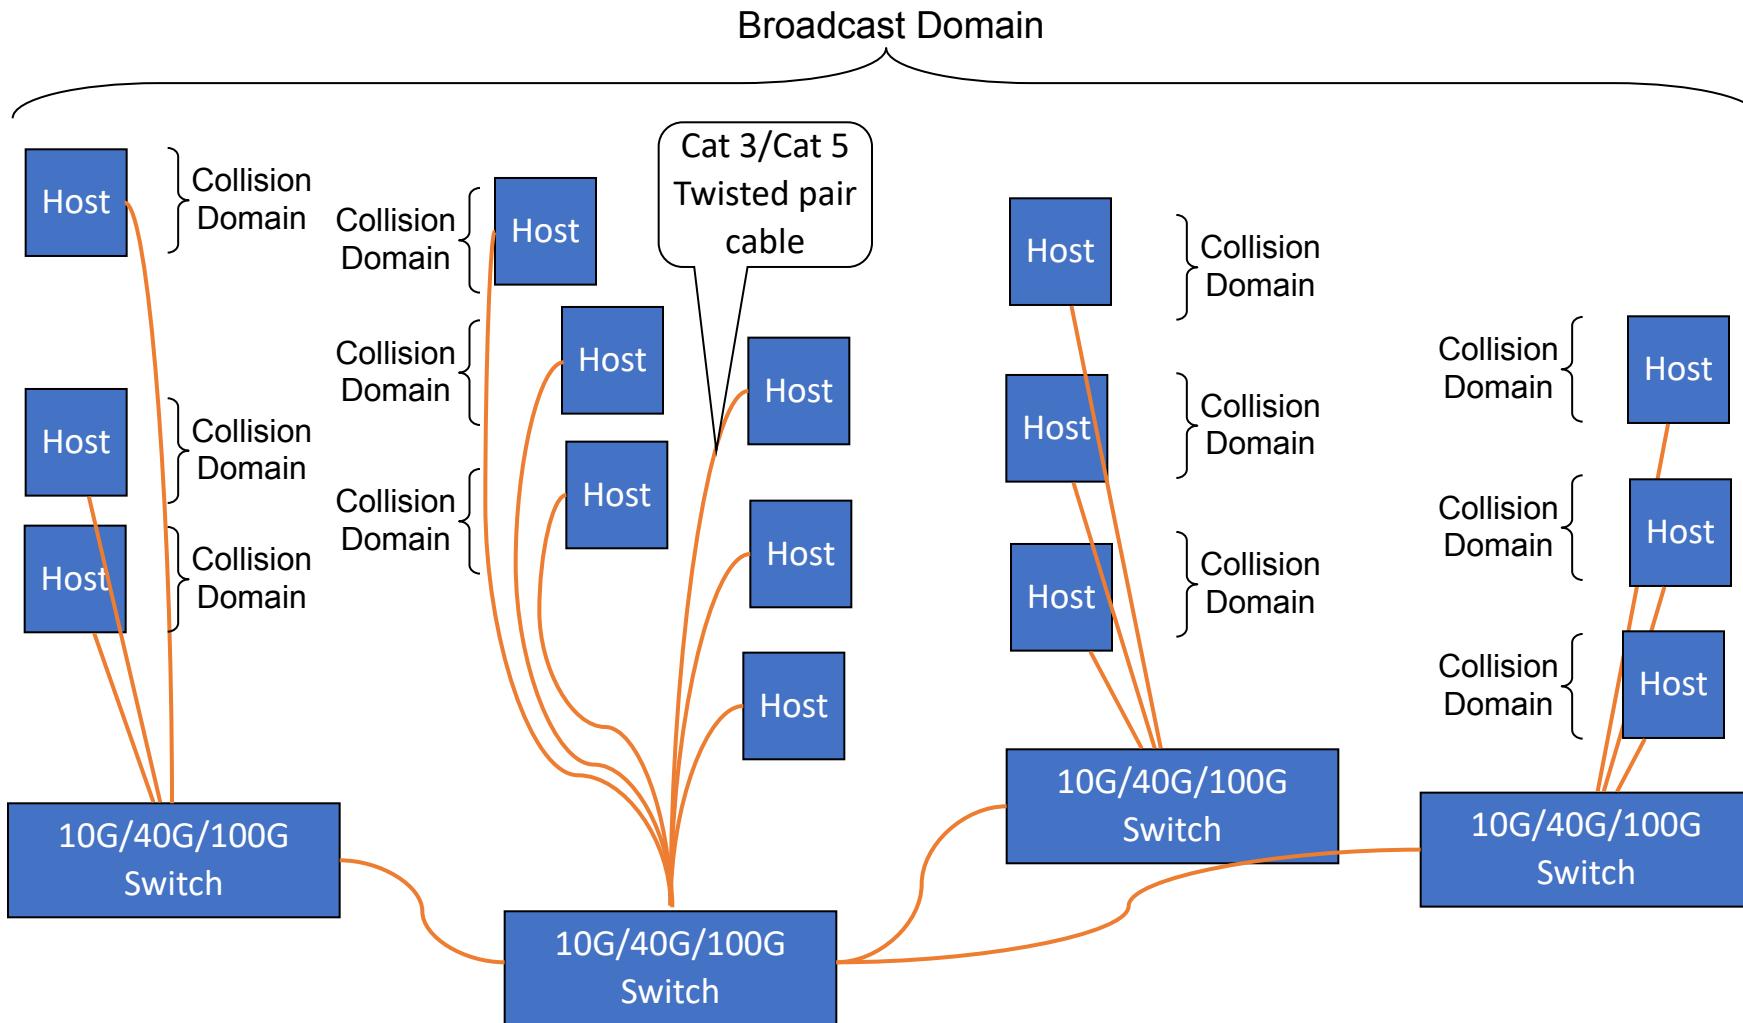
# Ethernet Topology (some history)



# Ethernet Topology (some history)



# Ethernet Topology (modern)



# How to deal with collisions during transmission?

---

CSMA/CD (Carrier Sense Multiple Access/Collision Detection)

# How to deal with collisions during transmission?

---

## CSMA/CD (Carrier Sense Multiple Access/Collision Detection)

- Check physical medium while transmitting to detect possible interference
- If interference, retry after exponential backoff
- Performance problem in large collision domains, alleviated by modern switches

# How to deal with collisions during transmission?

---

## CSMA/CD (Carrier Sense Multiple Access/Collision Detection)

- Check physical medium while transmitting to detect possible interference
- If interference, retry after exponential backoff
- Performance problem in large collision domains, alleviated by modern switches

## CSMA/CA (Collision Avoidance)

- First handshake with destination to ensure they are ready to receive transmission
- Still needed in wireless networks (WiFi)

# Ethernet Bridging Algorithm (simplified)

---

Bridge Table		
MAC Addr	Last Seen on Port	Last Used
<i>002080 00324a</i>	<i>FastEthernet 1/4</i>	<i>12 sec ago</i>
<i>0003ba 51733a</i>	<i>FastEthernet 2/9</i>	<i>72 sec ago</i>
...	...	...

- When a frame enters the bridge (switch):

# Ethernet Bridging Algorithm (simplified)

Bridge Table		
MAC Addr	Last Seen on Port	Last Used
<i>002080 00324a</i>	<i>FastEthernet 1/4</i>	<i>12 sec ago</i>
<i>0003ba 51733a</i>	<i>FastEthernet 2/9</i>	<i>72 sec ago</i>
...	...	...

- When a frame enters the bridge (switch):
  - If the **destination** MAC address is in the bridge table
    - Send the frame out the port listed in the bridge table entry
  - Otherwise
    - Flood the frame out all ports except the ingress port
  - Put the **source** MAC address in the bridge table, if necessary, along with the frame ingress port and mark that entry “recently used”

# Ethernet Bridging Algorithm (simplified)

Bridge Table		
MAC Addr	Last Seen on Port	Last Used
<i>002080 00324a</i>	<i>FastEthernet 1/4</i>	<i>12 sec ago</i>
<i>0003ba 51733a</i>	<i>FastEthernet 2/9</i>	<i>72 sec ago</i>
...	...	...

- When a frame enters the bridge (switch):
  - If the **destination** MAC address is in the bridge table
    - Send the frame out the port listed in the bridge table entry
  - Otherwise
    - Flood the frame out all ports except the ingress port
  - Put the **source** MAC address in the bridge table, if necessary, along with the frame ingress port and mark that entry “recently used”
  - Remove any “stale” entries in the bridge table

# Something Is Still Missing

---

# Something Is Still Missing

---

- Now more than one message can be in transit at a time!

# Something Is Still Missing

---

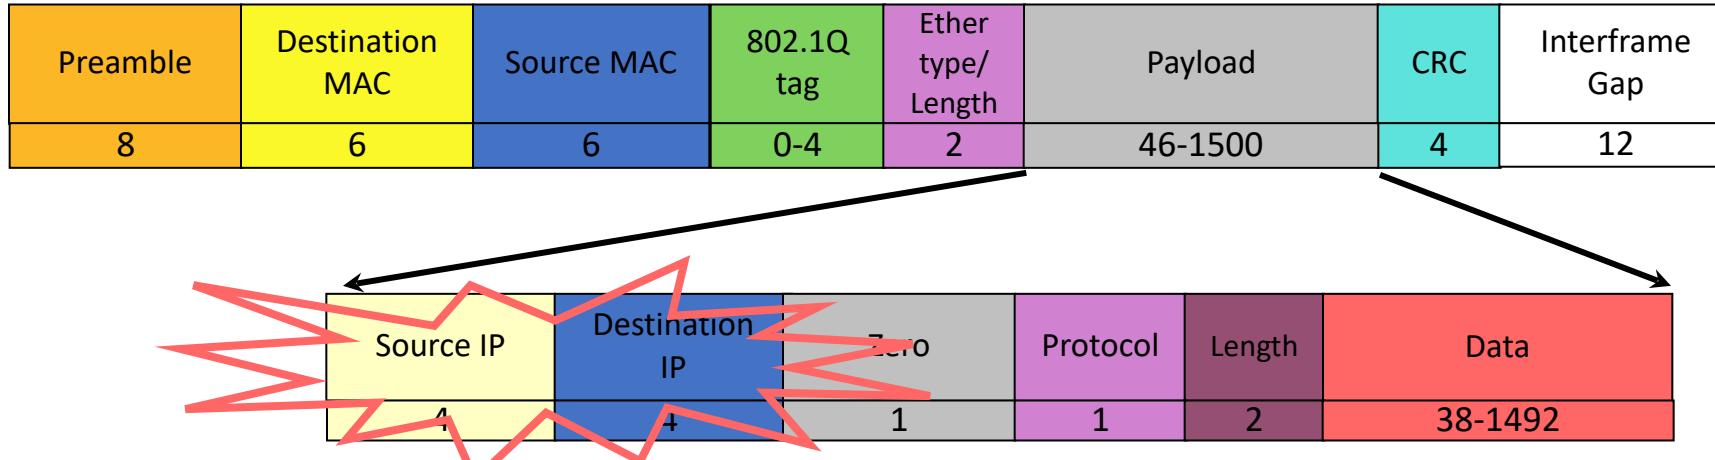
- Now more than one message can be in transit at a time!
- We've only talked about MAC addresses.

# Something Is Still Missing

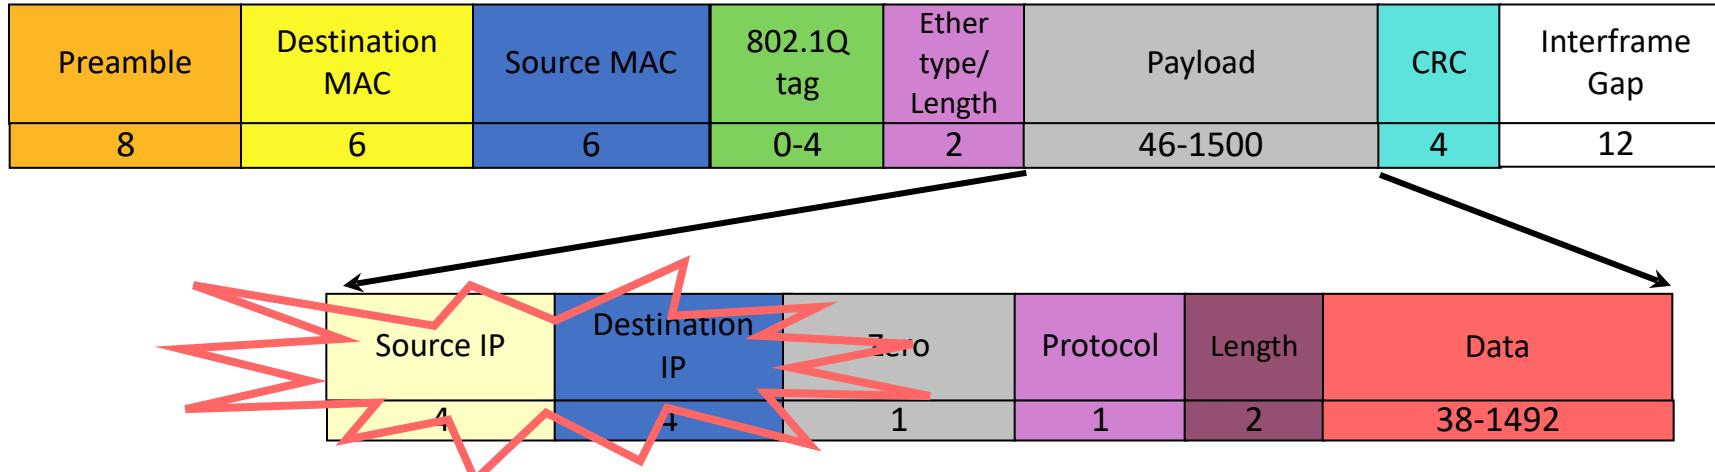
---

- Now more than one message can be in transit at a time!
- We've only talked about MAC addresses.
- Where are the IP addresses?

# Layer 3 (IP) packets

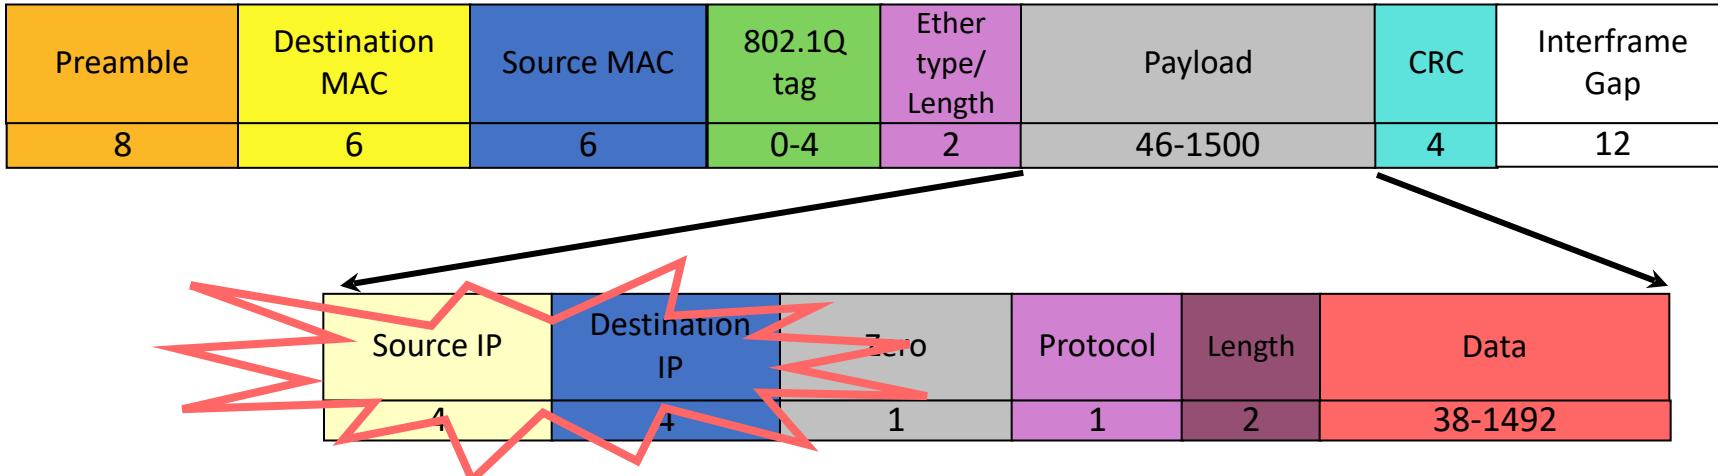


# Layer 3 (IP) packets



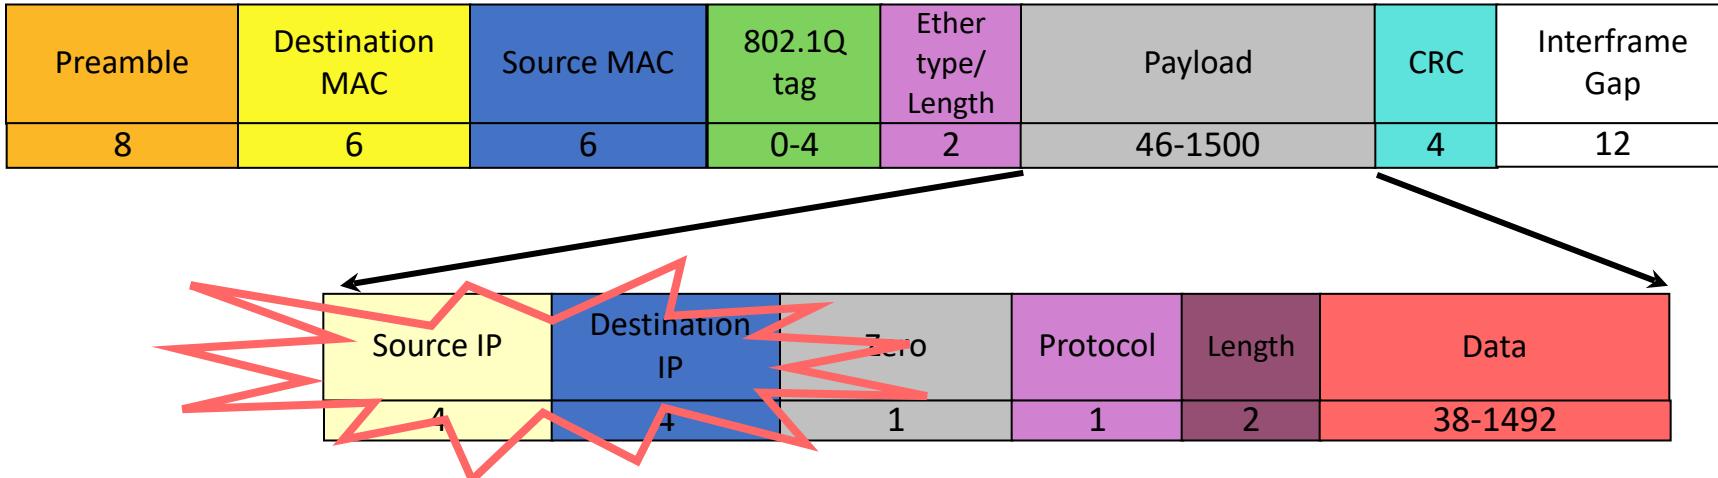
- IP addresses appear in the Ethernet frame payload, not the header

# Layer 3 (IP) packets



- IP addresses appear in the Ethernet frame payload, not the header
- But Ethernet hardware only examines the MAC addresses in the Ethernet frame(!)

# Layer 3 (IP) packets



- IP addresses appear in the Ethernet frame payload, not the header
- But Ethernet hardware only examines the MAC addresses in the Ethernet frame(!)
- How do we connect these two layers?

# ARP

---

- Address Resolution Protocol

# ARP

---

- Address Resolution Protocol
- Used with Ethernet to discover IP addresses

# ARP

---

- Address Resolution Protocol
- Used with Ethernet to discover IP addresses
- How?

# ARP

---

- Address Resolution Protocol
- Used with Ethernet to discover IP addresses
- How?
  - IP stack needs to send to IP address a, but doesn't know what destination MAC address to use.

# ARP

---

- Address Resolution Protocol
- Used with Ethernet to discover IP addresses
- How?
  - IP stack needs to send to IP address a, but doesn't know what destination MAC address to use.
  - Broadcast an ARP-request message: "Who has IP address a?"

# ARP

---

- Address Resolution Protocol
- Used with Ethernet to discover IP addresses
- How?
  - IP stack needs to send to IP address a, but doesn't know what destination MAC address to use.
  - Broadcast an ARP-request message: "Who has IP address a?"
  - Wait on unicast ARP-response: "I have IP address a."

# ARP

---

- Address Resolution Protocol
- Used with Ethernet to discover IP addresses
- How?
  - IP stack needs to send to IP address a, but doesn't know what destination MAC address to use.
  - Broadcast an ARP-request message: "Who has IP address a?"
  - Wait on unicast ARP-response: "I have IP address a."
  - Store the response in a local *ARP Table*, discarding entries when they become stale.

# ARP

---

- Address Resolution Protocol
- Used with Ethernet to discover IP addresses
- How?
  - IP stack needs to send to IP address a, but doesn't know what destination MAC address to use.
  - Broadcast an ARP-request message: "Who has IP address a?"
  - Wait on unicast ARP-response: "I have IP address a."
  - Store the response in a local *ARP Table*, discarding entries when they become stale.
  - Use the corresponding entry in the ARP table to set the destination MAC address of the frame.

# ARP Example

---

```
bash-3.00# arp -an
Net to Media Table: IPv4
Device      IP Address          Mask           Flags   Phys Addr
-----  -----
bge0        130.207.165.229    255.255.255.255   o       00:50:56:91:4a:2b
bge0        130.207.165.248    255.255.255.255   o       00:14:4f:f2:8c:ce
bge0        130.207.165.240    255.255.255.255   o       00:50:56:91:62:be
bge0        130.207.165.242    255.255.255.255   o       00:14:4f:21:14:f6
bge0        130.207.165.245    255.255.255.255   o       00:14:4f:94:d7:0c
bge0        130.207.165.194    255.255.255.255   o       00:50:56:91:59:25
bge0        130.207.165.197    255.255.255.255   o       00:50:56:91:46:a3
bge0        130.207.165.221    255.255.255.255   o       00:50:56:91:4c:af
bge0        130.207.165.163    255.255.255.255   o       00:14:4f:7d:7f:58
bge0        130.207.165.137    255.255.255.255
bge0        130.207.165.138    255.255.255.255
bge0        130.207.165.158    255.255.255.255   o       00:14:4f:1e:26:ec
bge0        130.207.165.96     255.255.255.255   o       00:1b:24:1d:eb:7c
bge0        130.207.165.103    255.255.255.255   o       00:50:56:a4:7e:df
bge0        130.207.165.122    255.255.255.255   o       00:50:56:91:43:f8
```

# VLANs

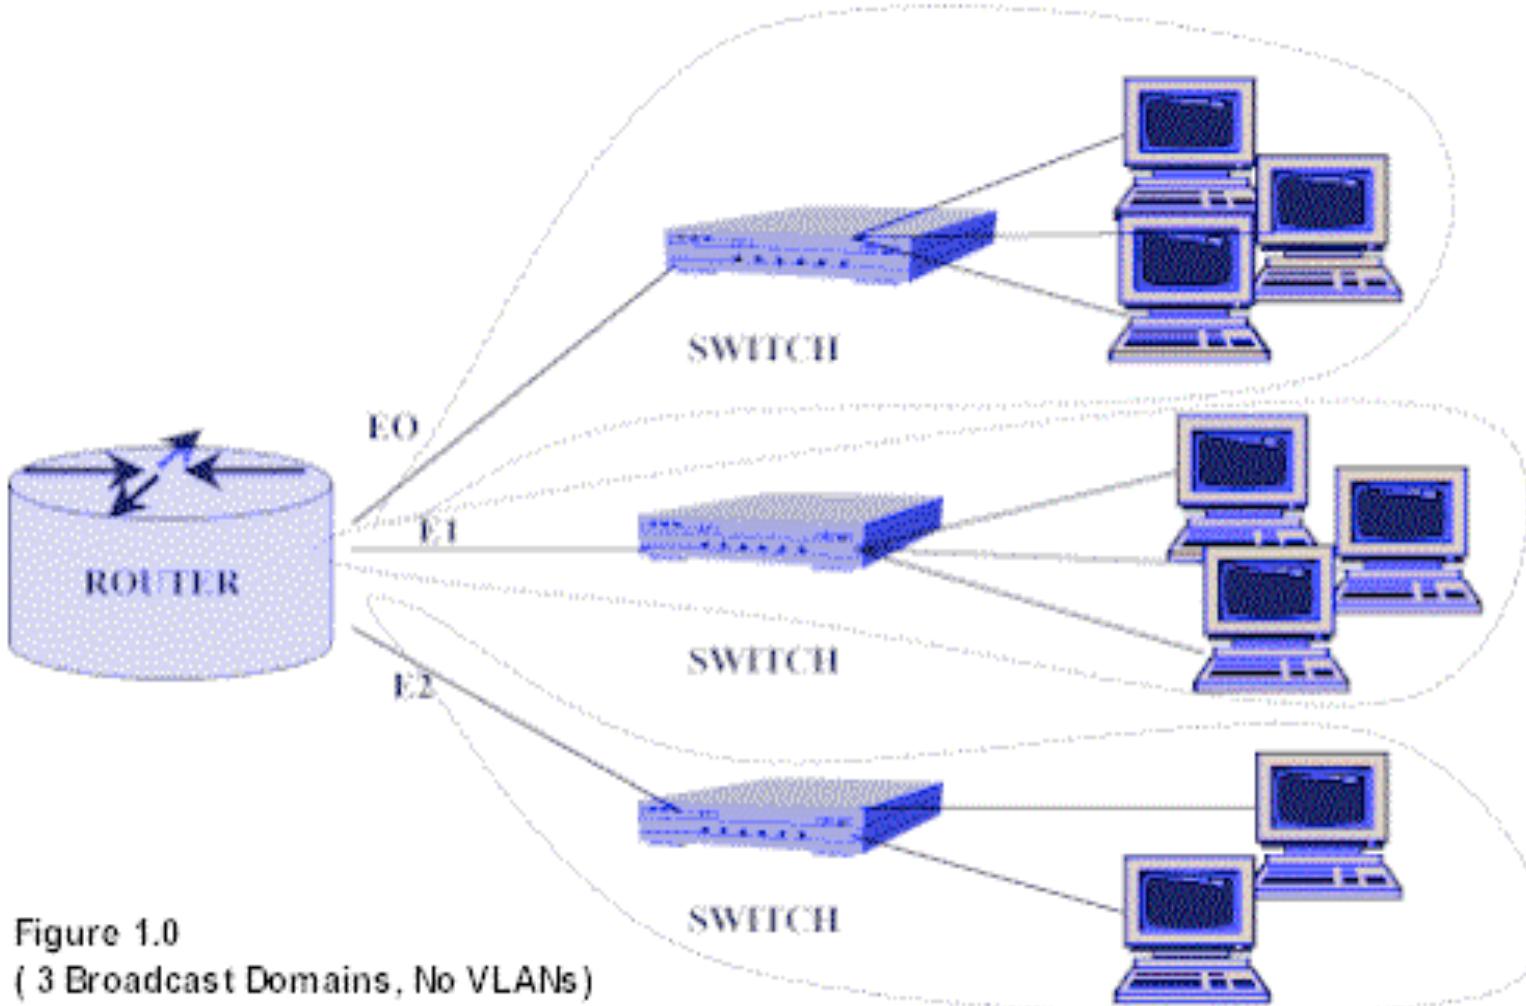
---

# VLANs

---

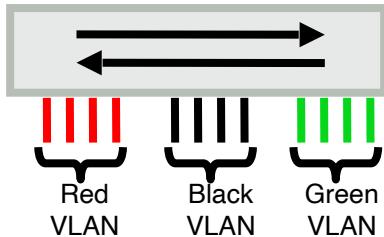
- Virtual LAN, a.k.a. VLAN
  - A group of hosts that can communicate as if they were attached to the same **broadcast domain**, regardless of their physical location.
  - A VLAN has the same attributes as a physical LAN, but it allows for hosts to be grouped together even if they are not connected to the same network switch.
  - Network reconfiguration can be done through software instead of physically relocating devices and wires.

# Geographic Addressing



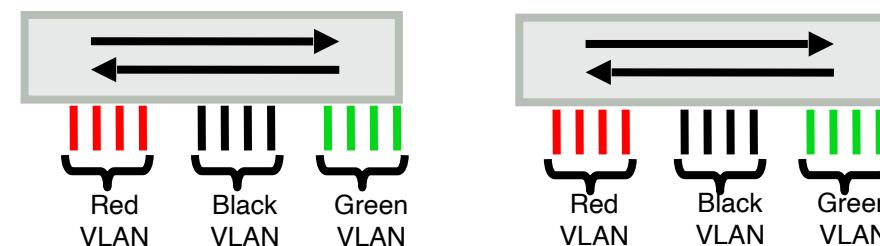
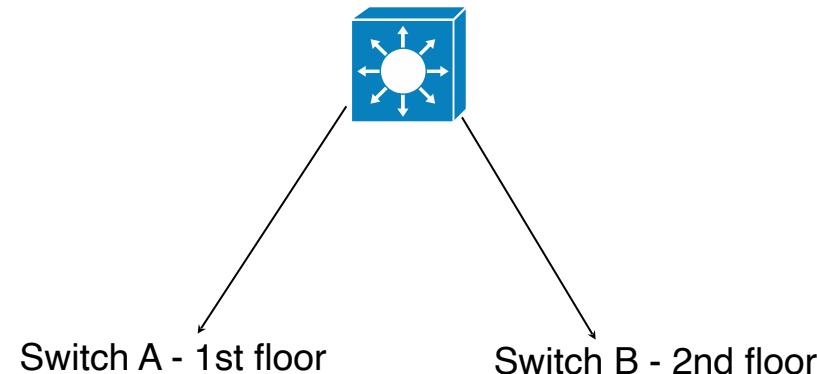
# After VLANs

Switch A



- Each logical VLAN is like a separate physical bridge

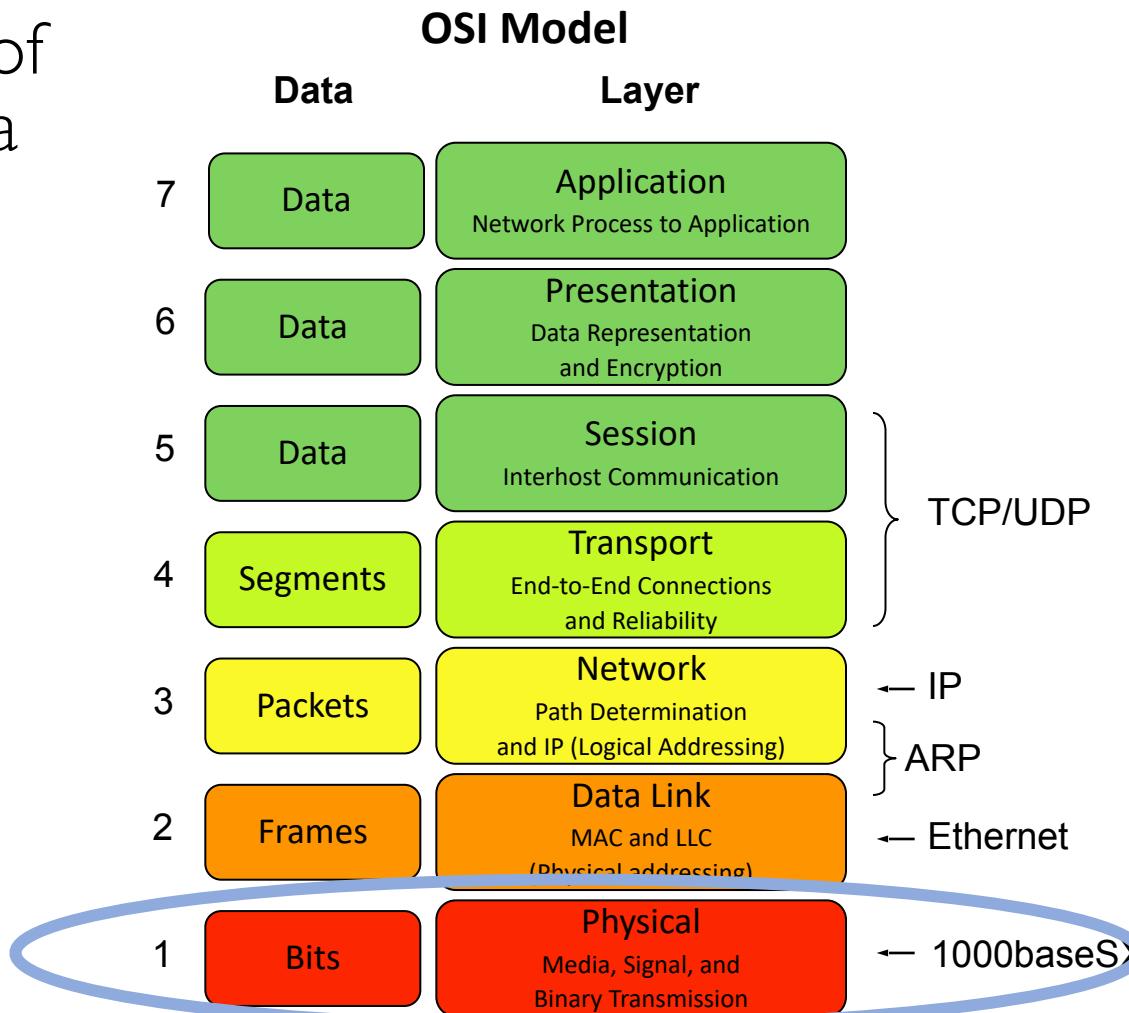
- College of Engineering (red)
- College of Computing (black)
- Classroom network (green)



- Each logical VLAN is like a separate physical bridge
- VLANs can span across multiple switches

# Now let's look at the last layer

- Ethernet standards allow use of several types of physical media

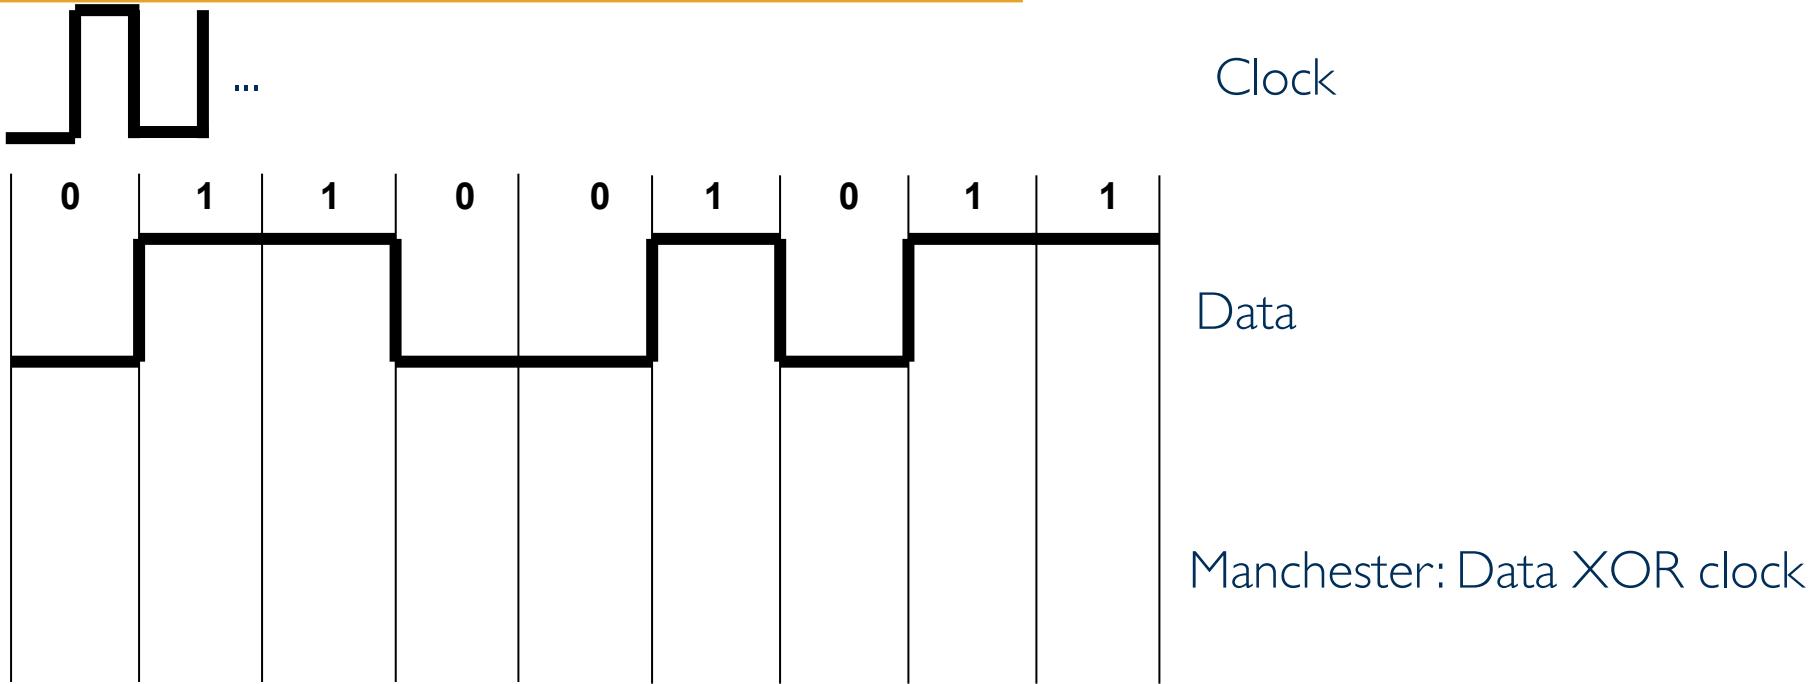


# Ethernet physical media standards

---

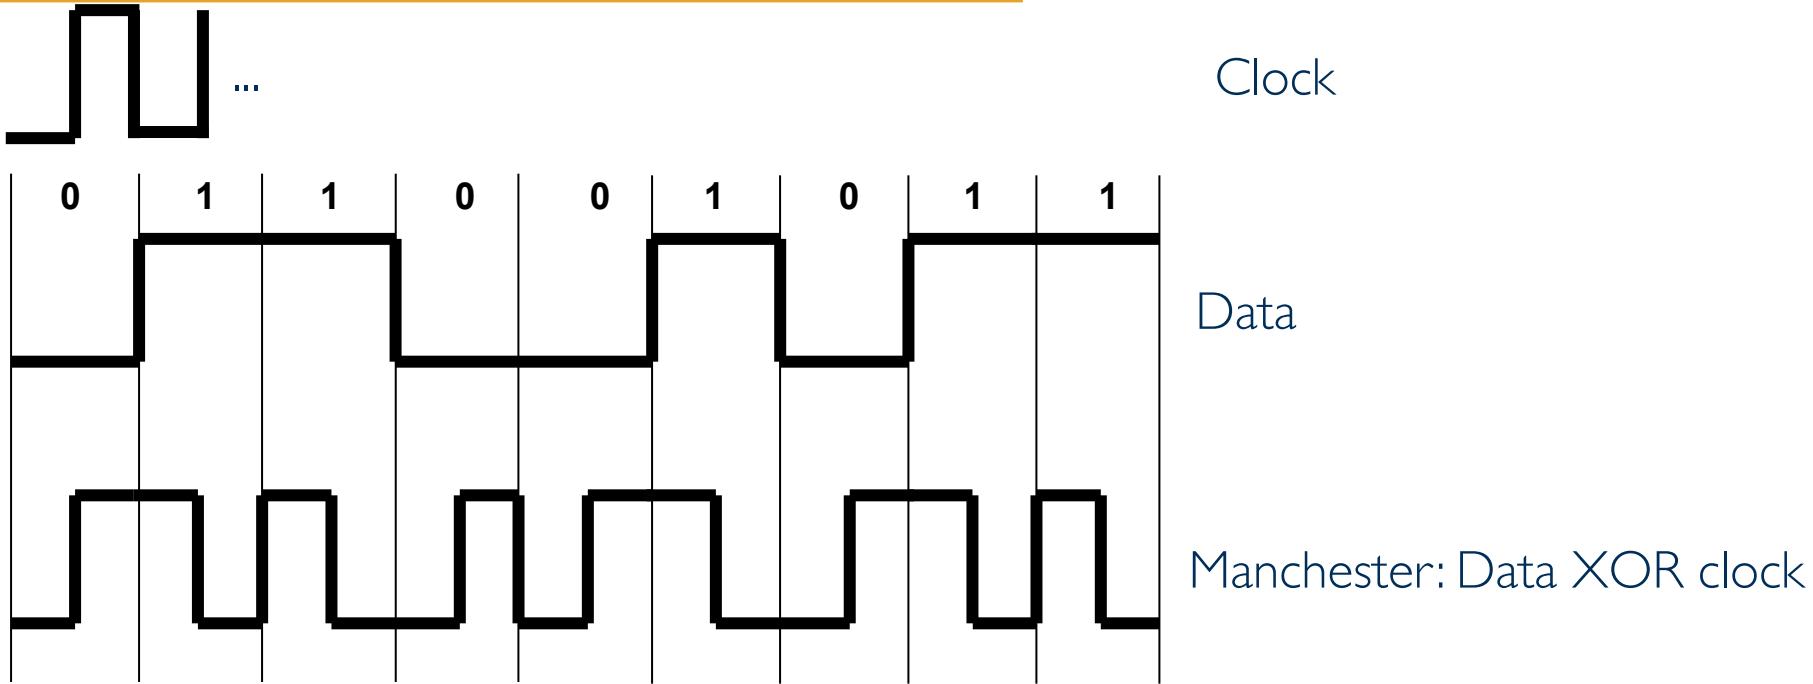
- Twisted pair copper (1m – 100m reach)
  - 10base-T
  - 100base-T
  - 1000base-TX
  - 10Gbase-T
- Fiber optic (100m – 40km reach)
  - 1000base-SX, 1000base-LX
  - 10Gbase-SR, 10Gbase-LR
  - 40Gbase-SR4, 40Gbase-LR4
  - 100Gbase-SR4, 100Gbase-SR10, 100Gbase-LR4
  - and many more

# Manchester encoding

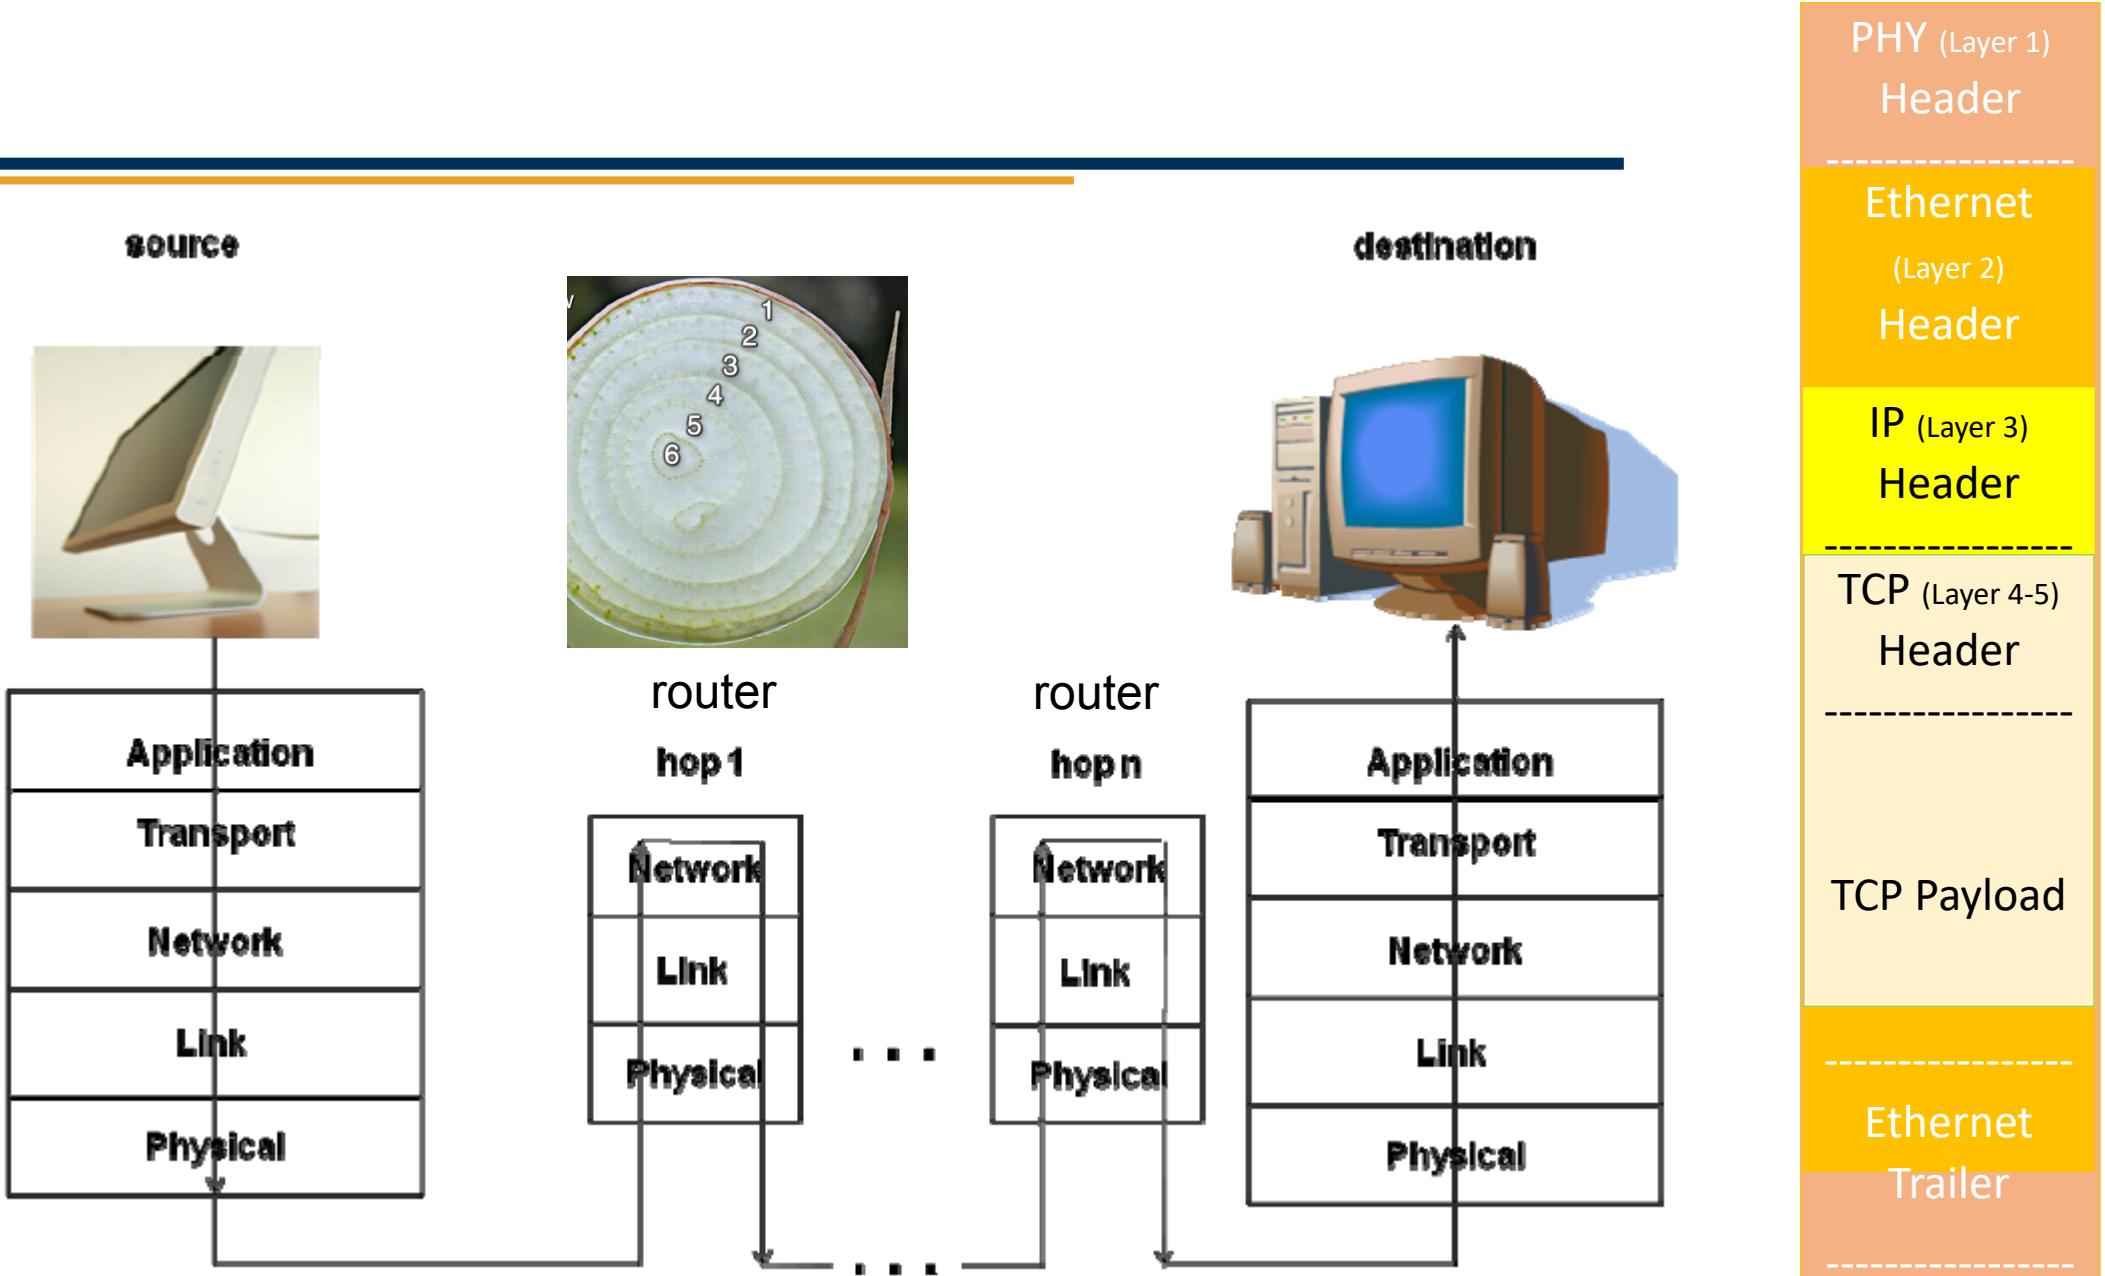


- Ethernet coax signaling uses Manchester encoding which is a scheme that guarantees at least one voltage transition during each clock tick

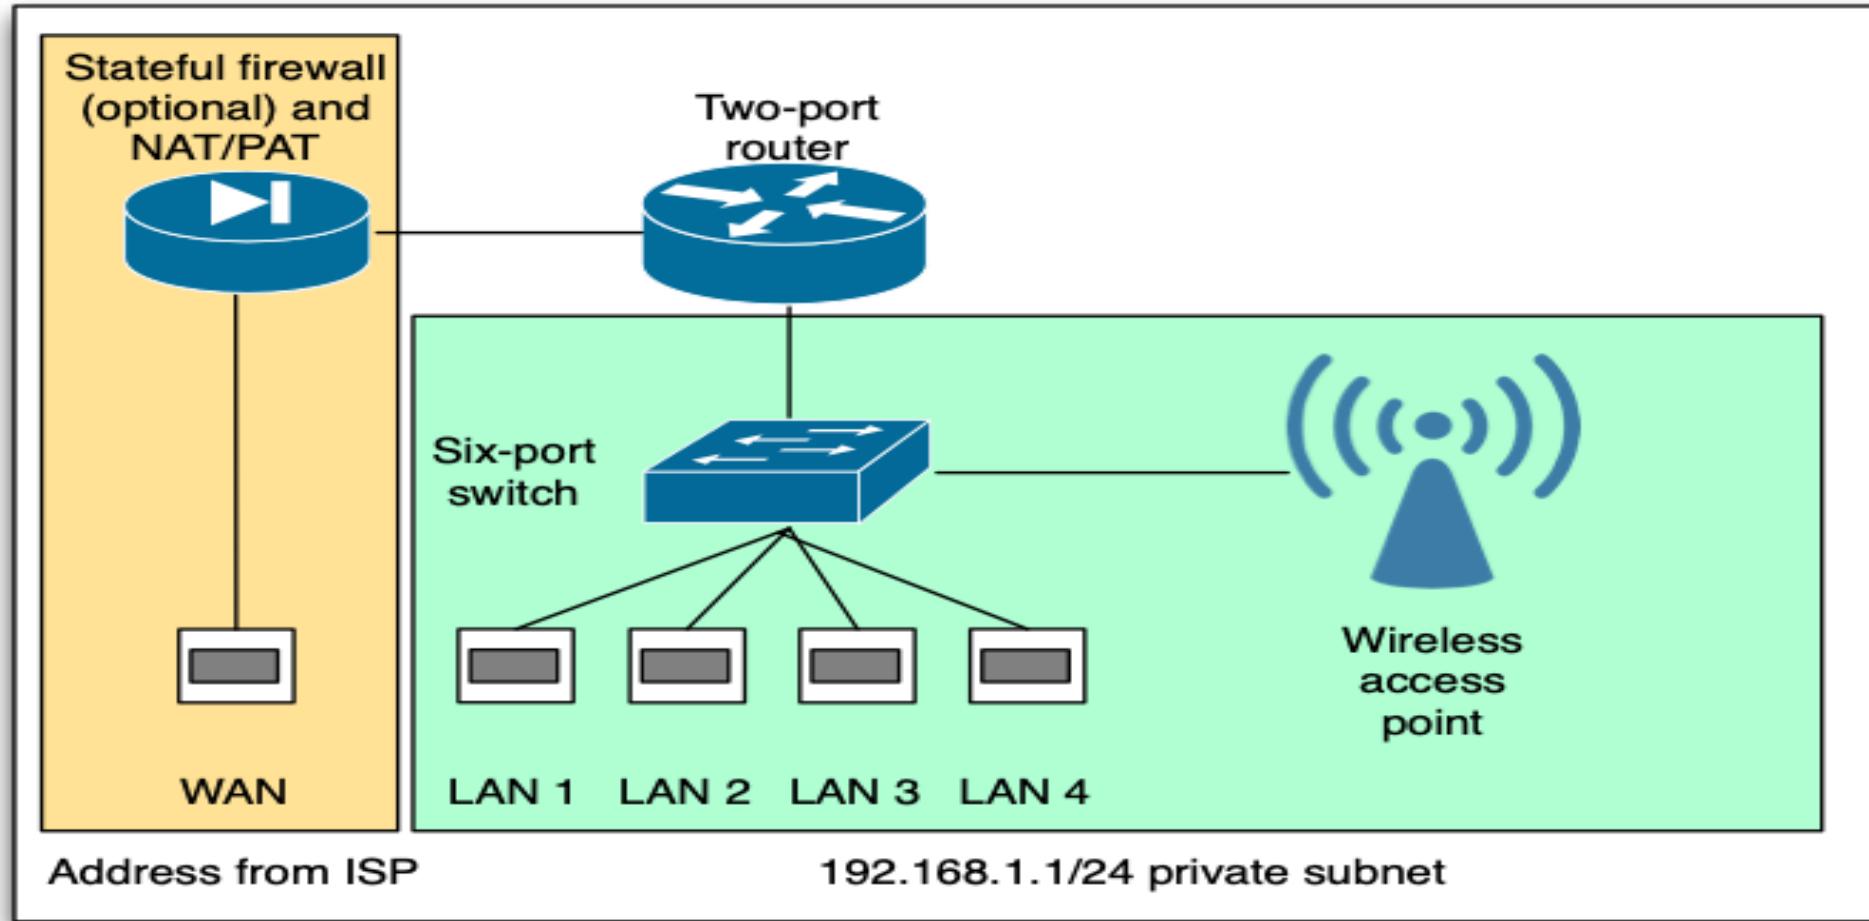
# Manchester encoding



- Ethernet coax signaling uses Manchester encoding which is a scheme that guarantees at least one voltage transition during each clock tick



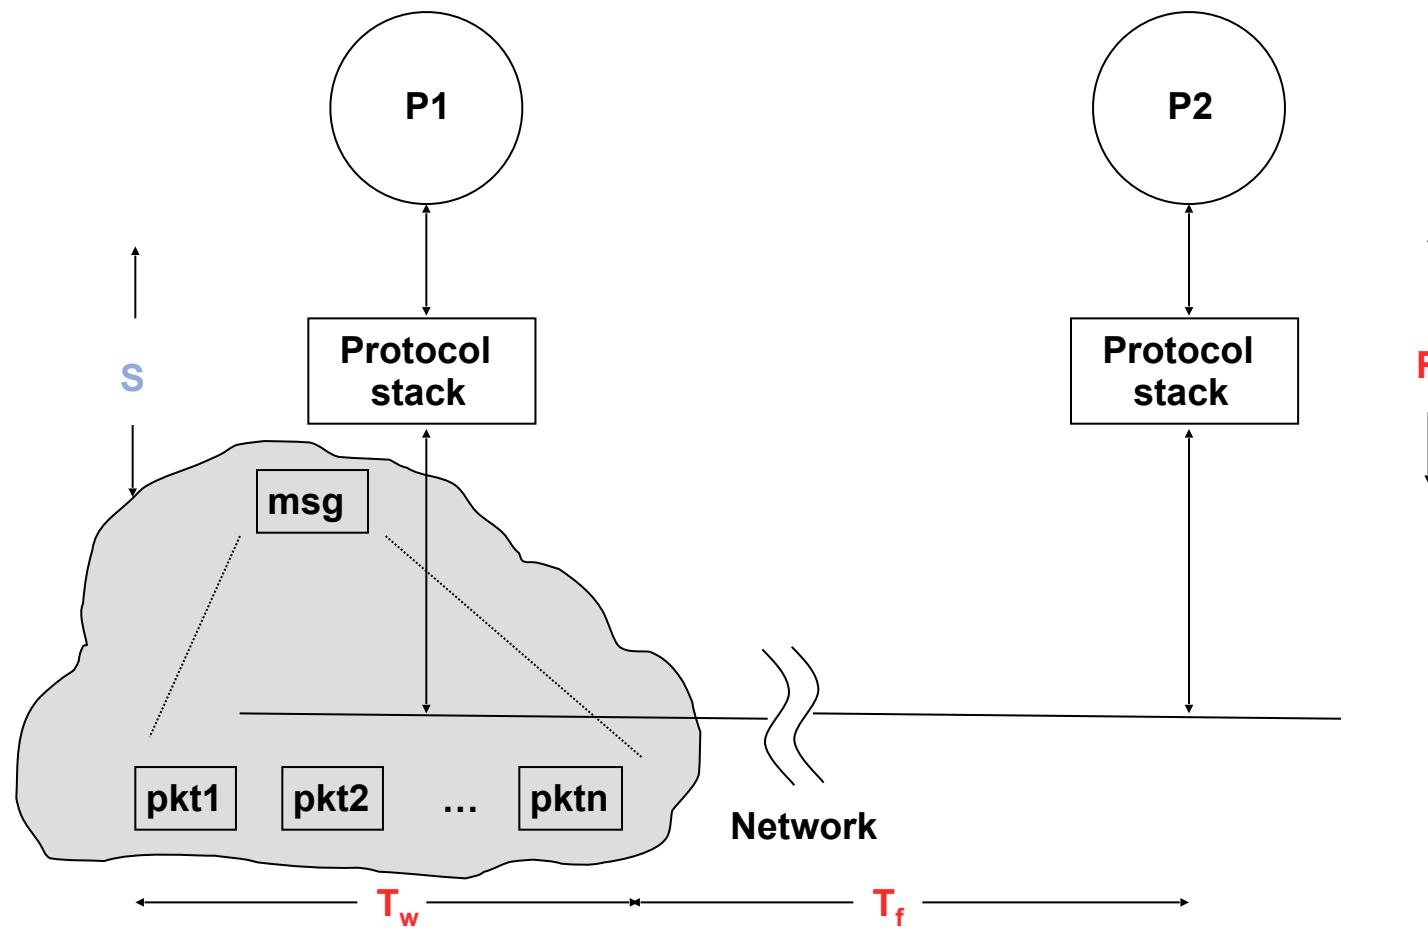
# What's in a Home Router?



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

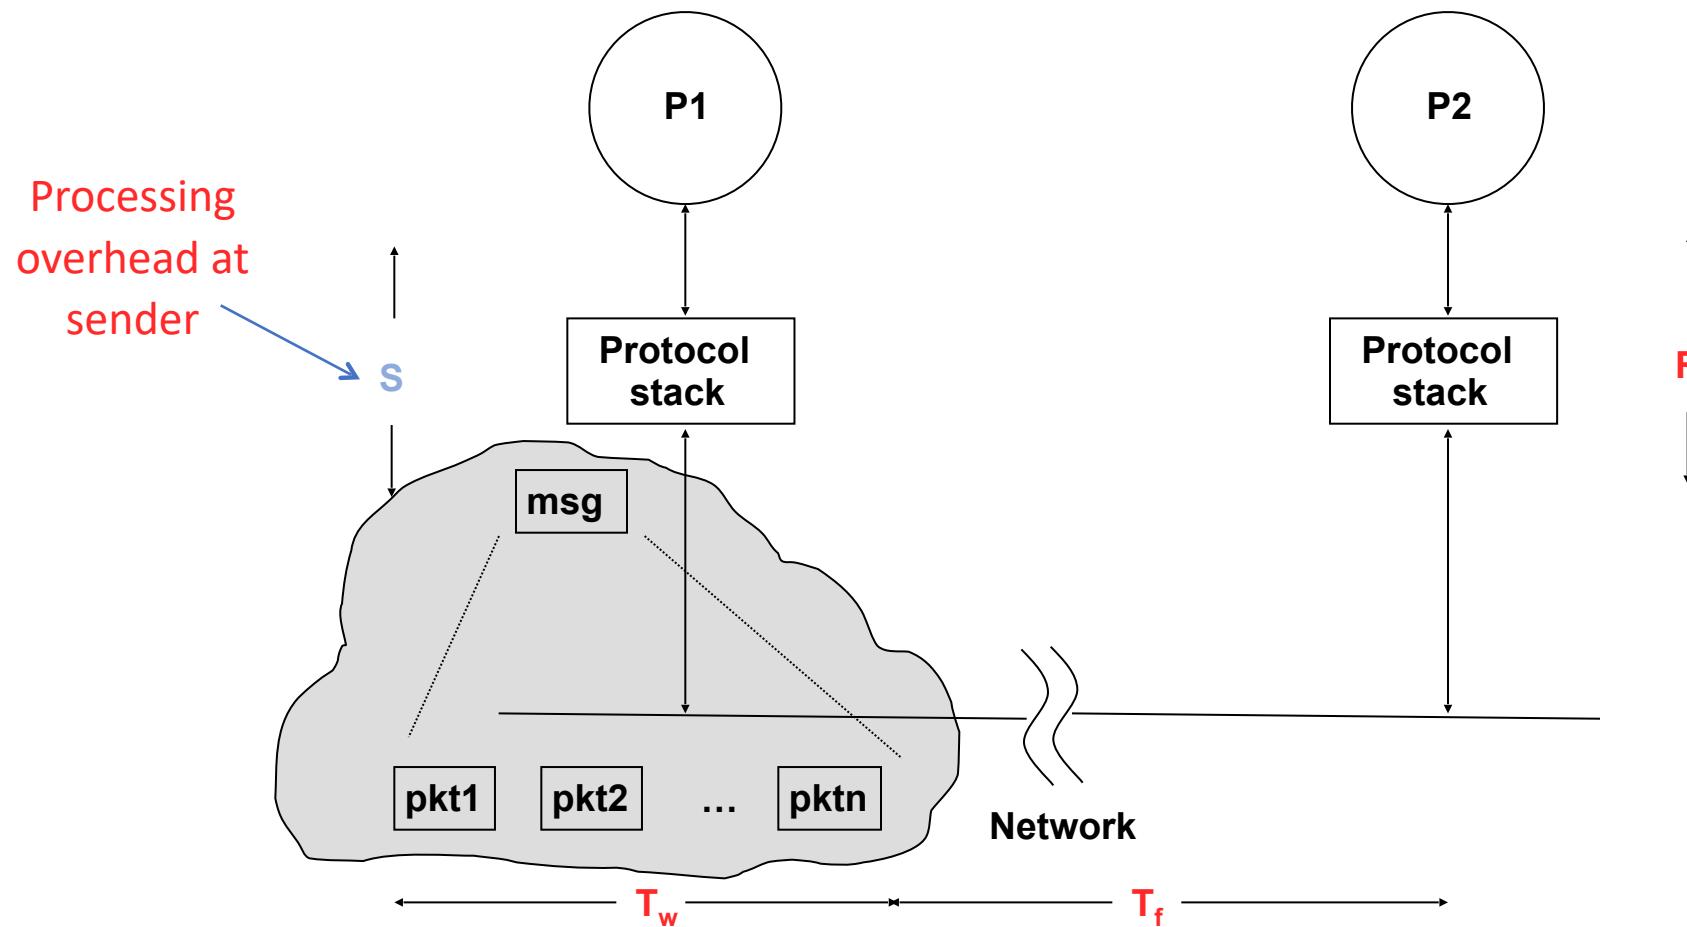
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

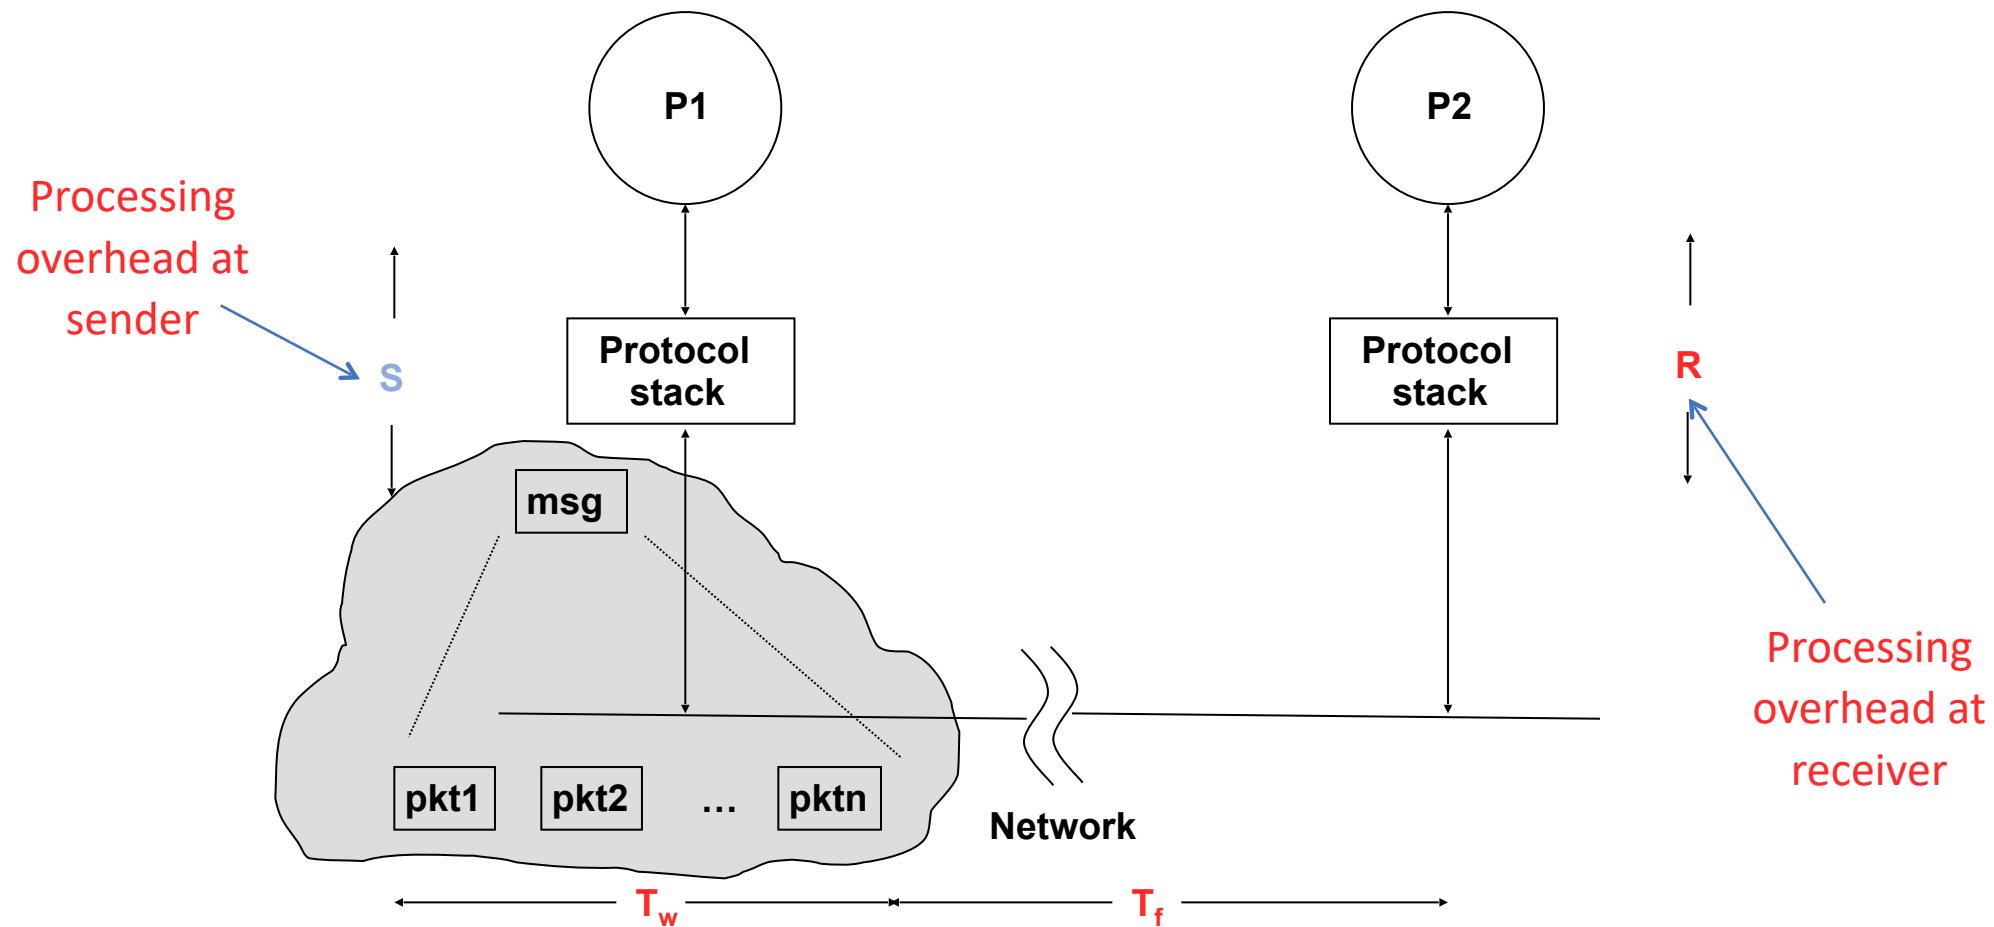
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

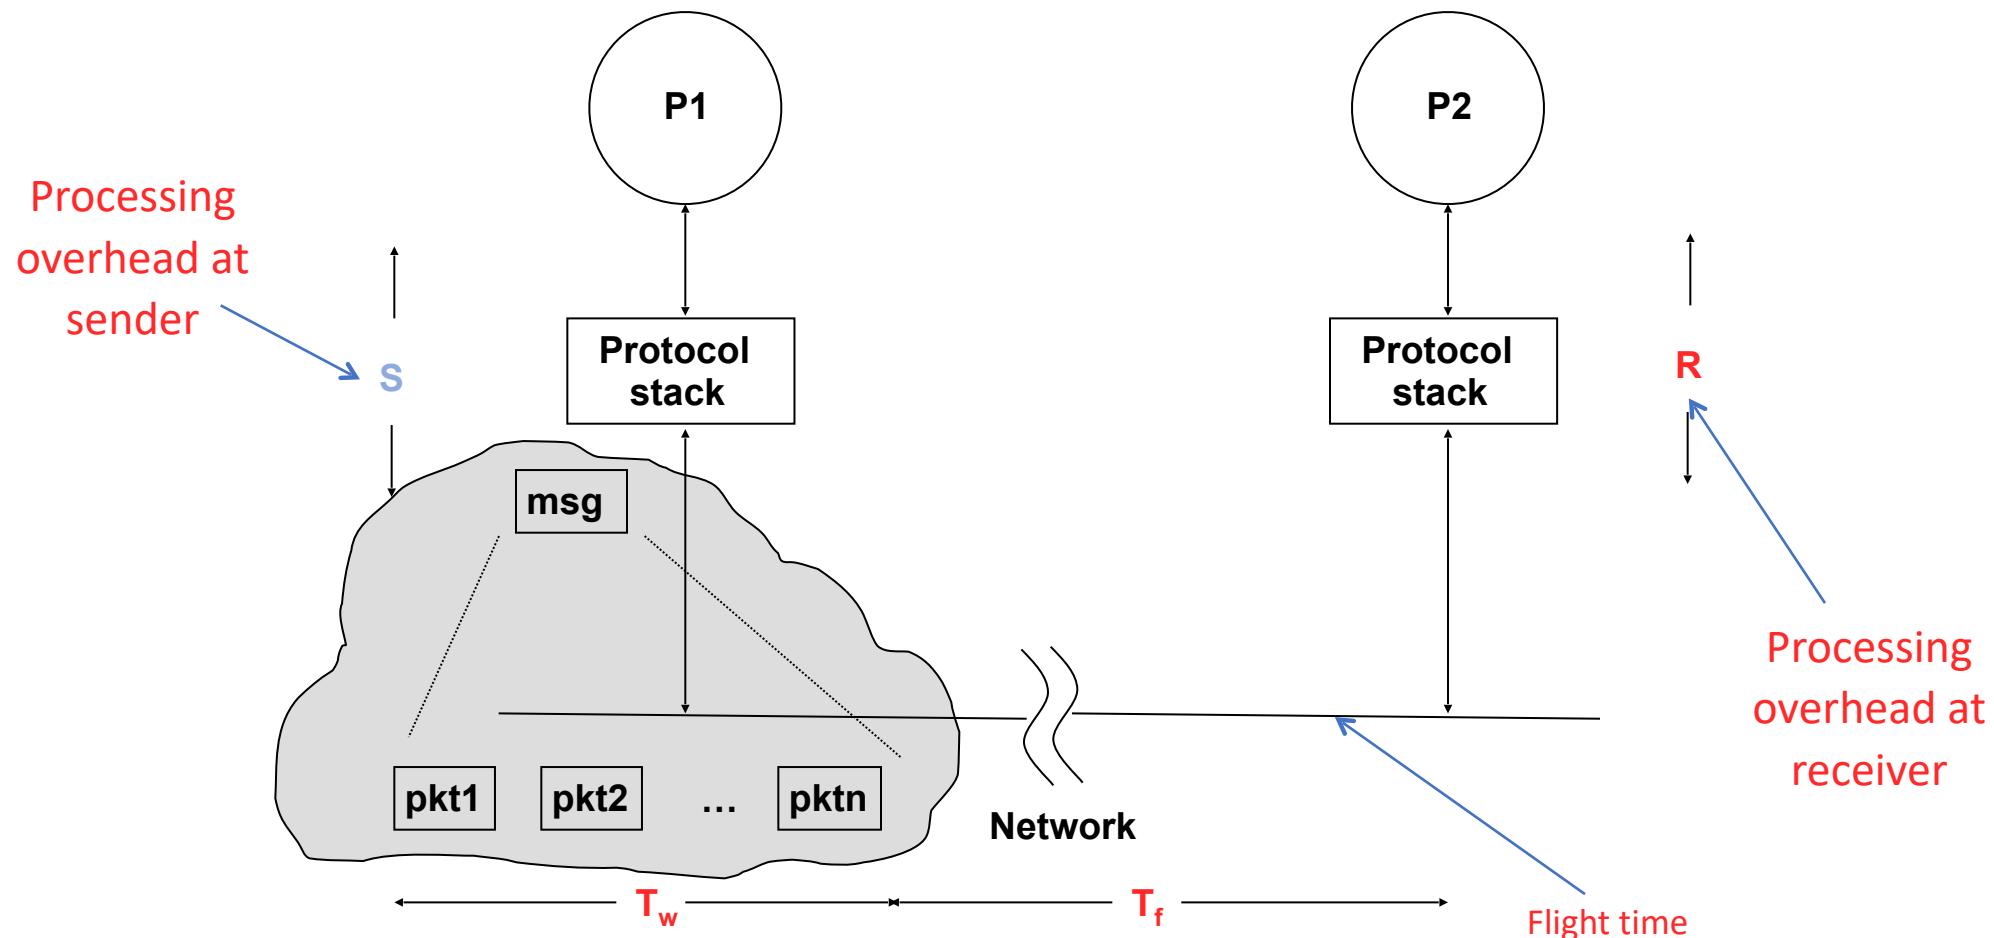
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

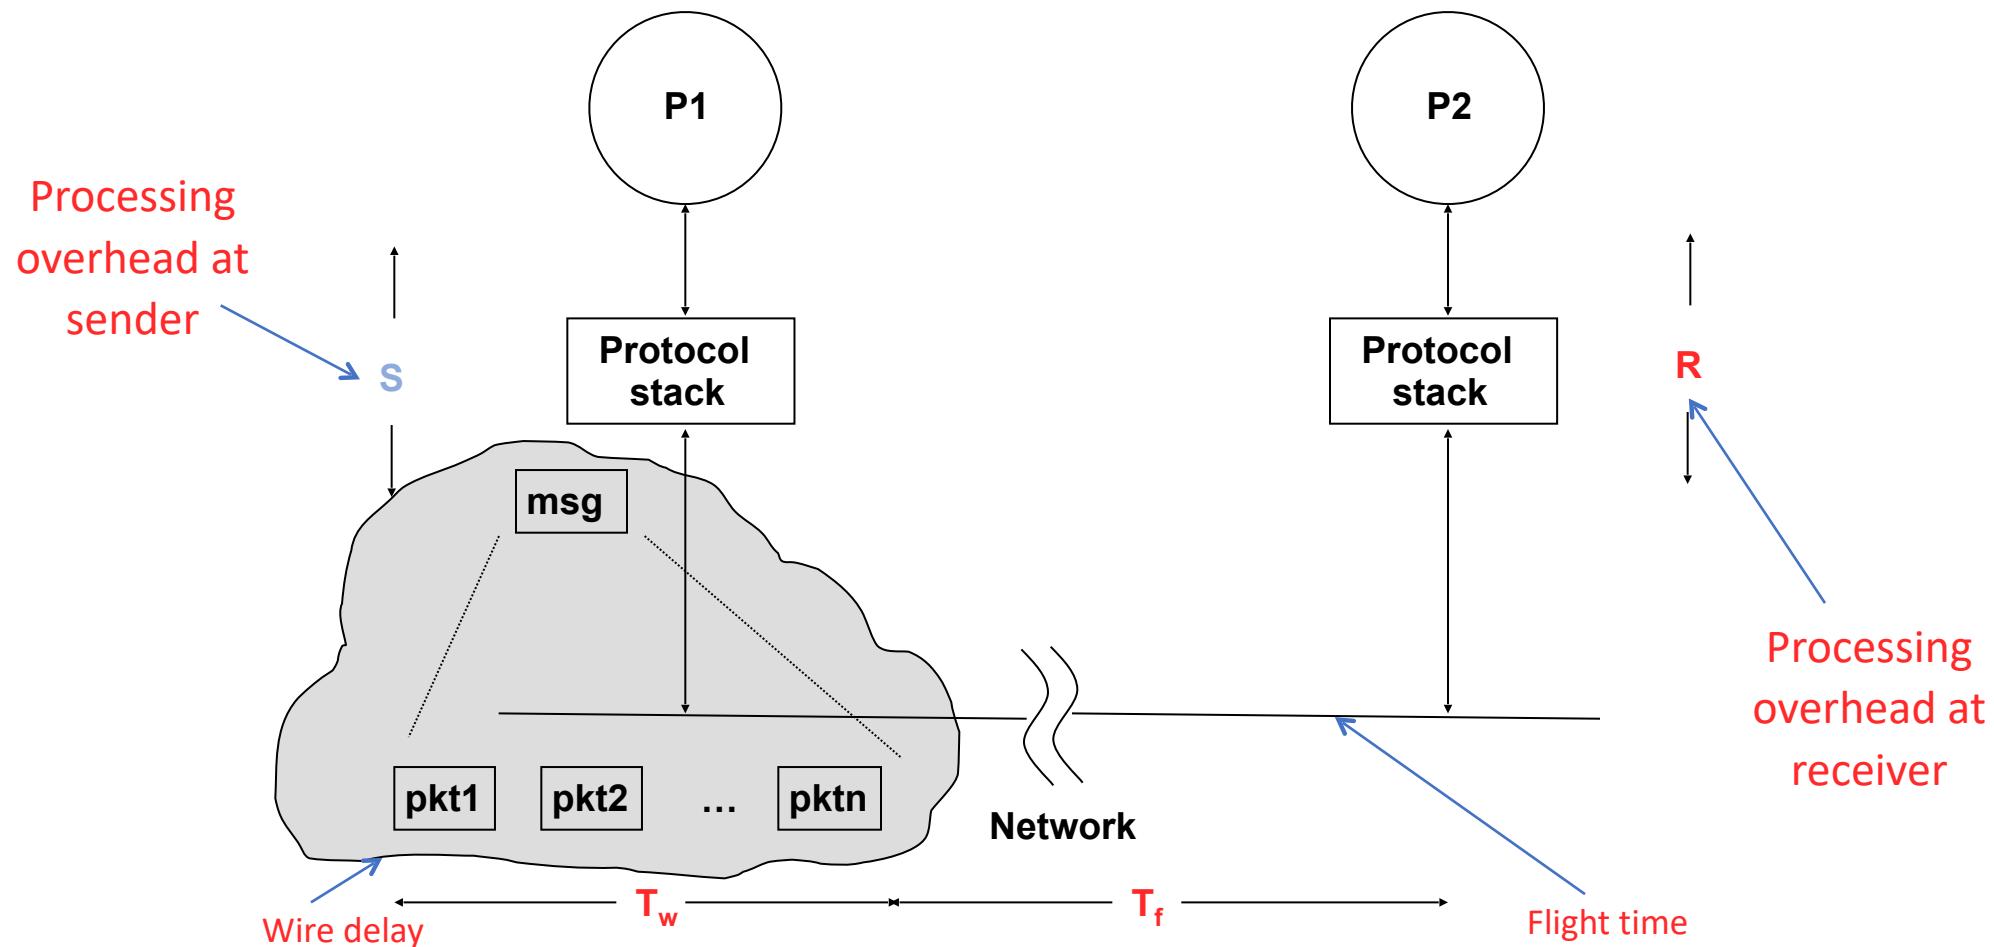
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

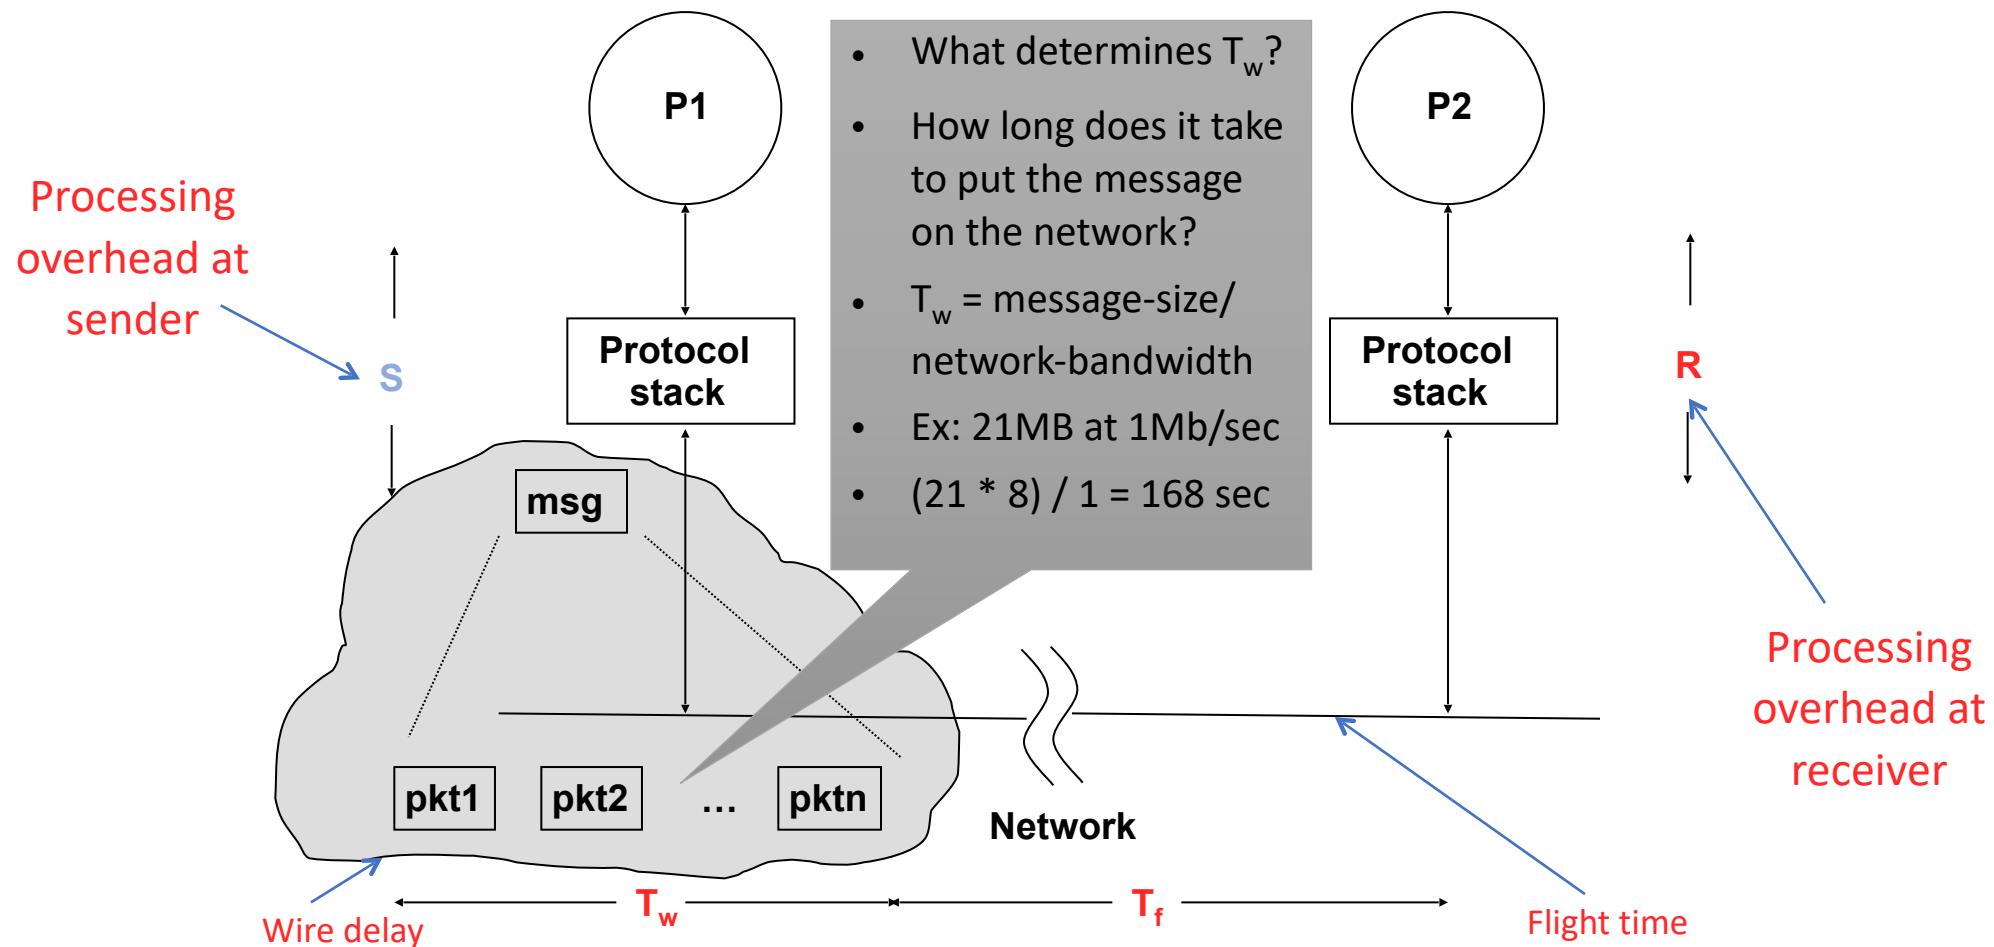
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

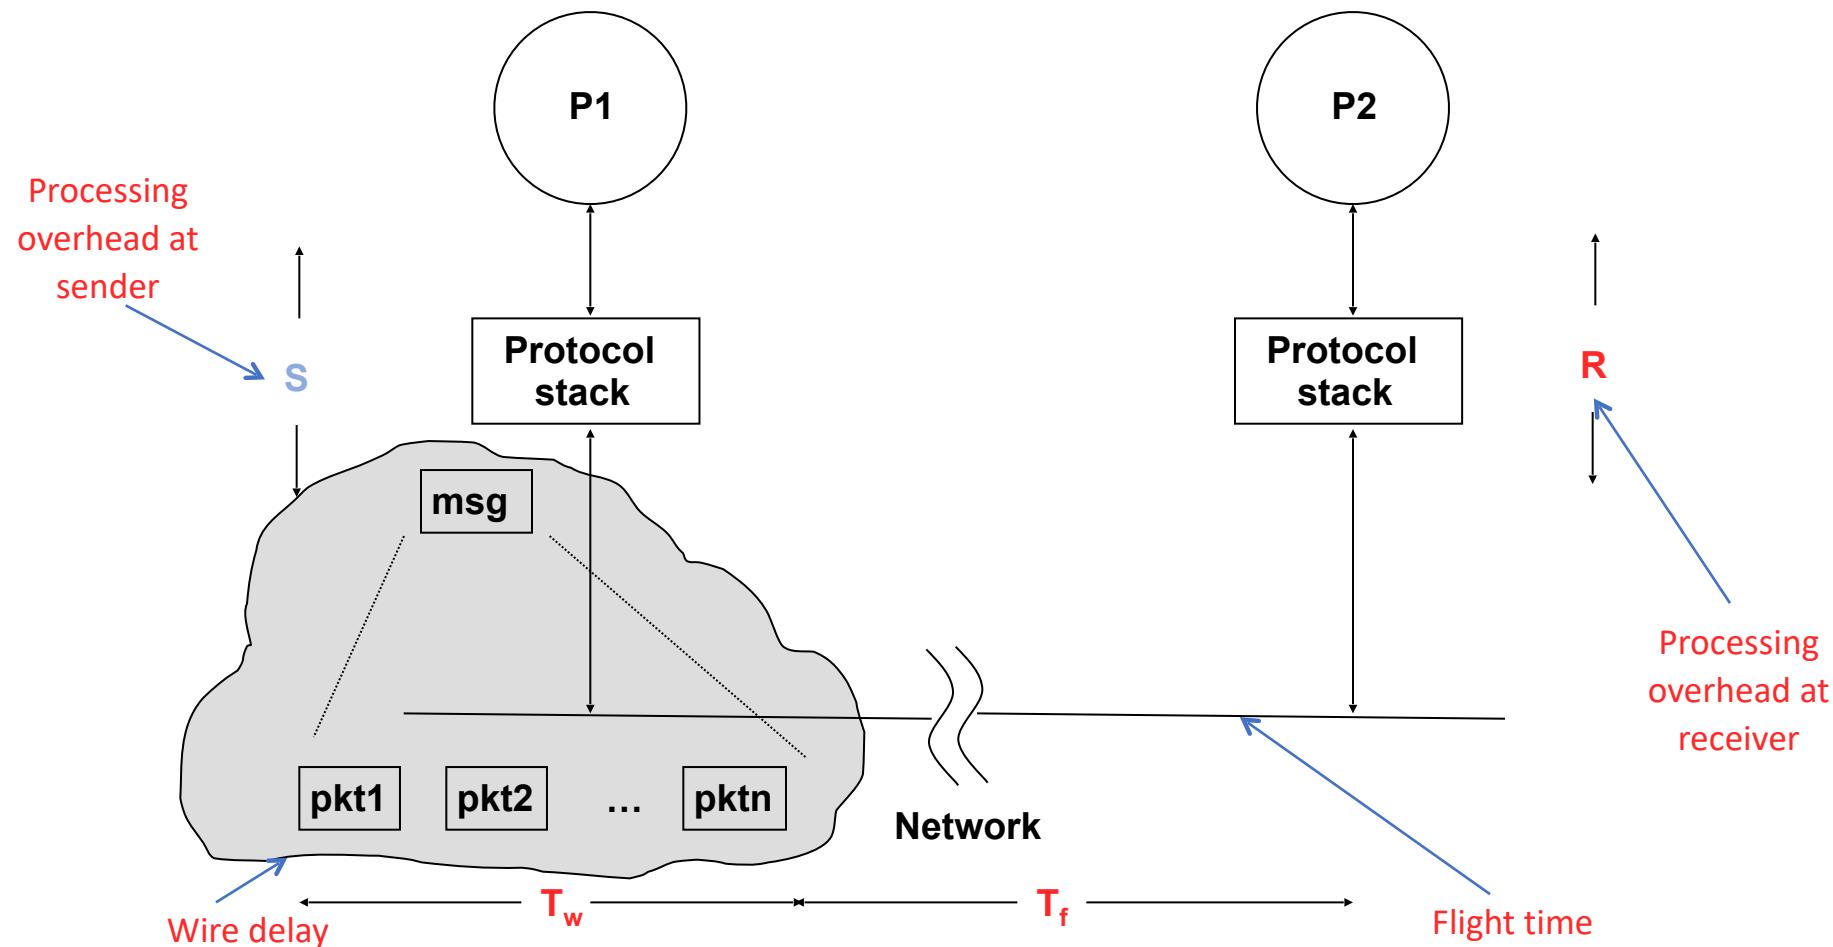
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

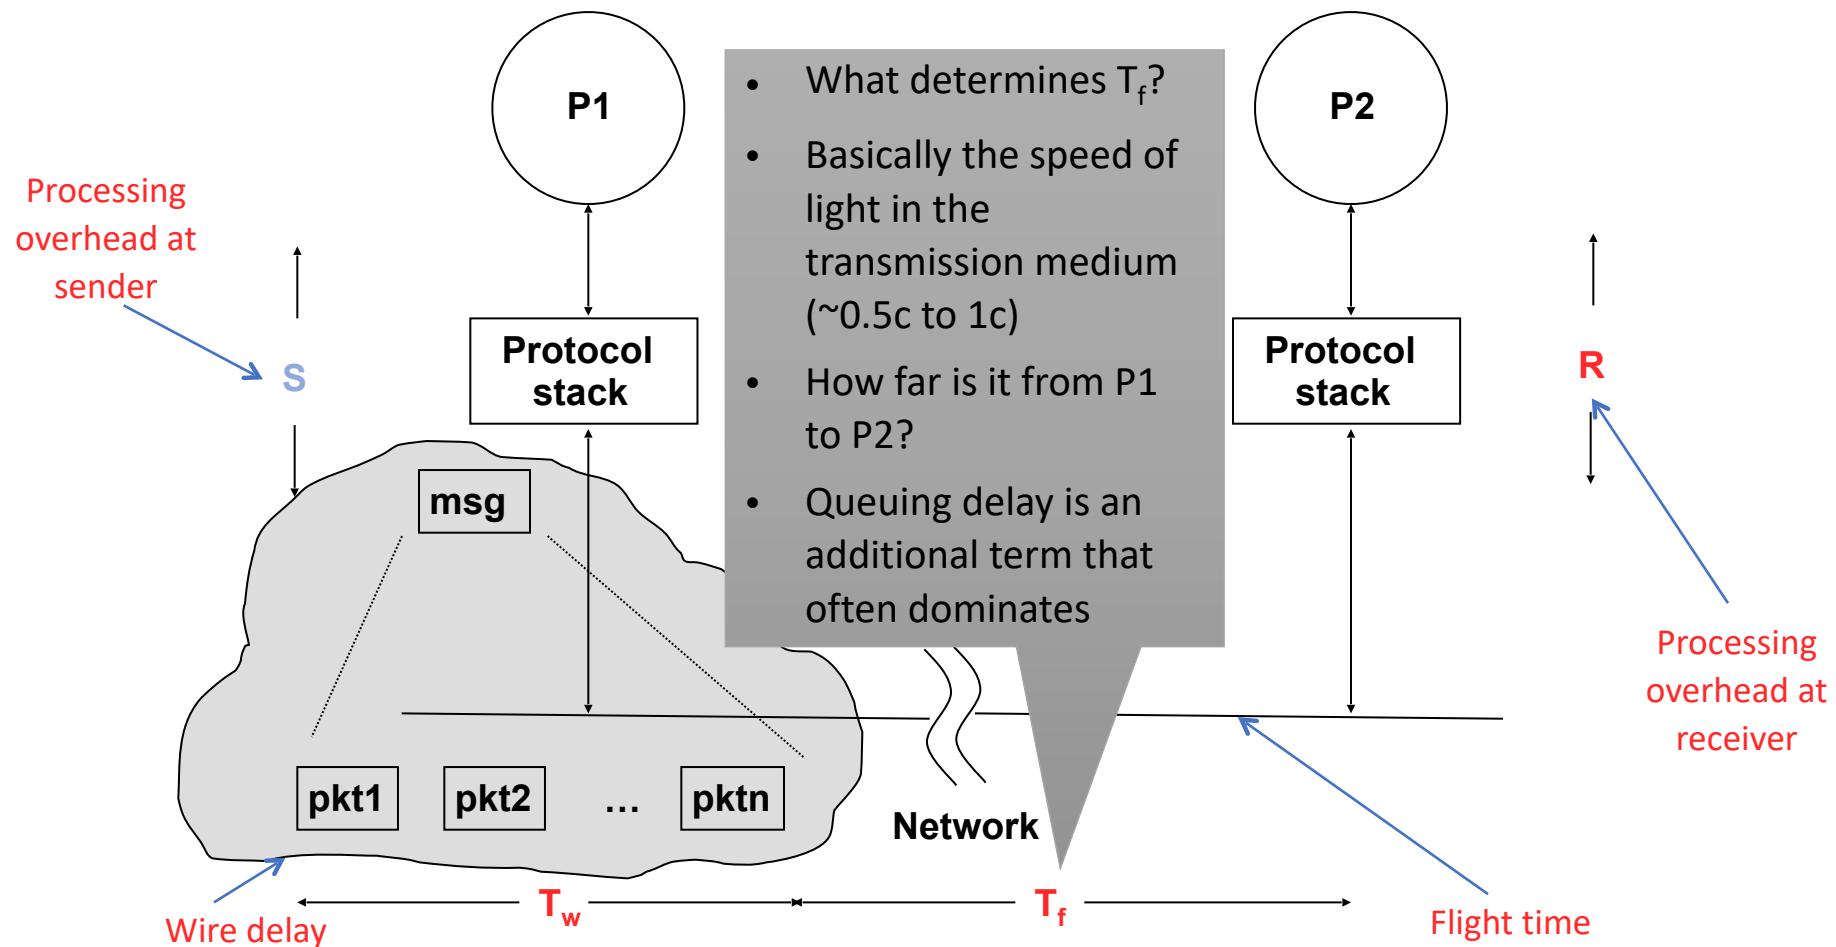
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

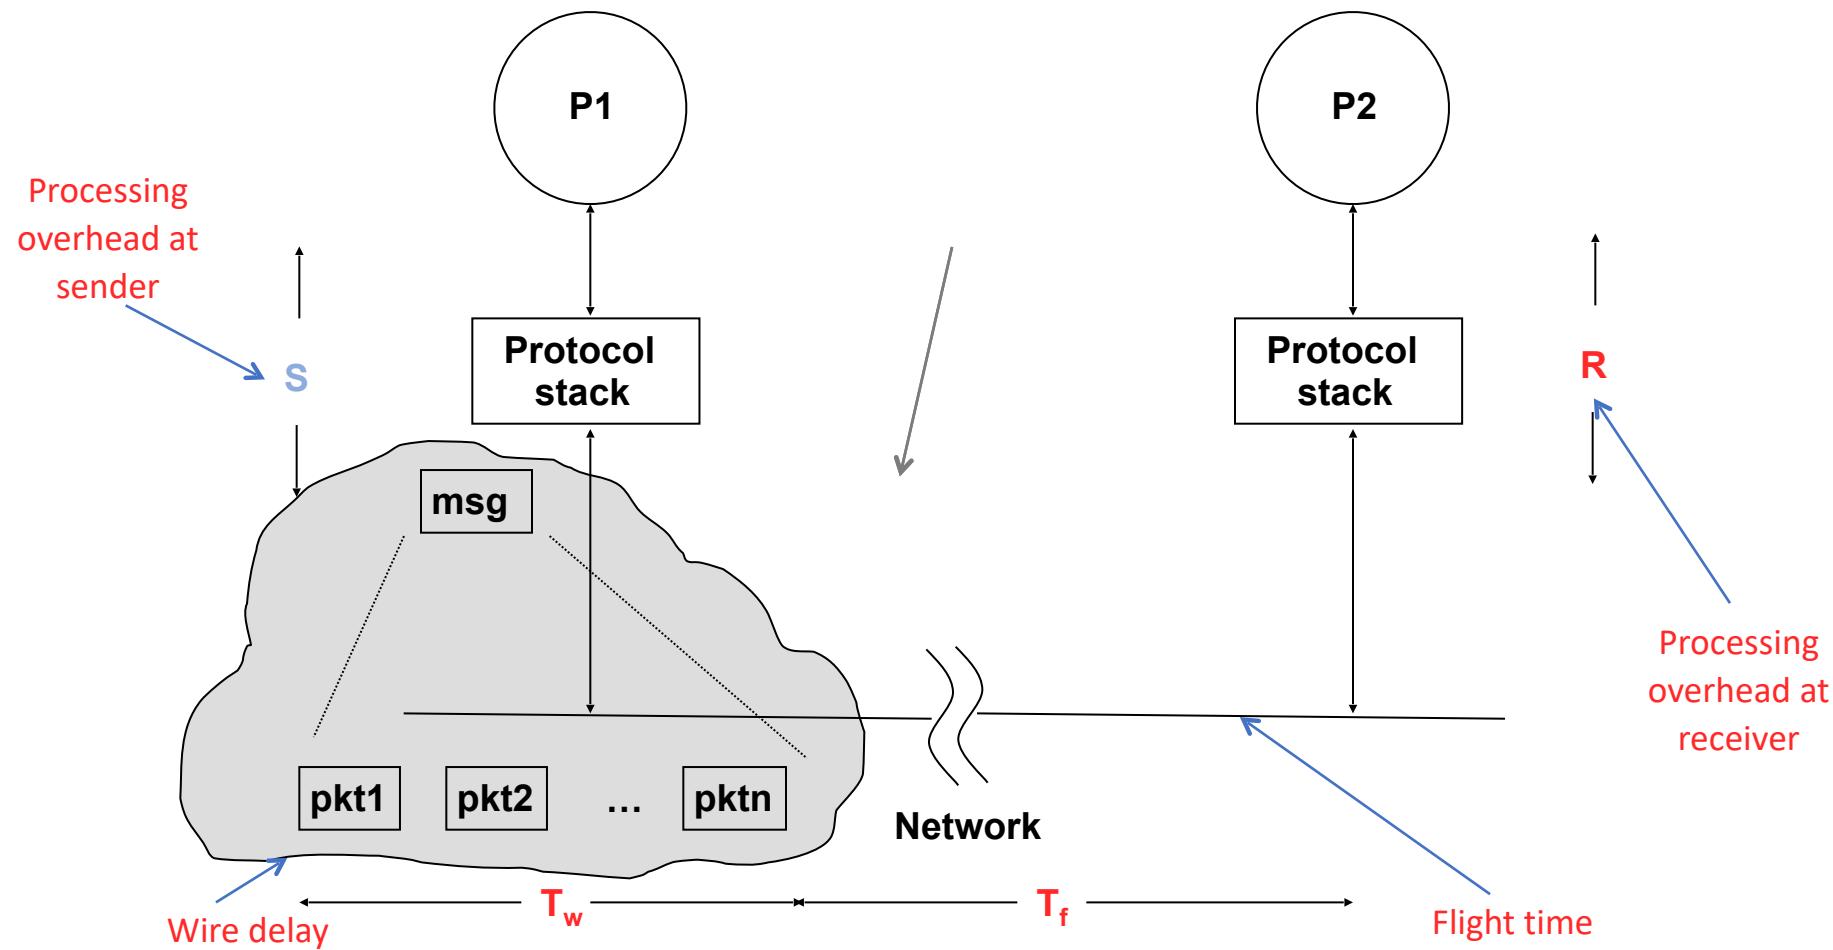
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

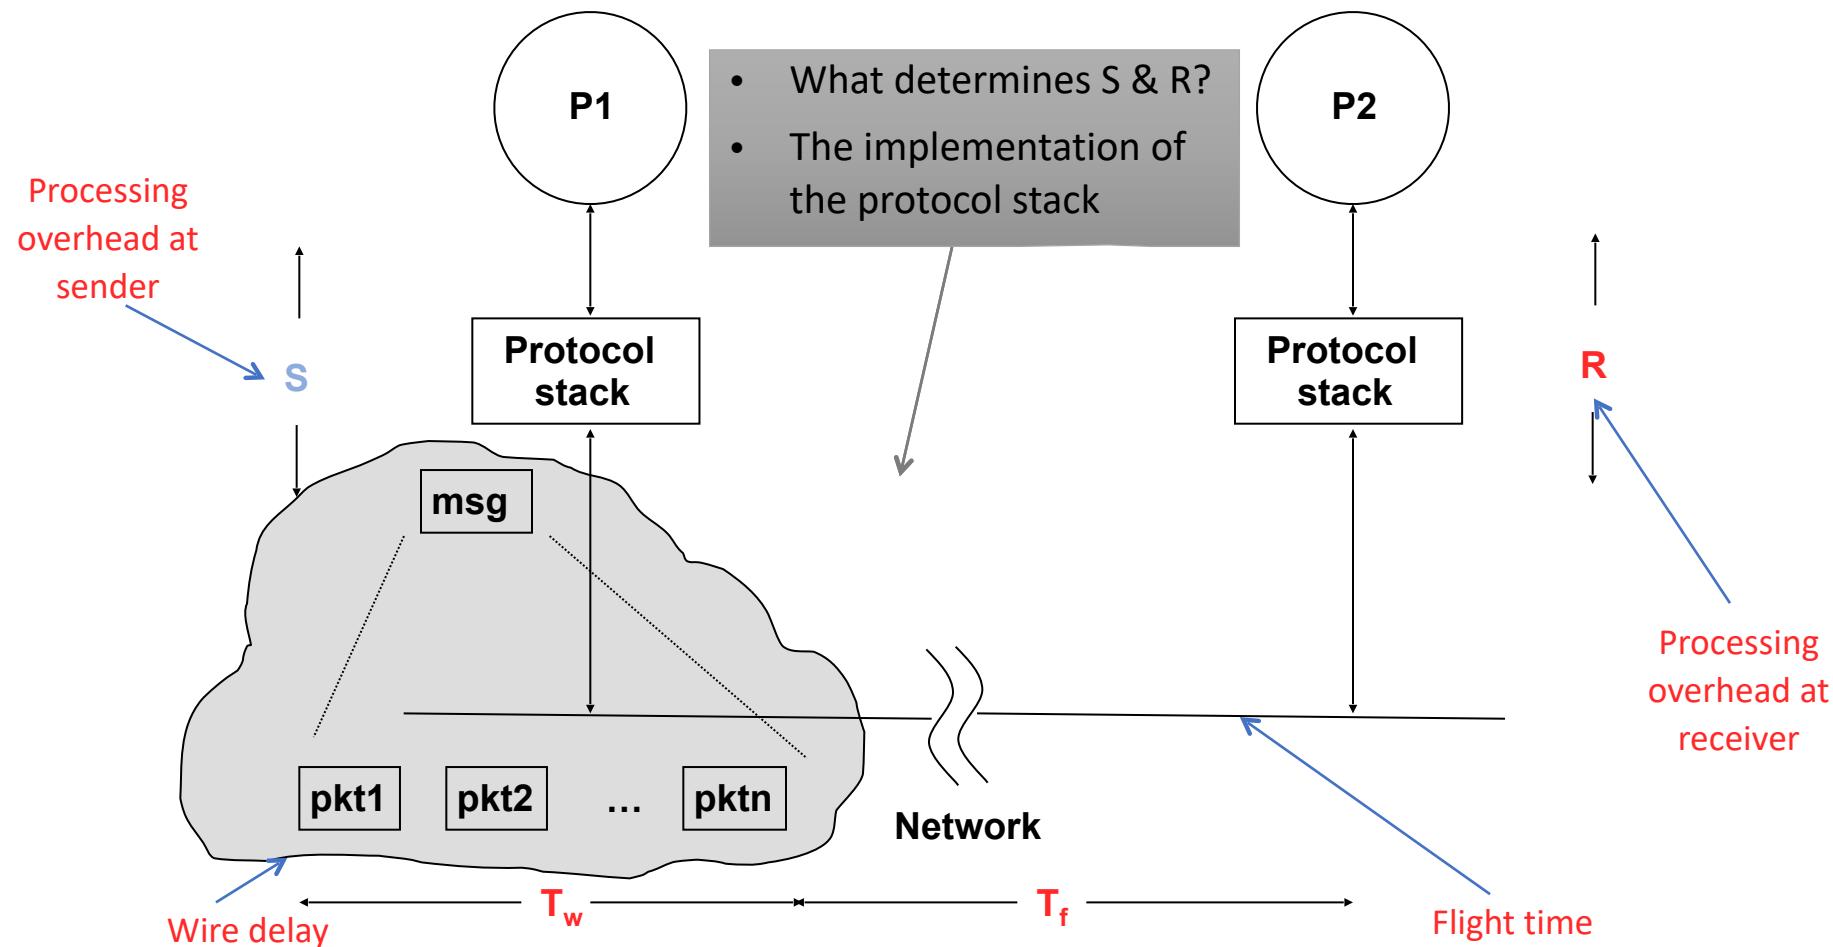
Message throughput = message-size/end-to-end-latency



# Performance metrics

Transmission time or end-to-end latency =  $S+T_w+T_f+R$

Message throughput = message-size/end-to-end-latency



# Timeline example

**100ms latency (flight time)**  
**10 packets**  
**Window size = 5**

**Wire delay (data) = 0.5ms**  
**Wire delay (ack) = ~0**  
**Receiver & Sender overhead = ~0**

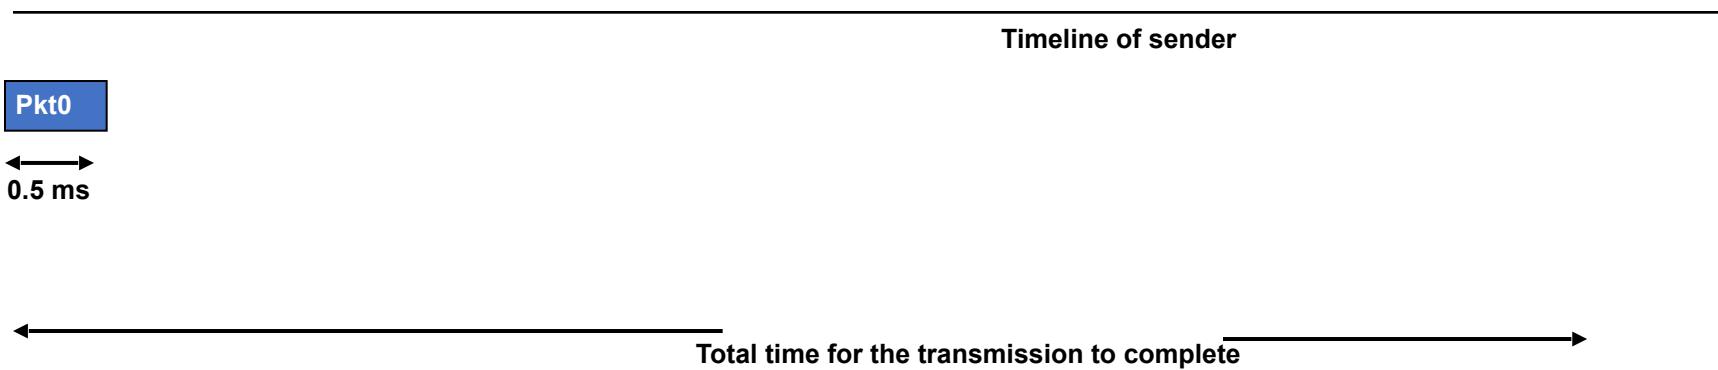
Timeline of sender

Total time for the transmission to complete

# Timeline example

**100ms latency (flight time)**  
**10 packets**  
**Window size = 5**

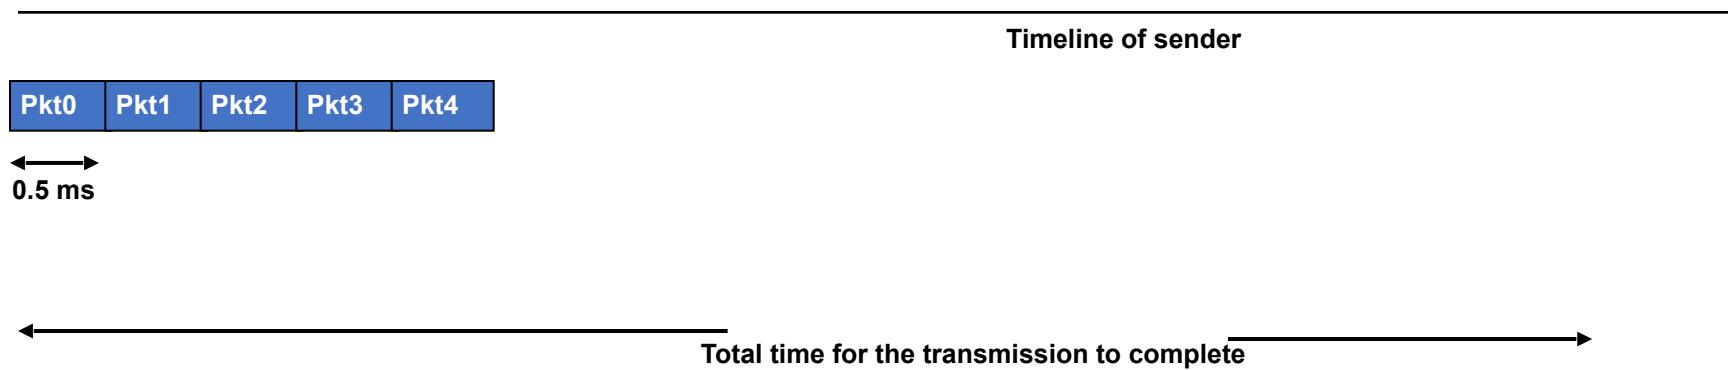
**Wire delay (data) = 0.5ms**  
**Wire delay (ack) = ~0**  
**Receiver & Sender overhead = ~0**



# Timeline example

**100ms latency (flight time)**  
**10 packets**  
**Window size = 5**

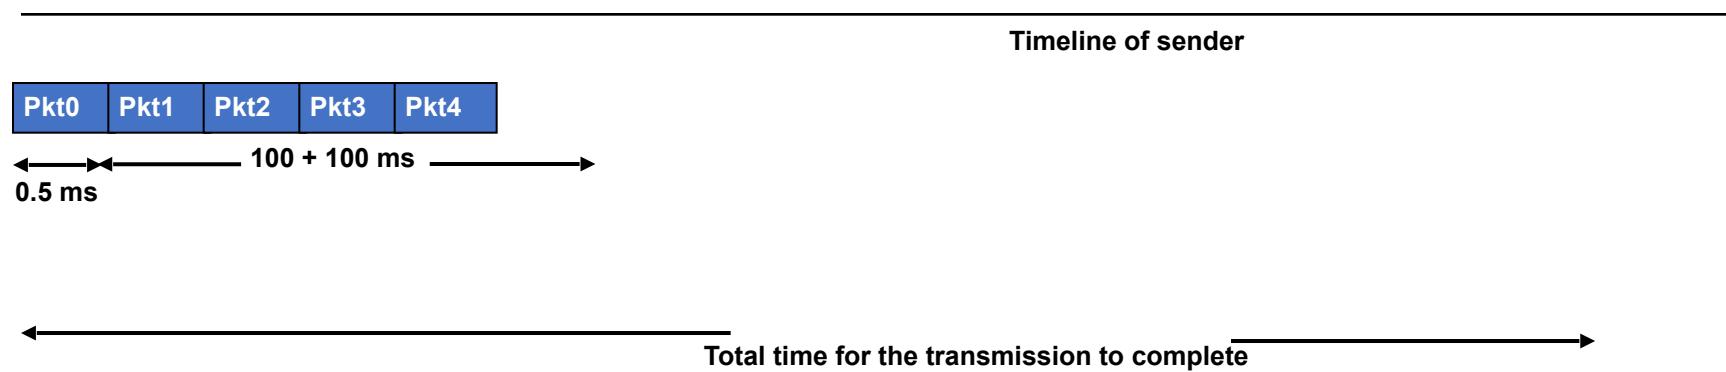
**Wire delay (data) = 0.5ms**  
**Wire delay (ack) = ~0**  
**Receiver & Sender overhead = ~0**



# Timeline example

**100ms latency (flight time)**  
**10 packets**  
**Window size = 5**

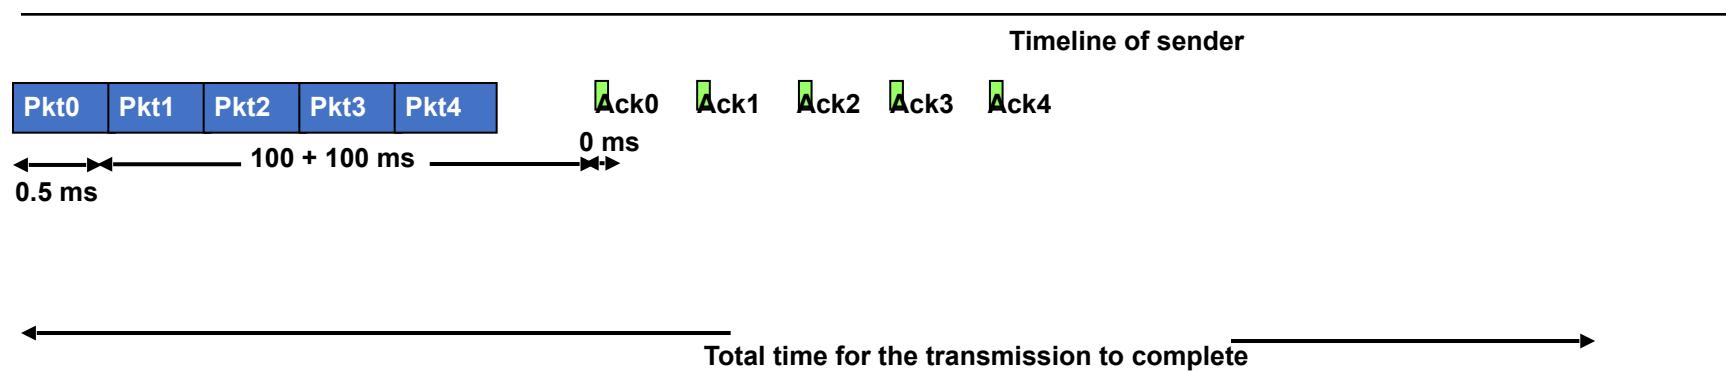
**Wire delay (data) = 0.5ms**  
**Wire delay (ack) = ~0**  
**Receiver & Sender overhead = ~0**



# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

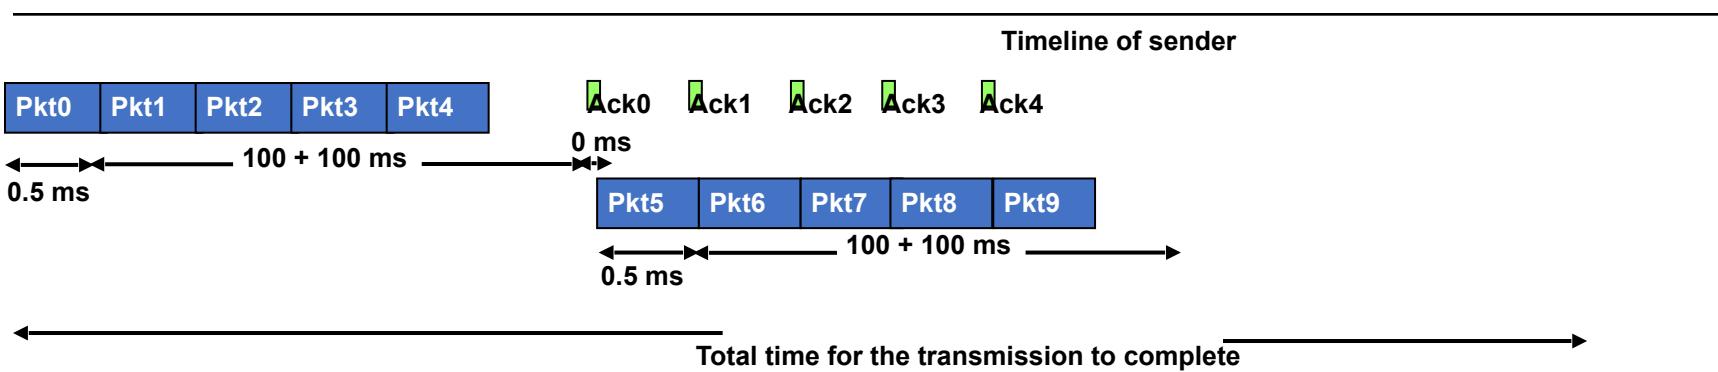
Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

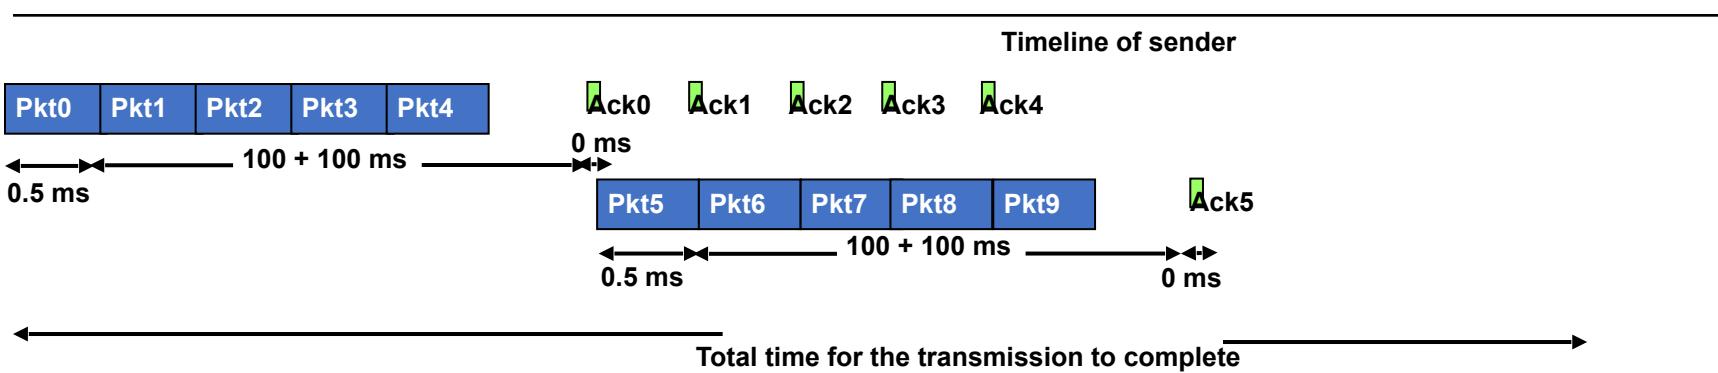
Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

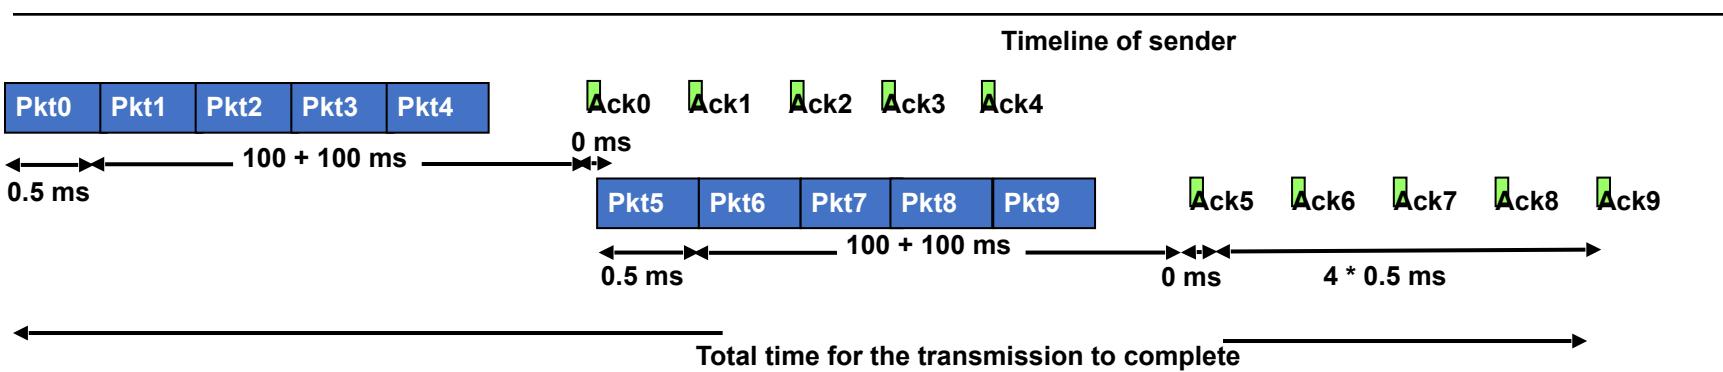
Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

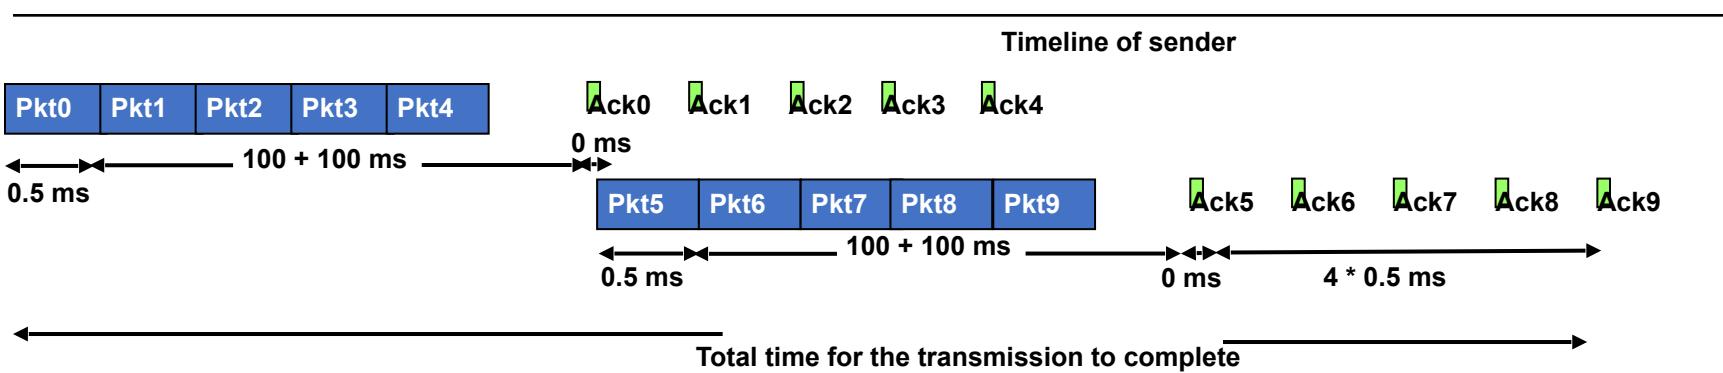
Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0

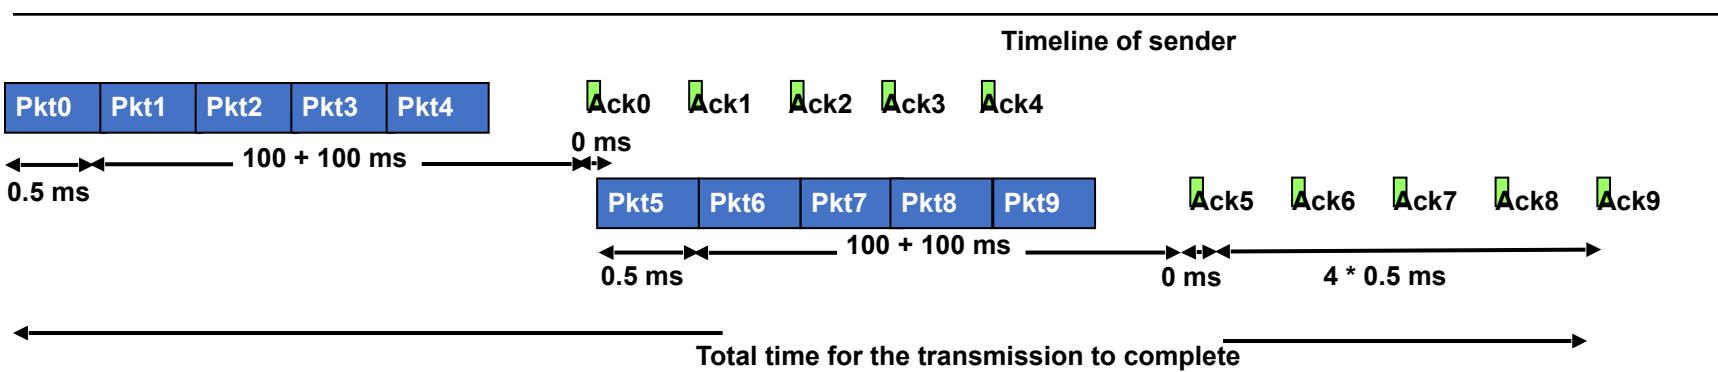


End-to-end latency for Pkt0  
 $= S + T_w + T_f + R$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

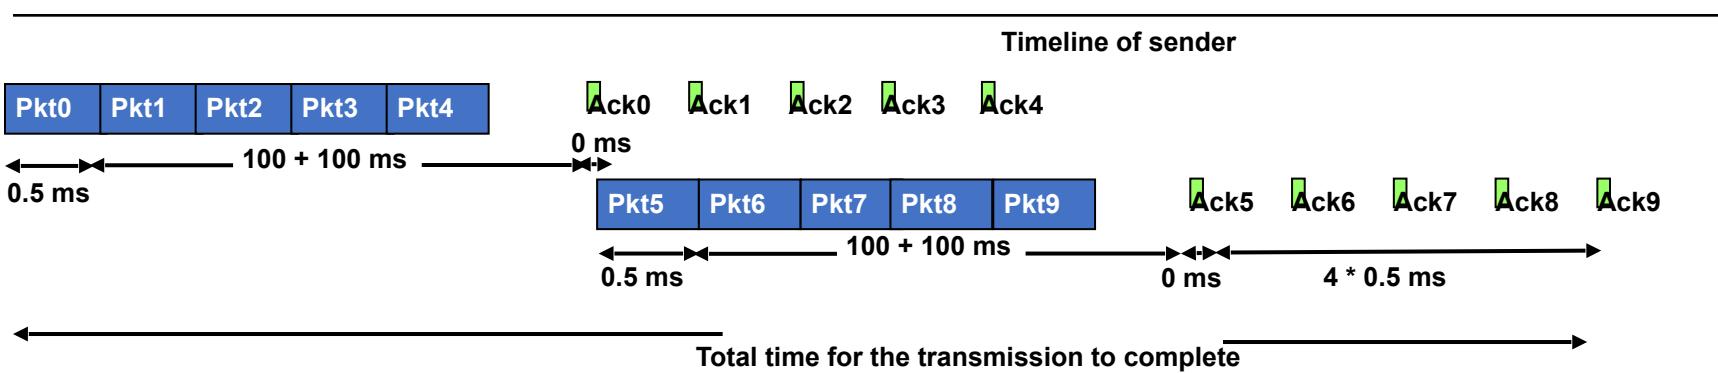
$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0

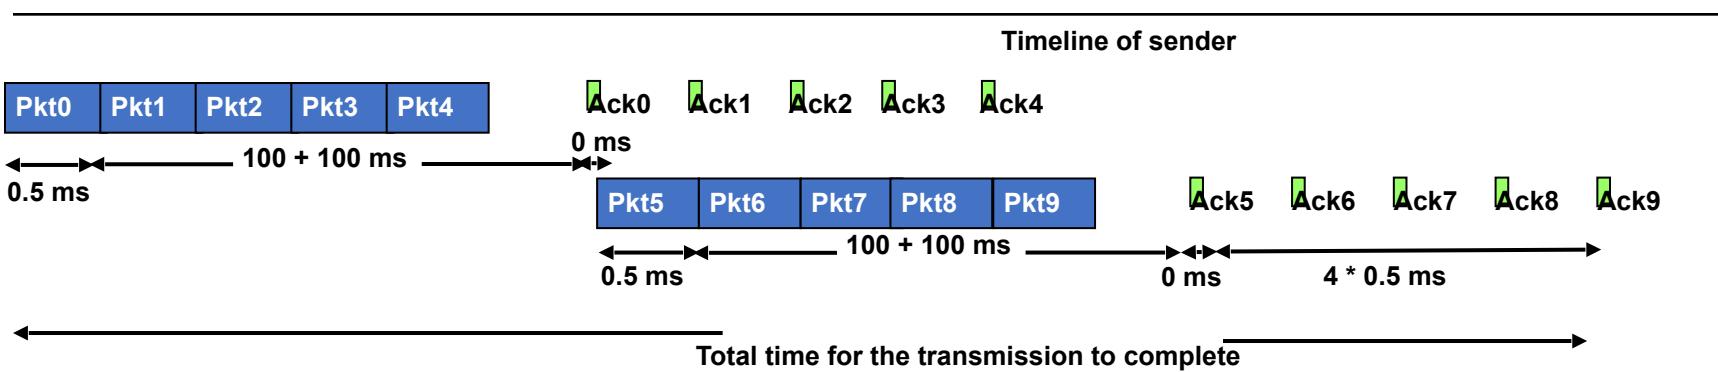


End-to-end latency for Pkt0  
 $= S + T_w + T_f + R$   
 $0 + 0.5 + 100 + 0$   
**100.5 ms**

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

$$100.5 \text{ ms}$$

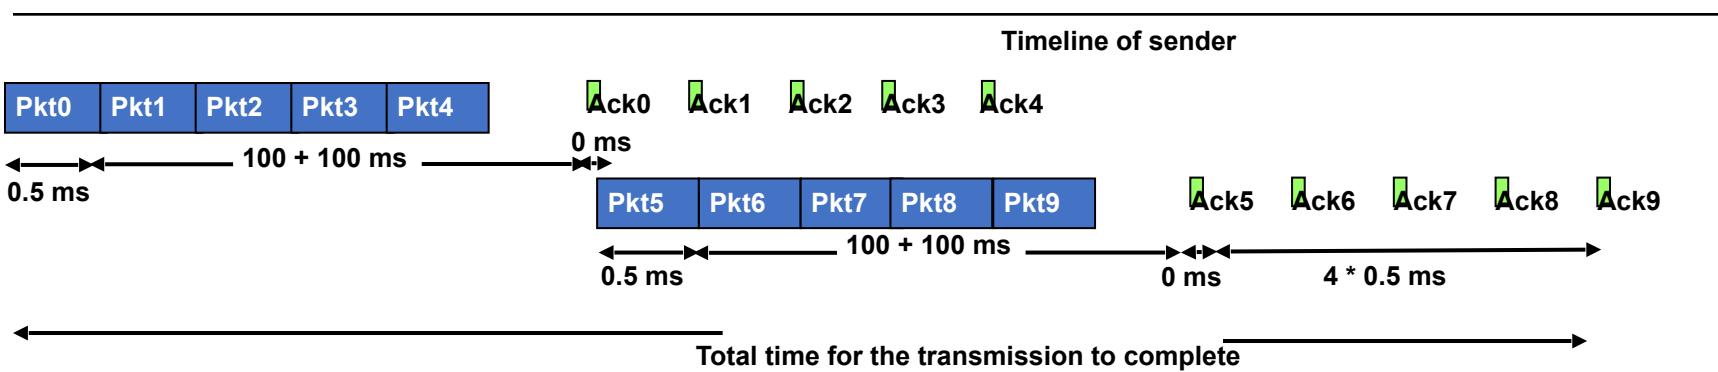
End-to-end latency for Ack0

$$= S + T_w + T_f + R$$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

$$100.5 \text{ ms}$$

End-to-end latency for Ack0

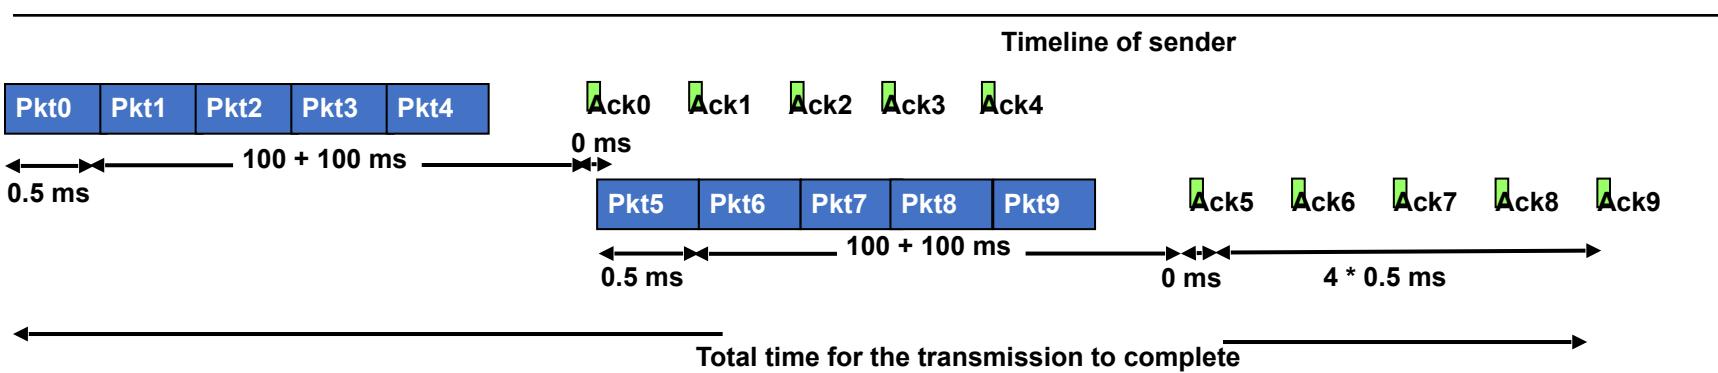
$$= S + T_w + T_f + R$$

$$0 + 0.0 + 100 + 0$$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

$$100.5 \text{ ms}$$

End-to-end latency for Ack0

$$= S + T_w + T_f + R$$

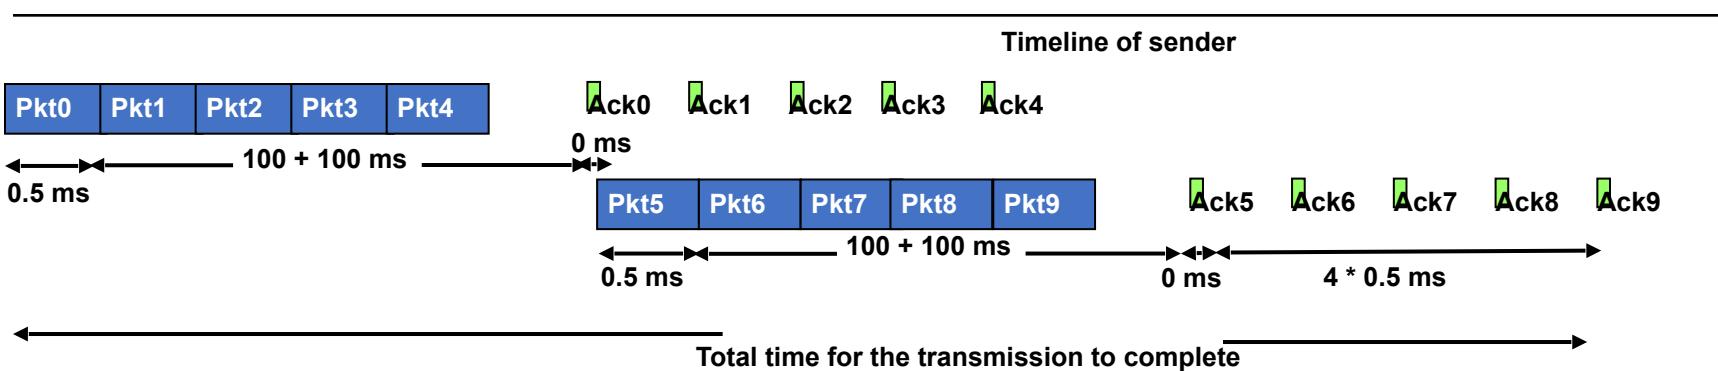
$$0 + 0.0 + 100 + 0$$

$$100.0 \text{ ms}$$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

$$100.5 \text{ ms}$$

Total transmission time

End-to-end latency for Ack0

$$= S + T_w + T_f + R$$

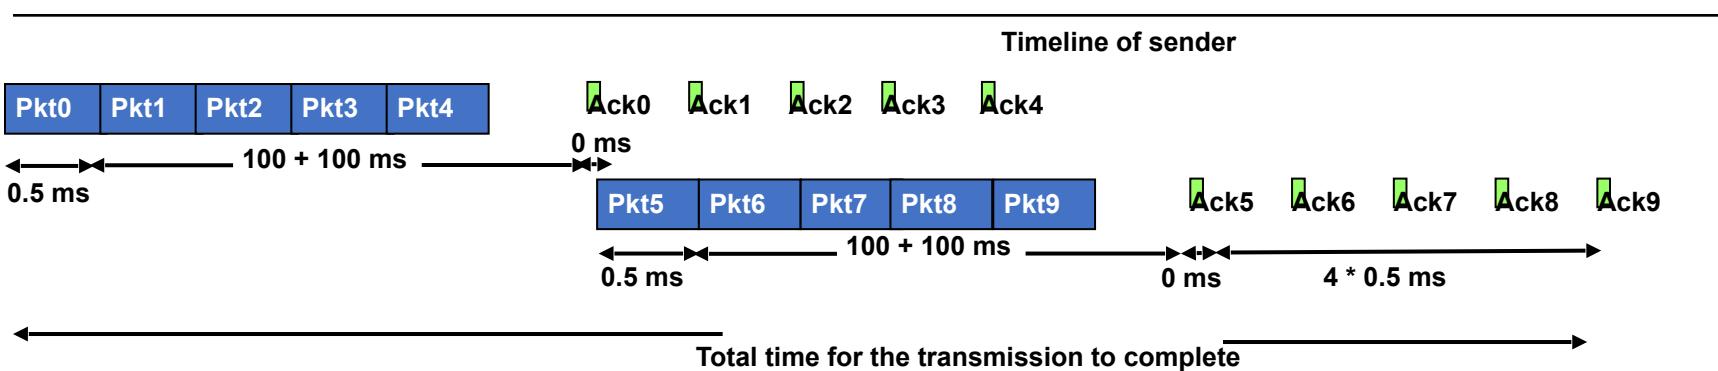
$$0 + 0.0 + 100 + 0$$

$$100.0 \text{ ms}$$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

$$100.5 \text{ ms}$$

End-to-end latency for Ack0

$$= S + T_w + T_f + R$$

$$0 + 0.0 + 100 + 0$$

$$100.0 \text{ ms}$$

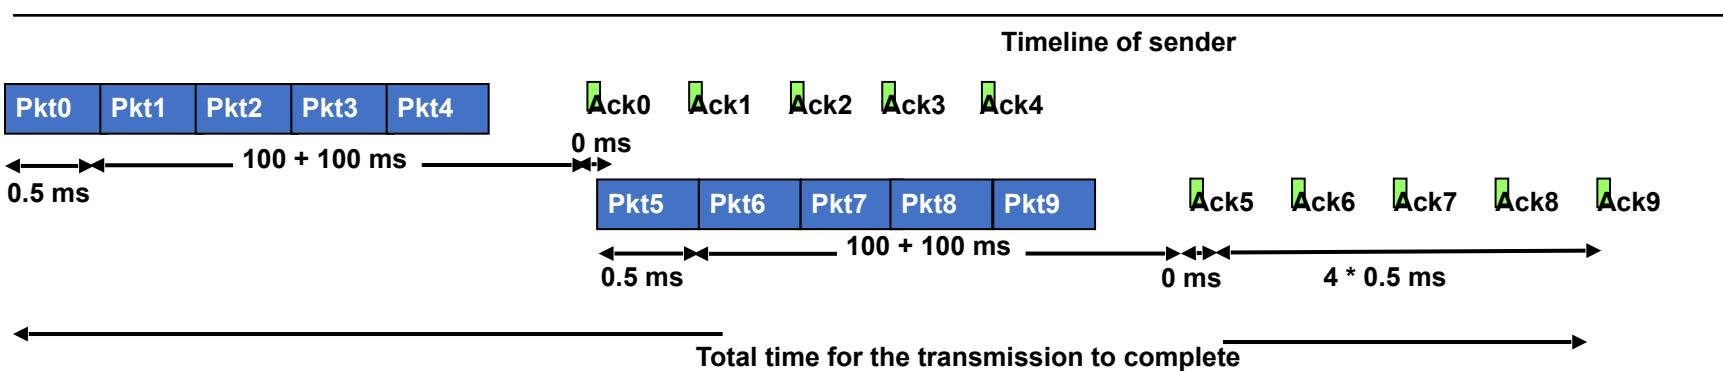
Total transmission time

$$= 0.5 + 100 + 100 + 0 + 0.5 + 100 + 100 + 0.5 + 0.5 + 0.5 + 0.5$$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

$$100.5 \text{ ms}$$

End-to-end latency for Ack0

$$= S + T_w + T_f + R$$

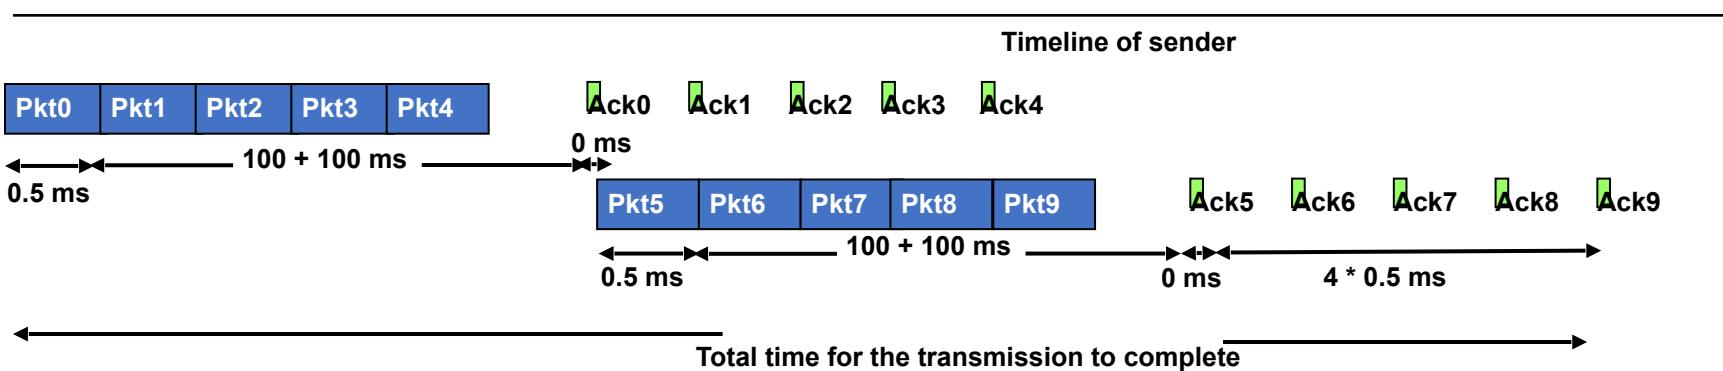
$$0 + 0.0 + 100 + 0$$

$$100.0 \text{ ms}$$

# Timeline example

100ms latency (flight time)  
10 packets  
Window size = 5

Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0



End-to-end latency for Pkt0

$$= S + T_w + T_f + R$$

$$0 + 0.5 + 100 + 0$$

$$100.5 \text{ ms}$$

End-to-end latency for Ack0

$$= S + T_w + T_f + R$$

$$0 + 0.0 + 100 + 0$$

$$100.0 \text{ ms}$$

Total transmission time

$$= 0.5 + 100 + 100 + 0 + 0.5 + 100 + 100 + 0.5 + 0.5 + 0.5 + 0.5$$

$$0.5 + 200 + 0.5 + 200 + 4 * 0.5$$

$$403 \text{ ms}$$

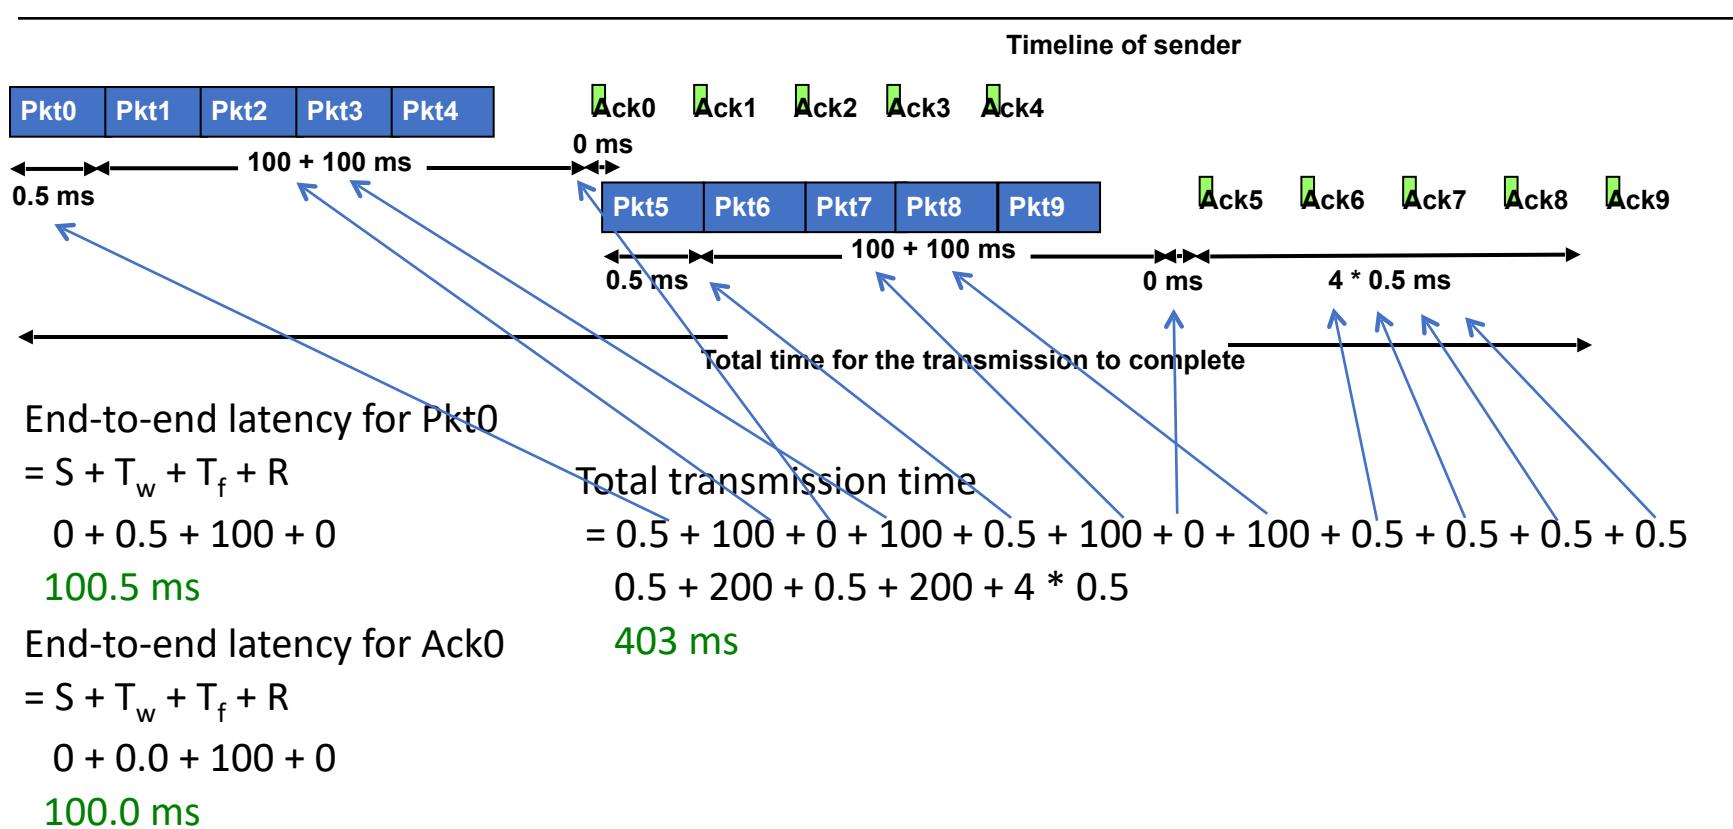


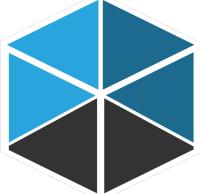
# Timeline example

Upcoming: dag2200

**100ms latency  
10 packets  
Window size = 5**

**Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0**



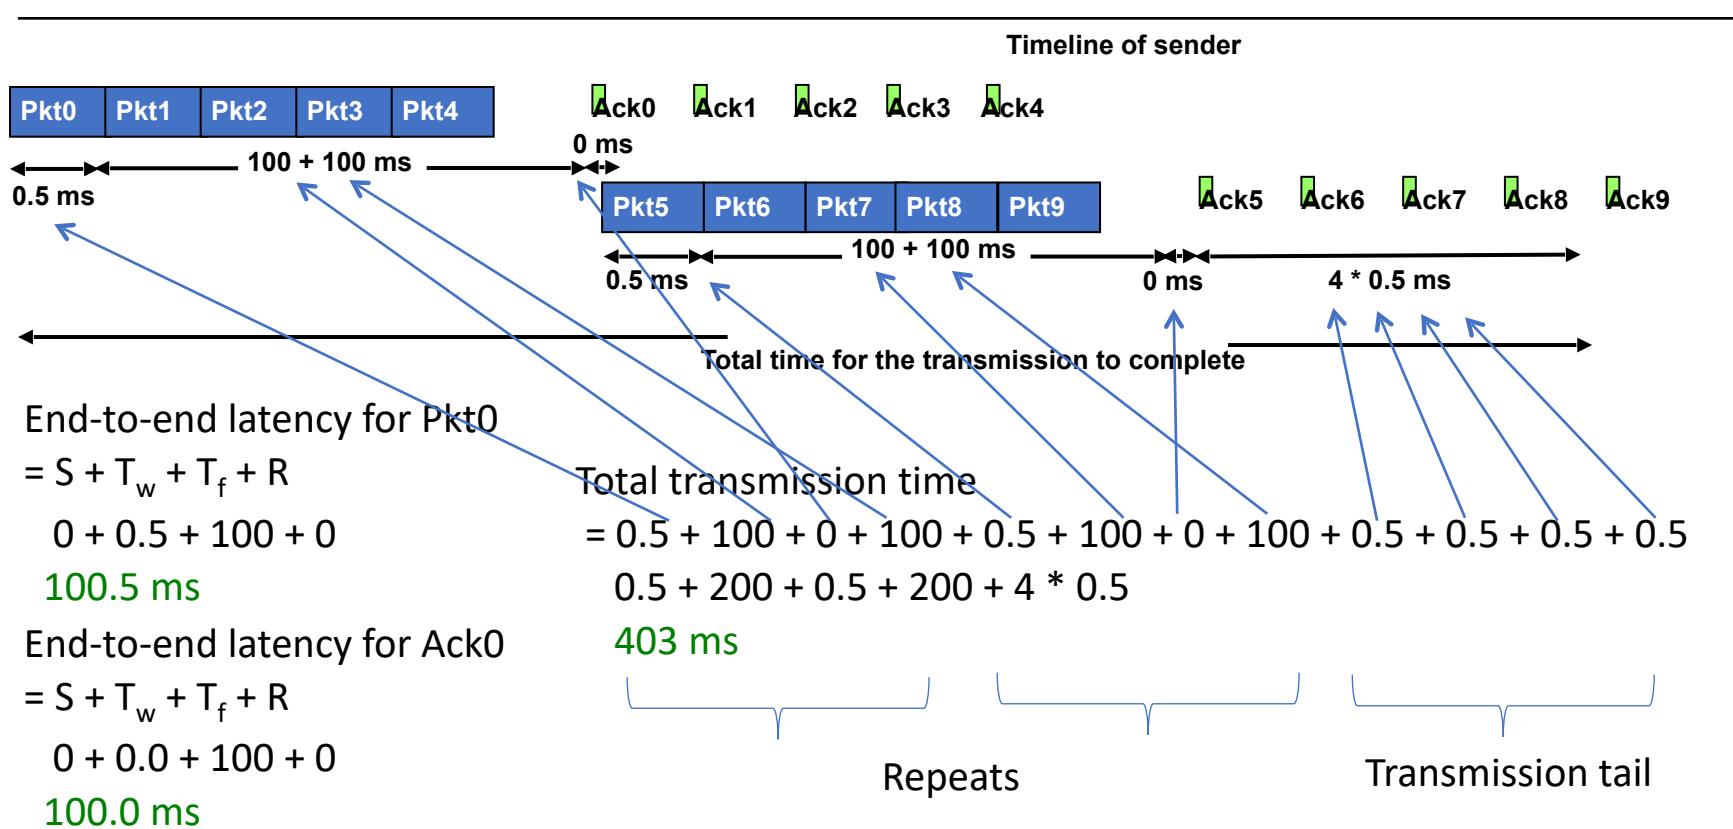


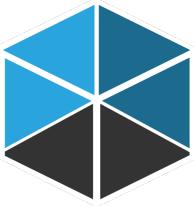
# Timeline example

Upcoming: dag2200

**100ms latency  
10 packets  
Window size = 5**

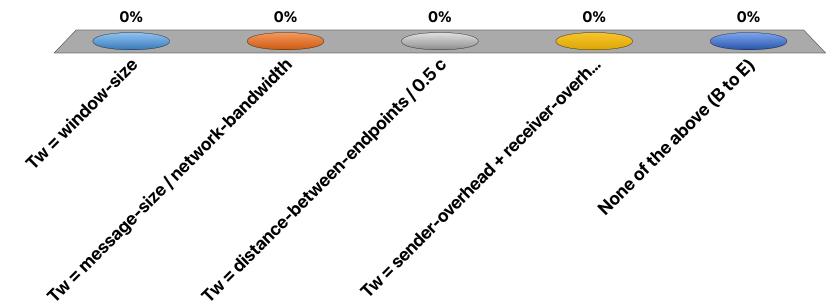
**Wire delay (data) = 0.5ms  
Wire delay (ack) = ~0  
Receiver & Sender overhead = ~0**

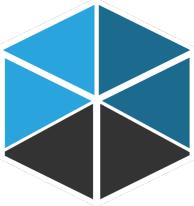




# In the expression $S + T_w + T_f + R$

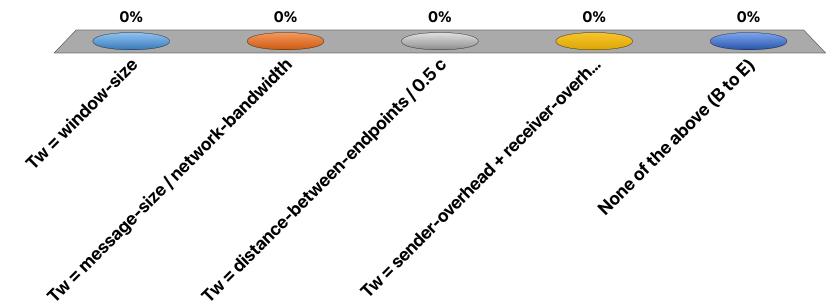
- A.  $T_w$  = window-size
- B.  $T_w$  = message-size / network-bandwidth
- C.  $T_w$  = distance-between-endpoints / 0.5 c
- D.  $T_w$  = sender-overhead + receiver-overhead
- E. None of the above (B to E)





# In the expression $S + T_w + T_f + R$

- A.  $T_w$  = window-size
- B.  $T_w$  = message-size / network-bandwidth
- C.  $T_w$  = distance-between-endpoints /  $0.5 c$
- D.  $T_w$  = sender-overhead + receiver-overhead  $T_f$   $S+R$
- E. None of the above (B to E)



# Examples

---

Assuming no loss, how many packets to transmit the message given?

- Message size = 100,000 bytes, header size per packet = 100 bytes, packet size = 1100 bytes

# Examples

---

Assuming no loss, how many packets to transmit the message given?

- Message size = 100,000 bytes, header size per packet = 100 bytes, packet size = 1100 bytes
  - 100 packets

# Examples

---

Assuming no loss, how many packets to transmit the message given?

- Message size = 100,000 bytes, header size per packet = 100 bytes, packet size = 1100 bytes
  - 100 packets

Assuming 10% loss, what is the number of packets to transmit the message?

# Examples

---

Assuming no loss, how many packets to transmit the message given?

- Message size = 100,000 bytes, header size per packet = 100 bytes, packet size = 1100 bytes
  - 100 packets

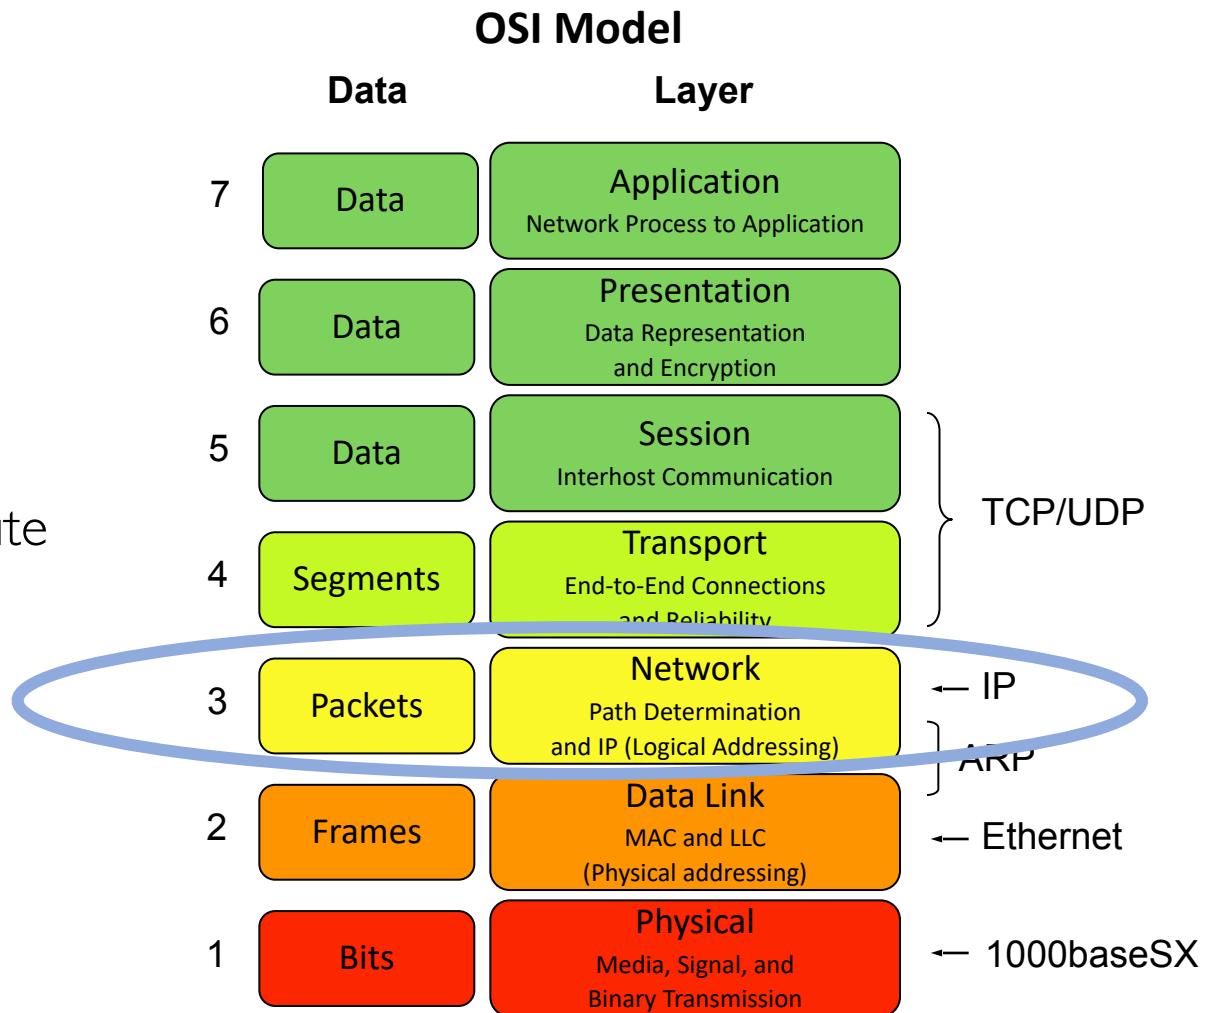
Assuming 10% loss, what is the number of packets to transmit the message?

- |||
- 100 + 10 + 1

More advanced performance estimation examples in the book!

# Back to the network layer

- Why are we back here?
- Recall, IP routing happens here
- We didn't answer a question:  
How does a router know where to route  
(i.e. who sets up the routing table)?



# Network layer

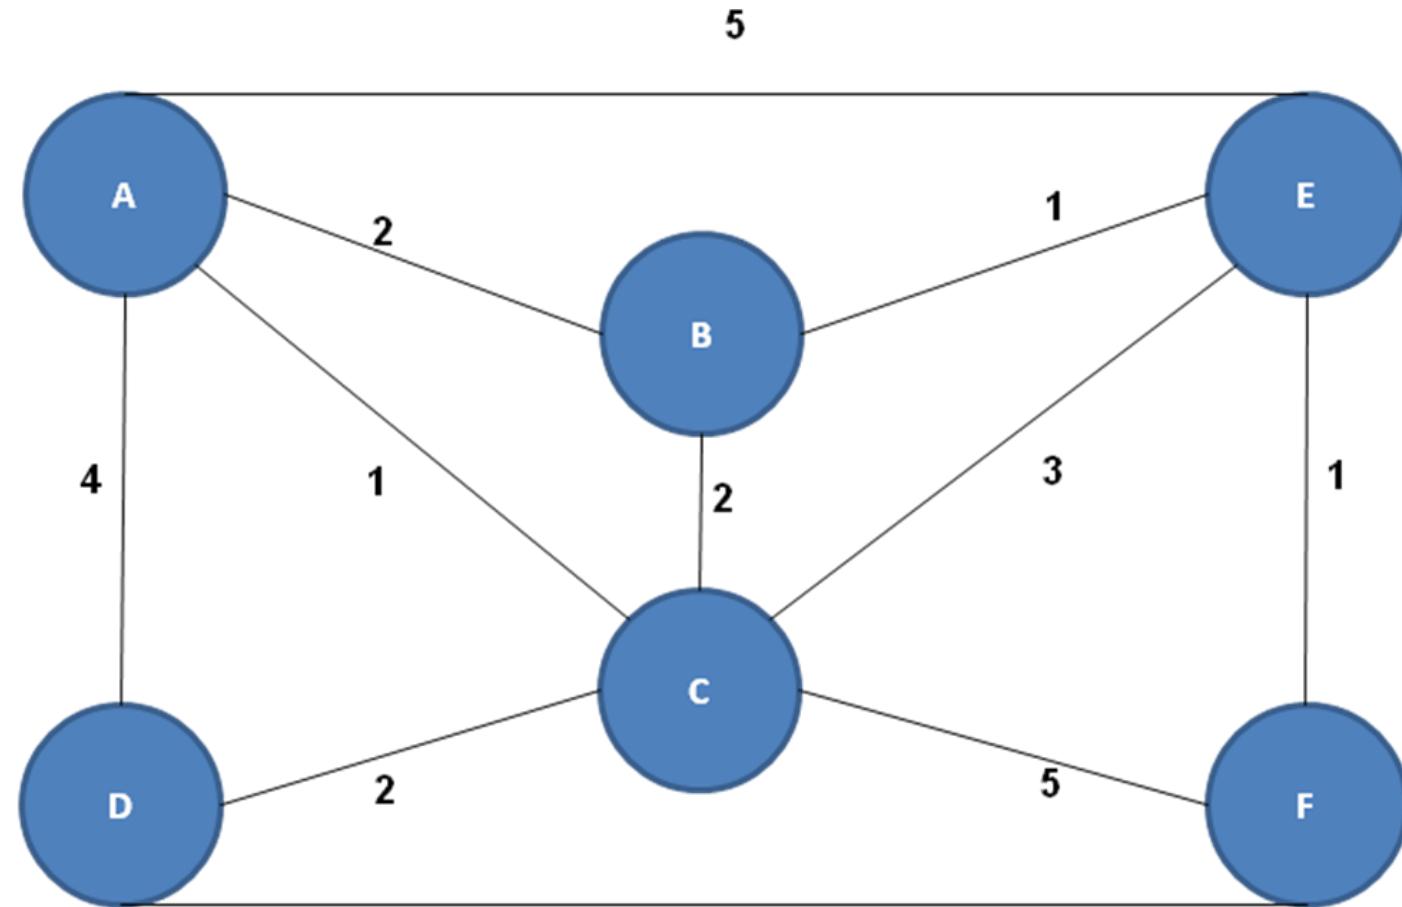
---

- Routing information protocols
  - Link state protocols
  - Distance vector protocols
  - Hierarchical routing
- Addressing

# Intuition behind routing algorithms (link state, distance vector)

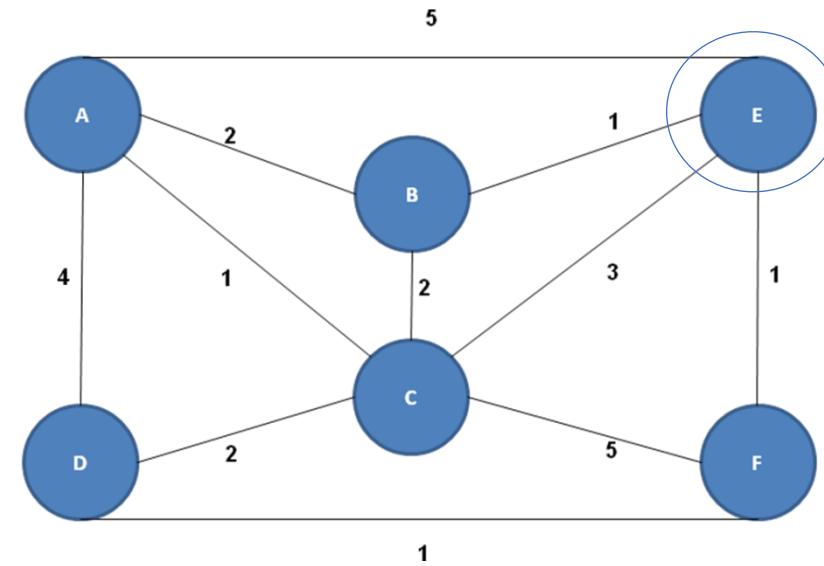
---

What's the latency to get to your neighbor?



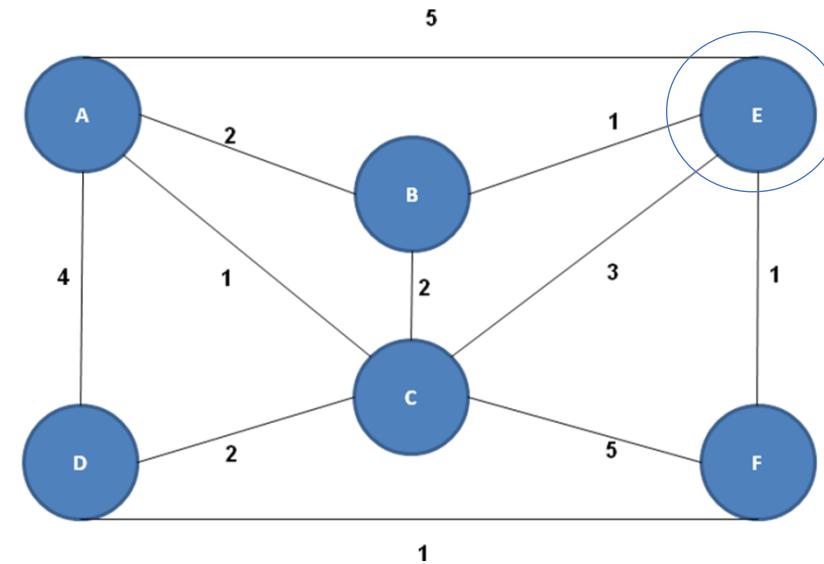
# Distance vector in action

- Uses neighbor info, local algorithm (i.e. every node only sees its neighbors)
- Each node knows the cost to each of its neighbors
- Each node sends its idea of routing to each of its neighbors
- Known colloquially as "routing by rumor"
- For every other node, each node determines the lowest cost next-hop
- Algorithm runs simultaneously on every node
- How does it work?



# Distance vector in action

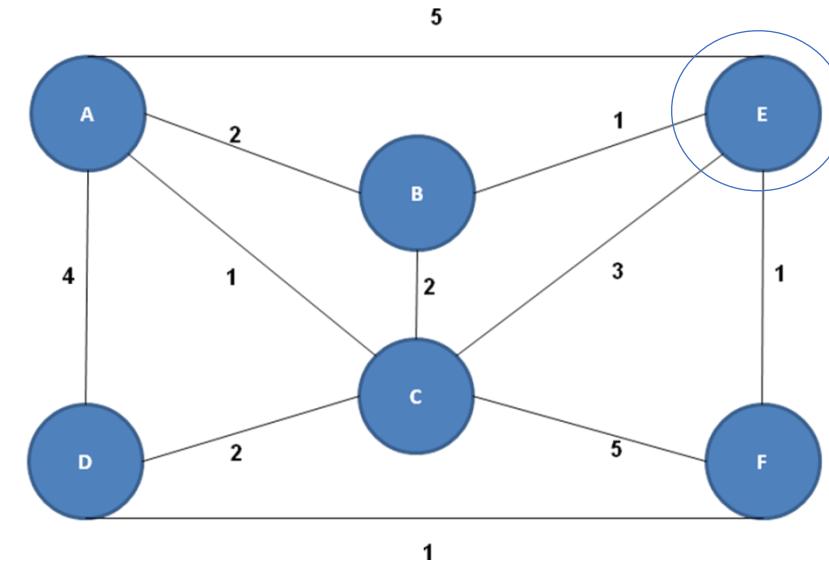
- Uses neighbor info, local algorithm (i.e. every node only sees its neighbors)
- Each node knows the cost to each of its neighbors
- Each node sends its idea of routing to each of its neighbors
- Known colloquially as "routing by rumor"
- For every other node, each node determines the lowest cost next-hop
- Algorithm runs simultaneously on every node
- How does it work?
- Node E knows its costs to its neighbors:



Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F

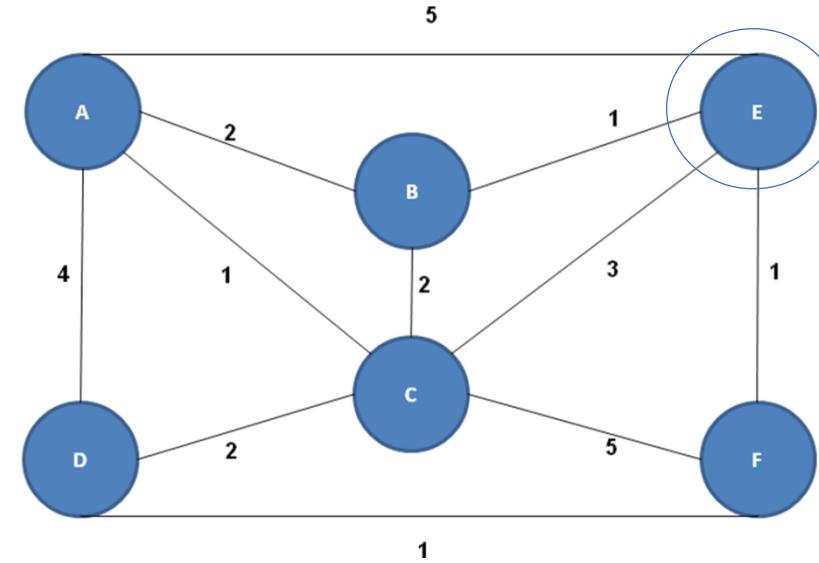
# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F



# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F

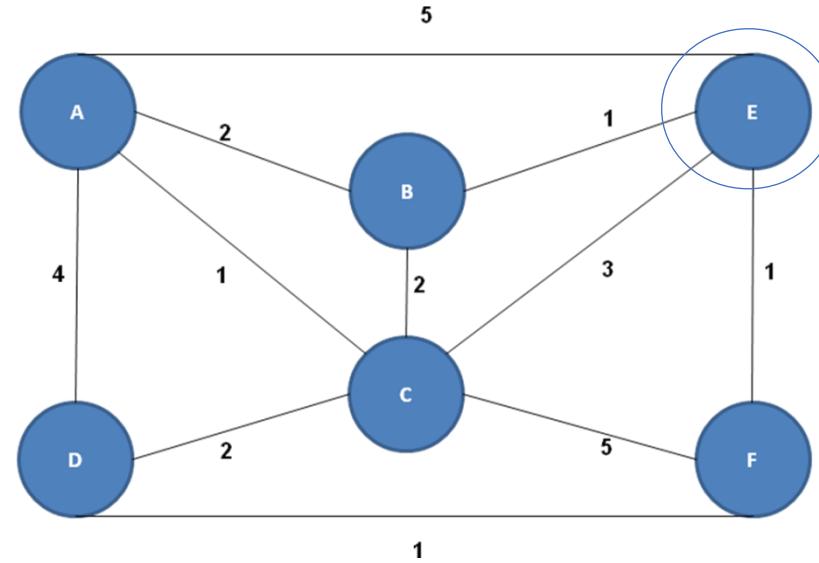


- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F

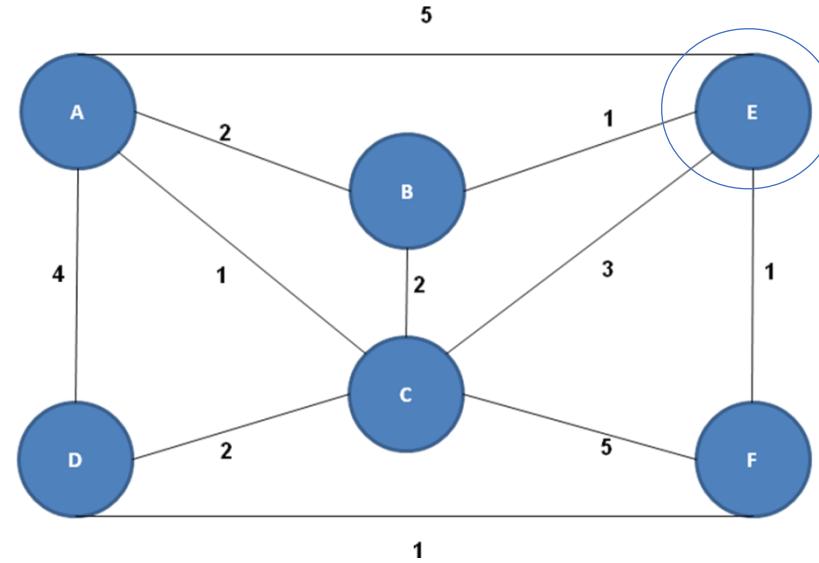
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.



# Distance vector in action

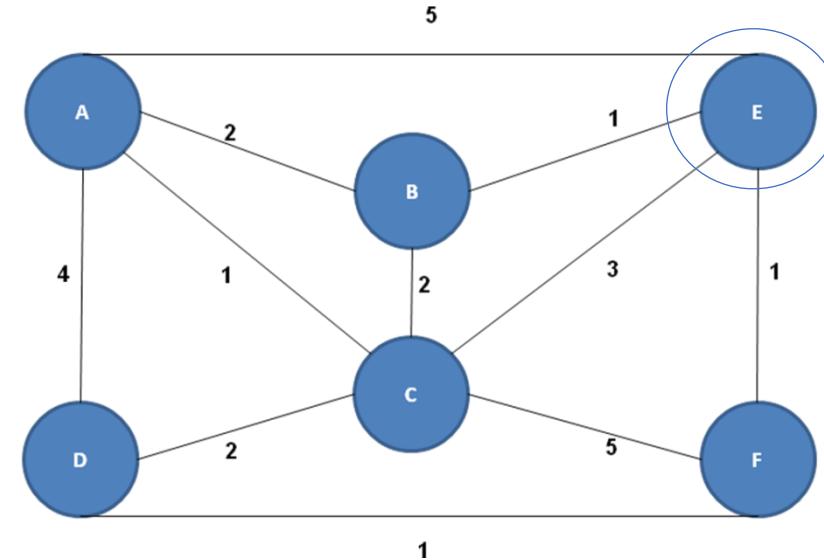
Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F



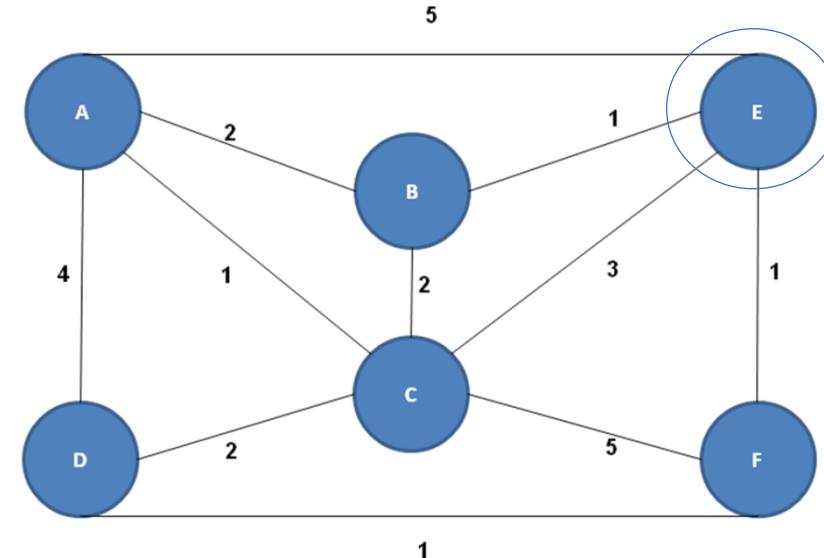
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing node A's message
  - $5+2/B$ ,  $5+1/C$ ,  $5+4/D$ ,  $5+5/F$
  - Only  $9/D$  is less
  - We'll set the next hop to A

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	$\infty/$	1/F

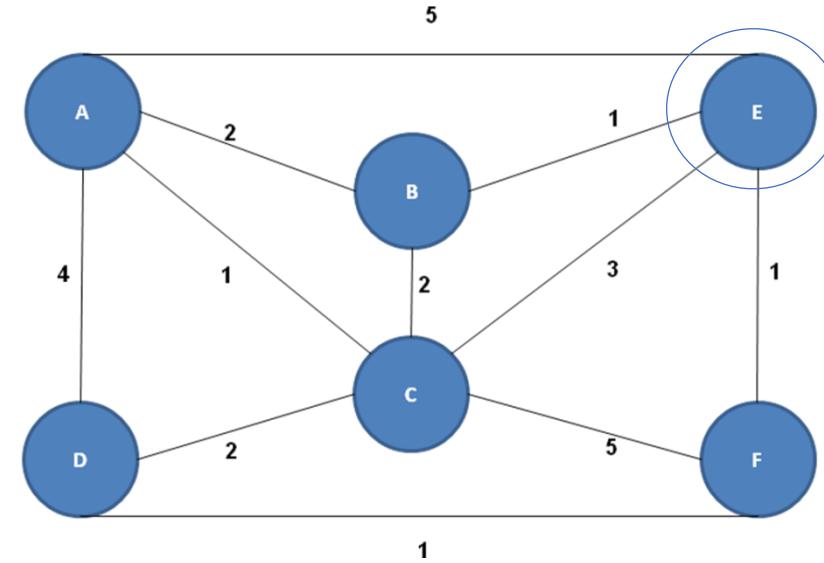
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



- Reviewing node A's message
  - $5+2/B, 5+1/C, 5+4/D, 5+5/F$
  - Only 9/D is less
  - We'll set the next hop to A

# Distance vector in action

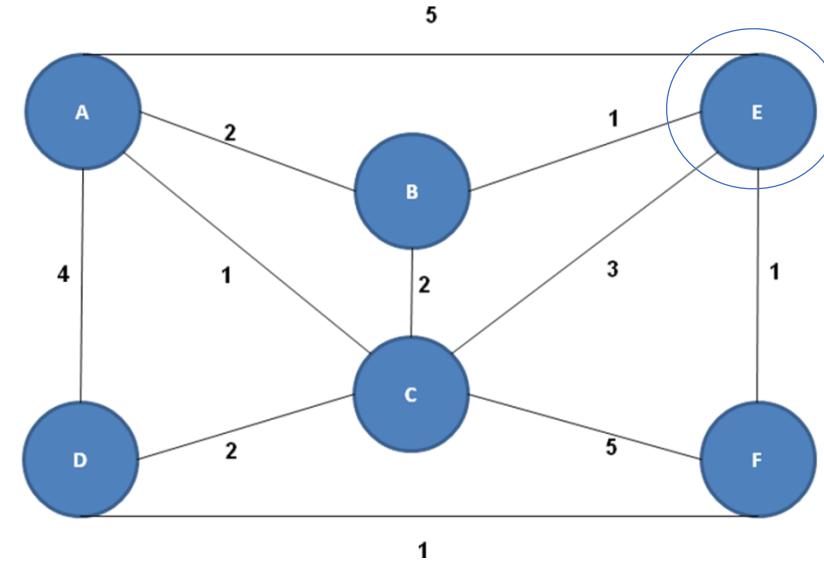
Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F



- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F

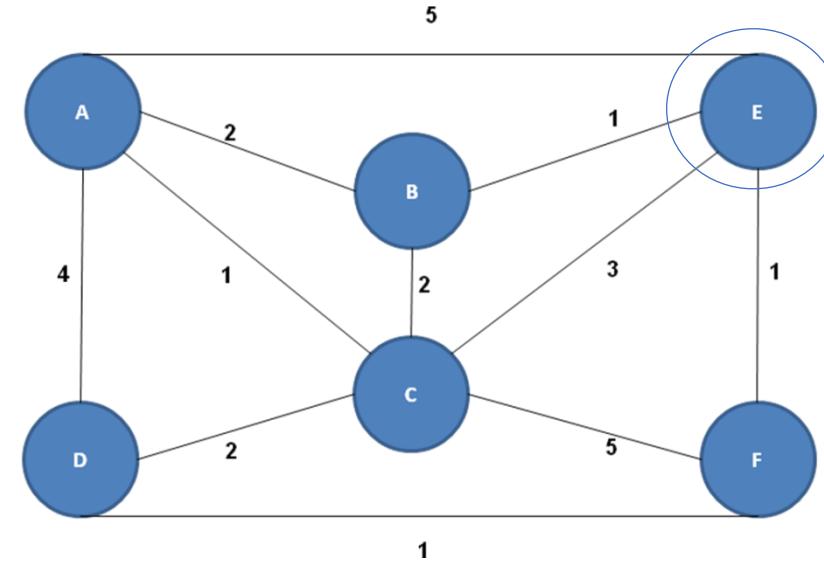


- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node B's message

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F

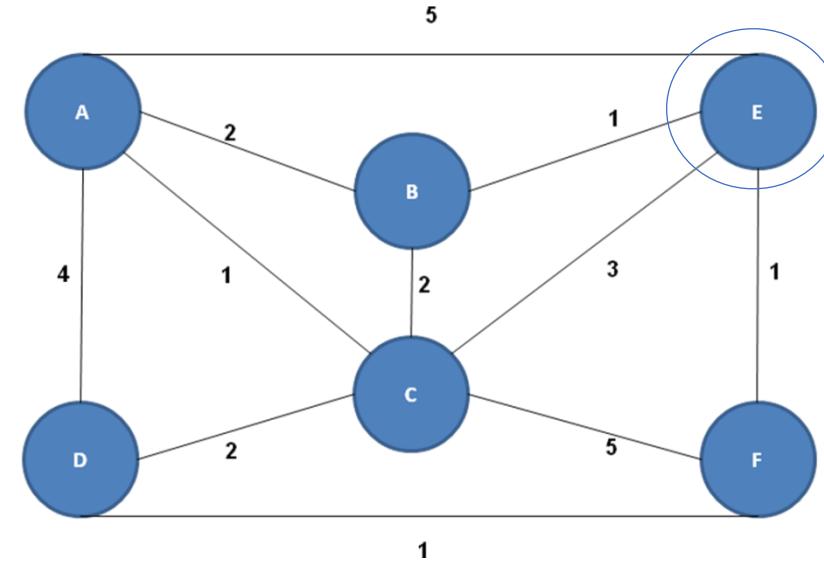


- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node B's message
  - 1+2/C, 1+2/A, 1+5/F, 1+4/D

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F

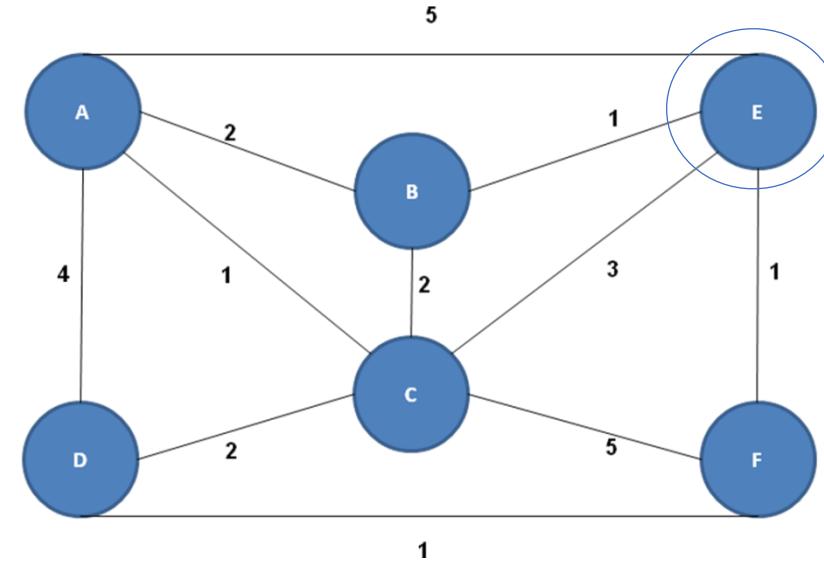


- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node B's message
  - $1+2/C, 1+2/A, 1+5/F, 1+4/D$
  - 3/A and 5/D are lower cost

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F

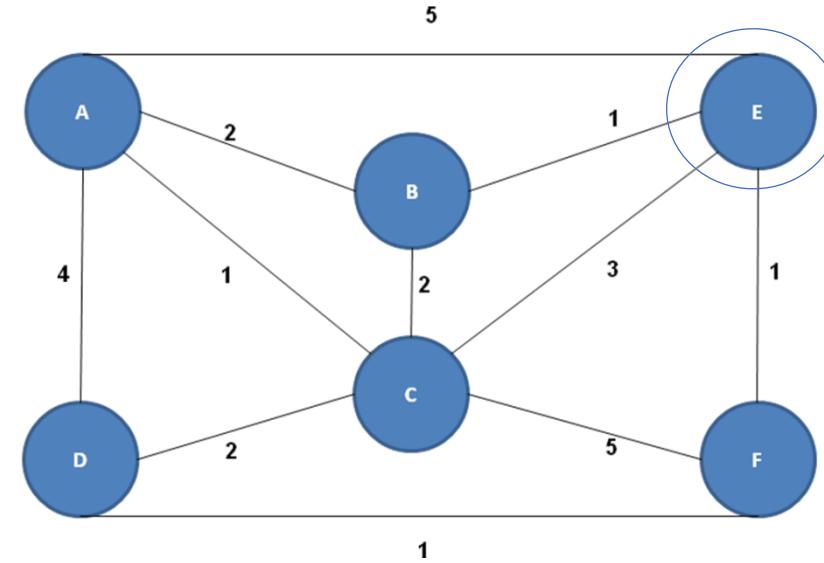


- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node B's message
  - $1+2/C, 1+2/A, 1+5/F, 1+4/D$
  - 3/A and 5/D are lower cost
  - We set their next hop to B

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F

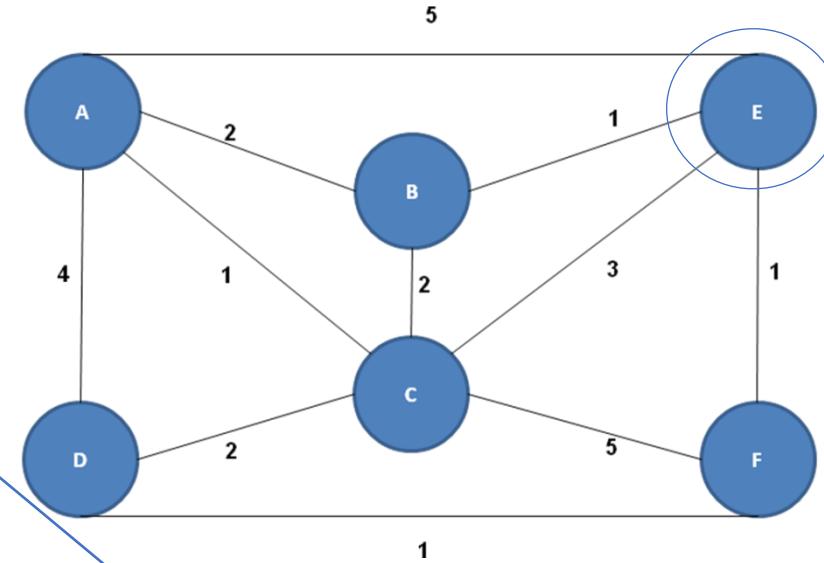


- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node B's message
  - $1+2/C, 1+2/A, 1+5/F, 1+4/D$
  - 3/A and 5/D are lower cost
  - We set their next hop to B

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	5/A	1/B	3/C	9/A	1/F



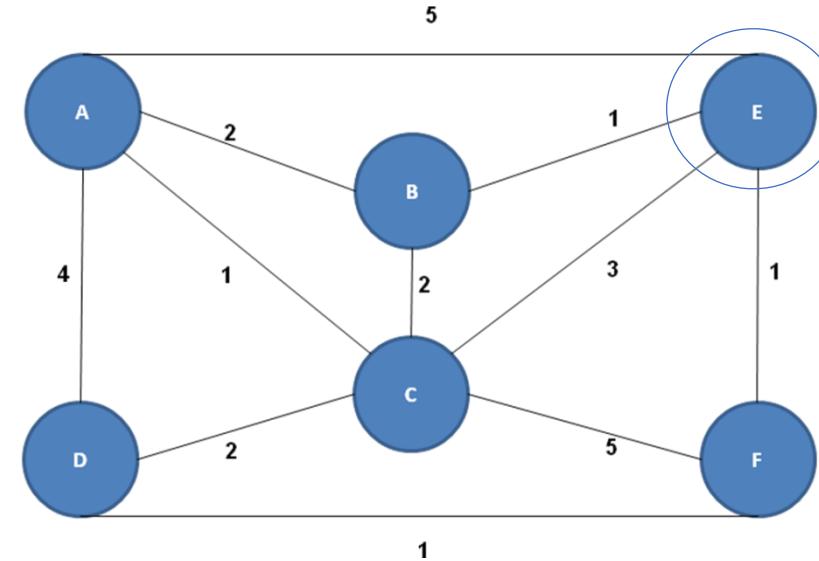
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node B's message
- $1+2/C, 1+2/A, 1+5/F, 1+4/D$
  - 3/A and 5/D are lower cost
  - We set their next hop to B

# Distance vector in action

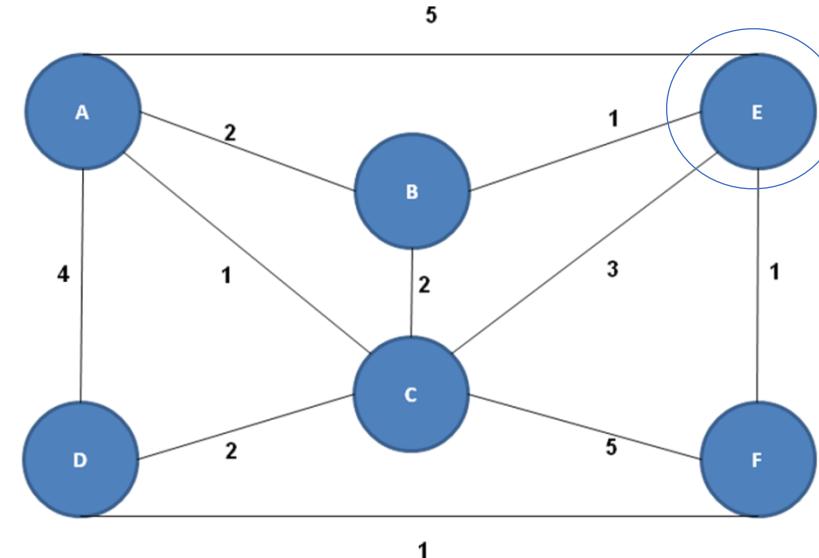
Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F



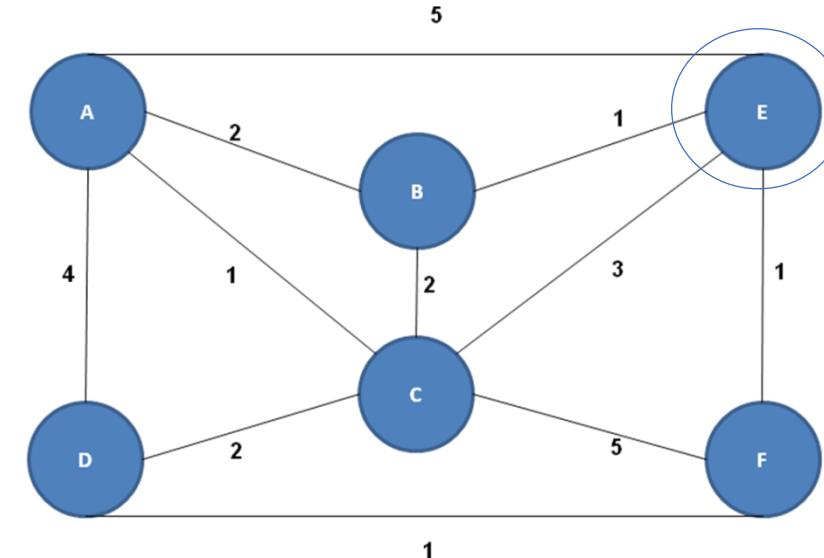
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

- Reviewing Node C's message

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

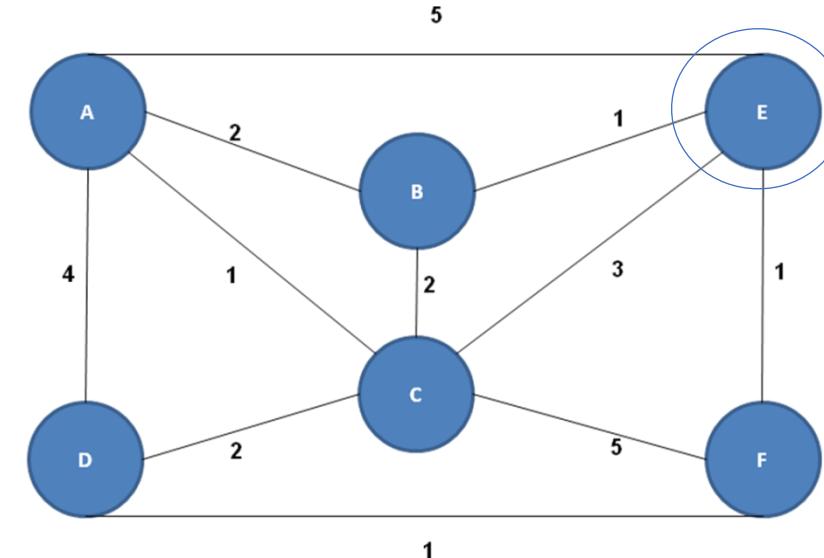


- Reviewing Node C's message
  - $3+1/A, 3+2/D, 3+5/F, 3+2/B$

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

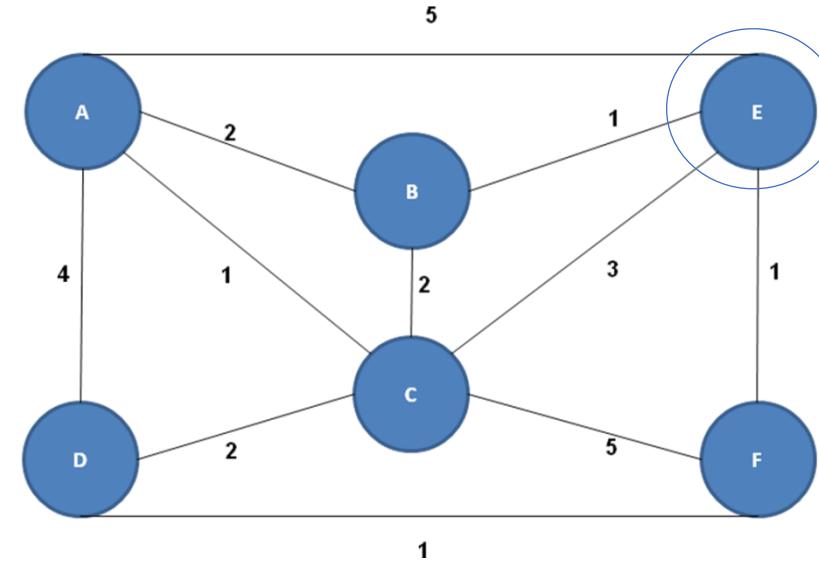


- Reviewing Node C's message
  - $3+1/A$ ,  $3+2/D$ ,  $3+5/F$ ,  $3+2/B$
  - None are lower cost

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

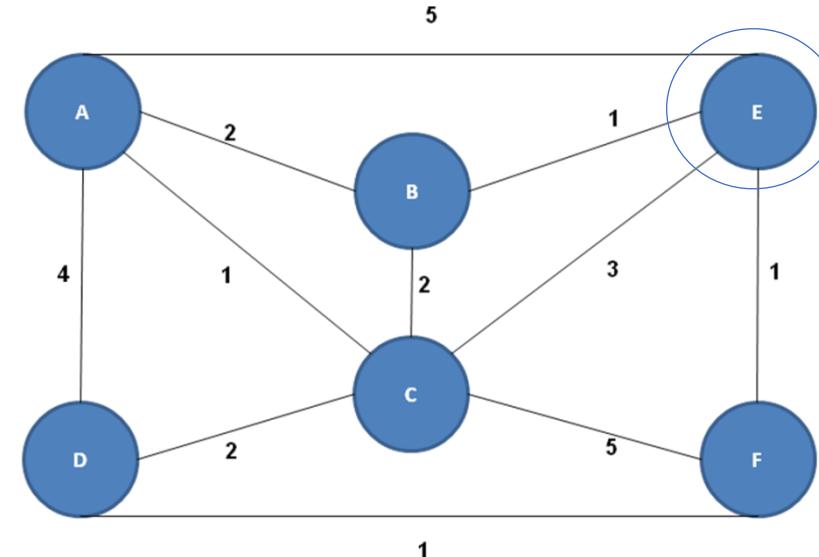
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

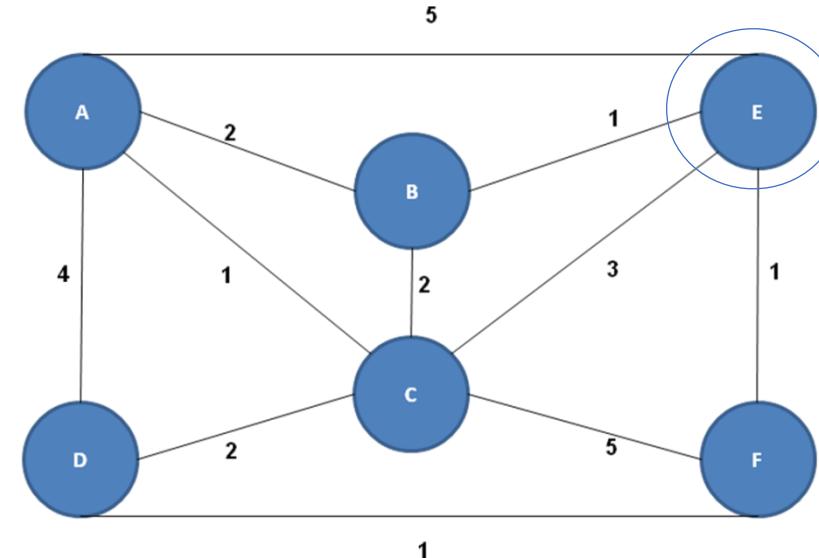


- Reviewing Node F's message

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

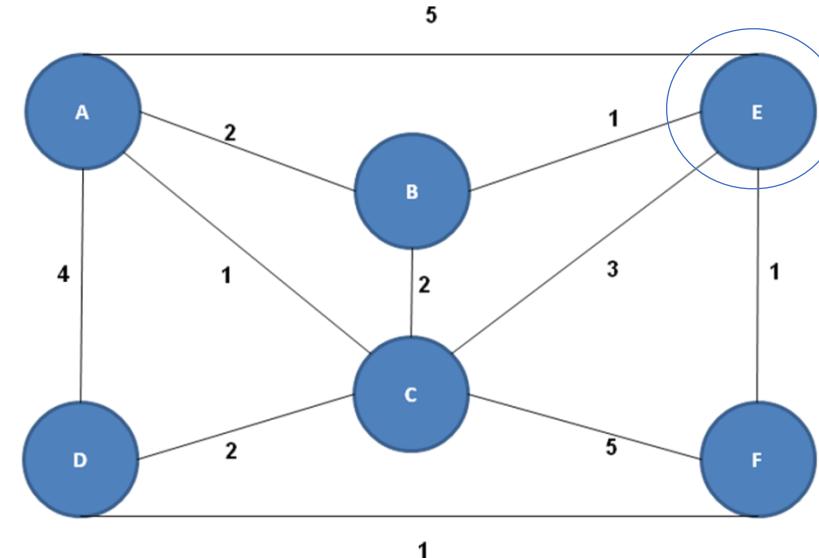


- Reviewing Node F's message
  - $1+5C, 1+1/D, 1+2/B, 1+4/A$

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

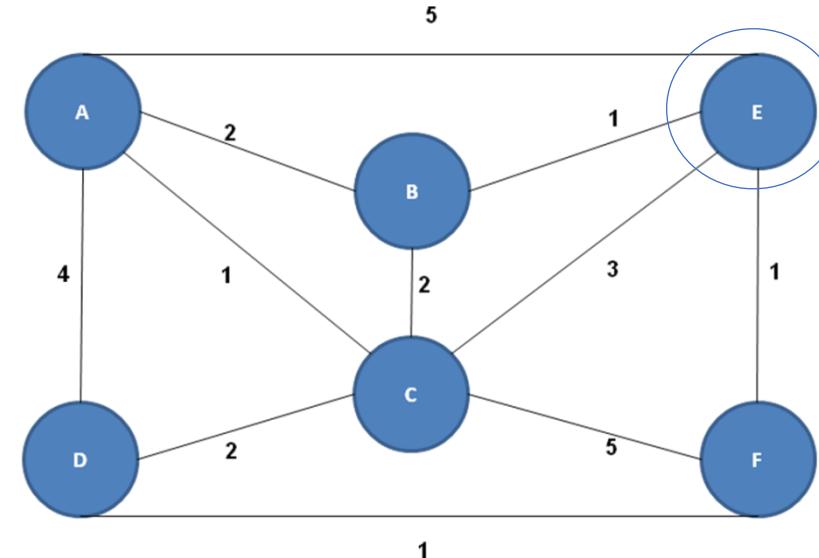


- Reviewing Node F's message
  - $1+5C, 1+1/D, 1+2/B, 1+4/A$
  - $2/D$  is lower cost

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F

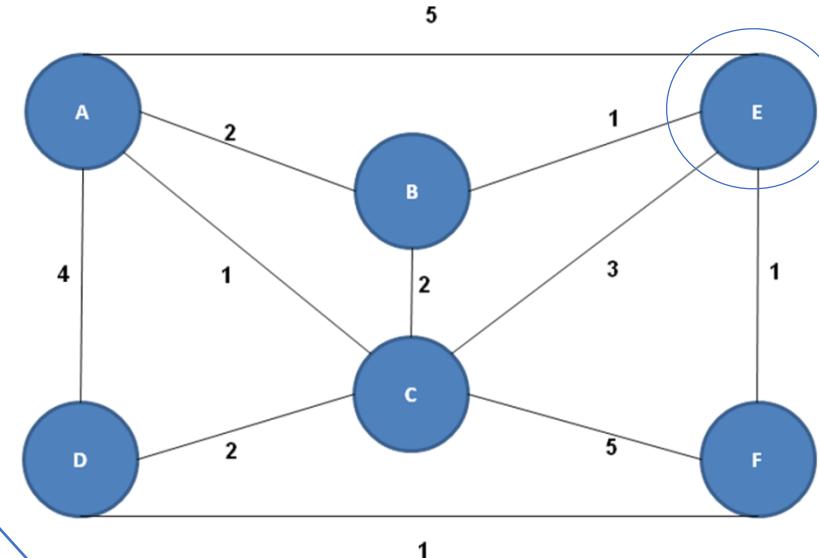
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



- Reviewing Node F's message
  - $1+5C, 1+1/D, 1+2/B, 1+4/A$
  - $2/D$  is lower cost
  - We'll set the next hop to F

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	5/B	1/F



- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

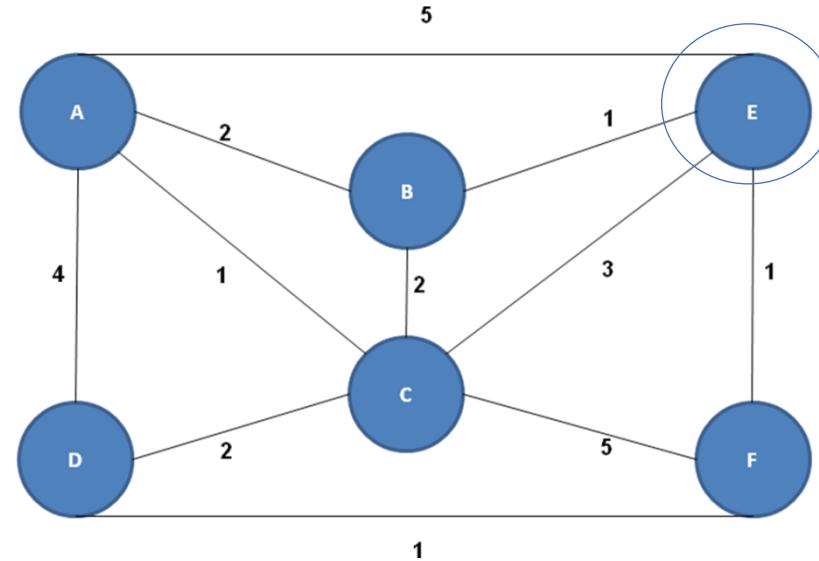
## Reviewing Node F's message

- $1+5C, 1+1/D, 1+2/B, 1+4/A$
- 2/D is lower cost
- We'll set the next hop to F

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	2/F	1/F

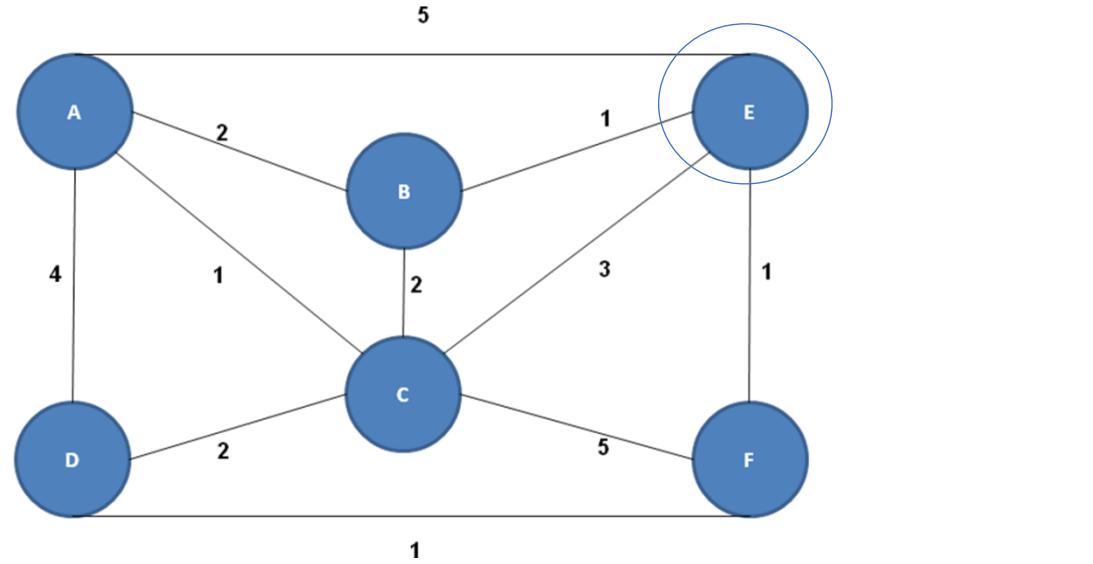
- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	2/F	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

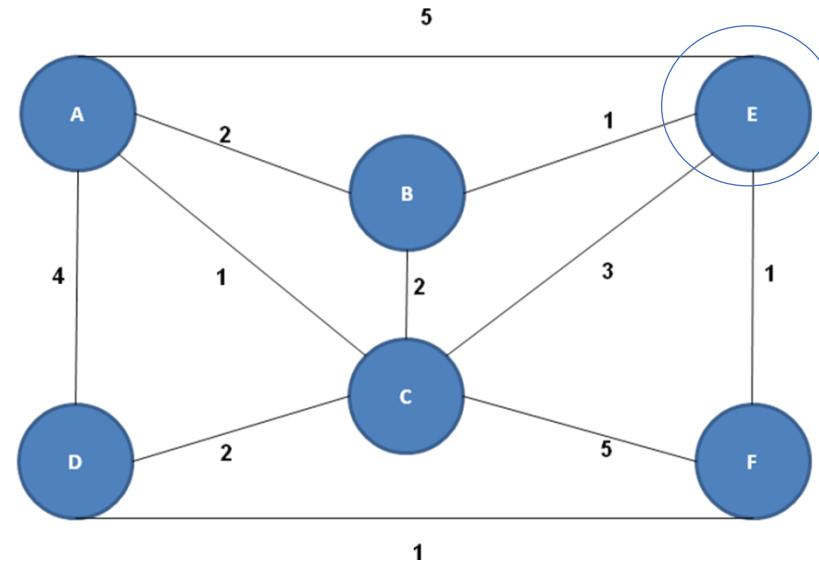


- Do you agree that E now knows the next hop for the lowest cost routes to the other nodes?

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	2/F	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor

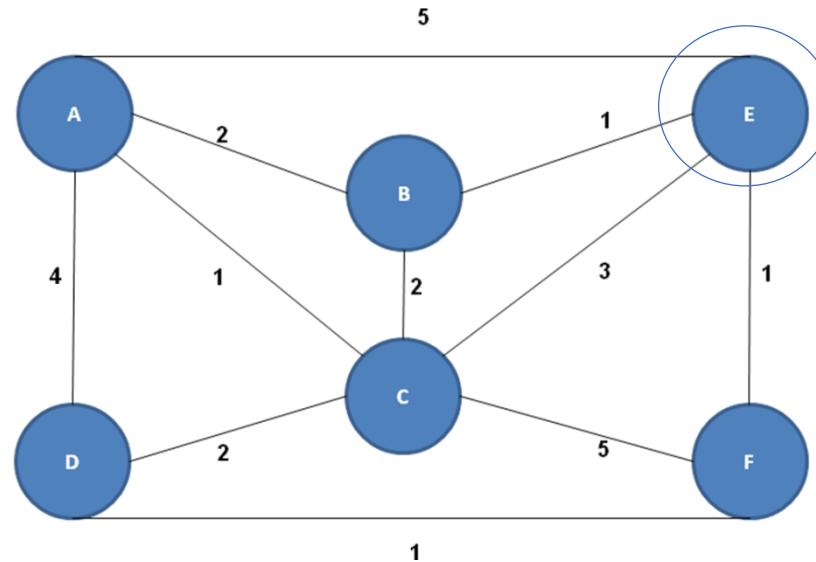


- Do you agree that E now knows the next hop for the lowest cost routes to the other nodes?
- This of course presumed the other nodes knew their lowest cost routes

# Distance vector in action

Routes from E	A	B	C	D	F
Cost/next hop	3/B	1/B	3/C	2/F	1/F

- Node E listens to messages from its neighbors:
  - A says its routes are 0/A, 5/E, 2/B, 1/C, 4/D, 5/F
  - B says its routes are 0/B, 1/E, 2/C, 2/A, 5/F, 4/D
  - C says its routes are 0/C, 1/A, 2/D, 5/F, 3/E, 2/B
  - F says its routes are 0/F, 5/C, 1/D, 1/E, 2/B, 4/A
  - D isn't heard because it isn't connected to E
- Node E looks at each route from each neighbor N; it adds its cost to get to N to N's cost to get to each other node.
- For each node, if the sum is lower cost than what E has in its route table, E changes its own route to run through the neighbor



- Do you agree that E now knows the next hop for the lowest cost routes to the other nodes?
- This of course presumed the other nodes knew their lowest cost routes
- When this runs on all nodes simultaneously, they will **converge** on this solution in a relatively short time (convergence depends on diameter of the network)

# Distance Vector vs Link State Routing

---

# Distance Vector vs Link State Routing

---

- Distance Vector
  - Each node knows the cost to reach each of its neighbors
  - Each node sends its routing table to only its neighbors
  - Each node revises its own routing table every time it receives a routing table from a neighbor
  - Slower convergence but lower traffic

# Distance Vector vs Link State Routing

---

- Distance Vector
  - Each node knows the cost to reach each of its neighbors
  - Each node sends its routing table to only its neighbors
  - Each node revises its own routing table every time it receives a routing table from a neighbor
  - Slower convergence but lower traffic
- Link State
  - Each node advertises its neighbor links and costs to every other node in the network
  - Each node collects the entire topology of the network from these advertisements
  - Using a “shortest path” algorithm, each node calculates its own routing table
  - Faster convergence but more traffic

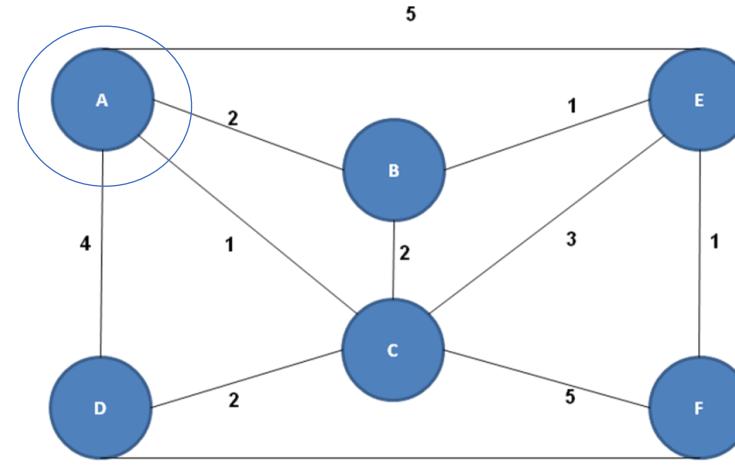
# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations (for  $n$ -node network)
- One least-cost route to a destination
- Runs simultaneously on each node

## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1						
2						
3						
4						
5						



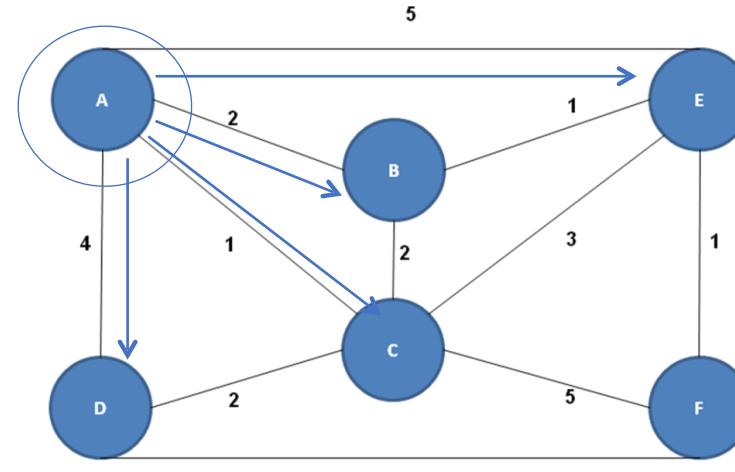
# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations (for  $n$ -node network)
- One least-cost route to a destination
- Runs simultaneously on each node

## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1						
2						
3						
4						
5						



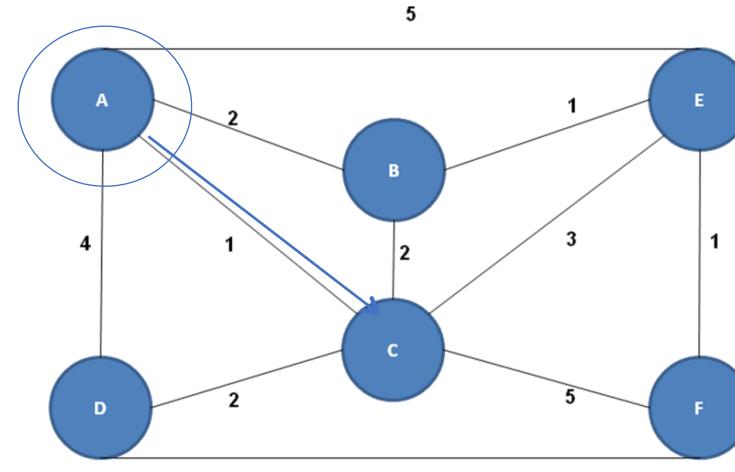
# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations (for  $n$ -node network)
- One least-cost route to a destination
- Runs simultaneously on each node

## Calculating on Node A

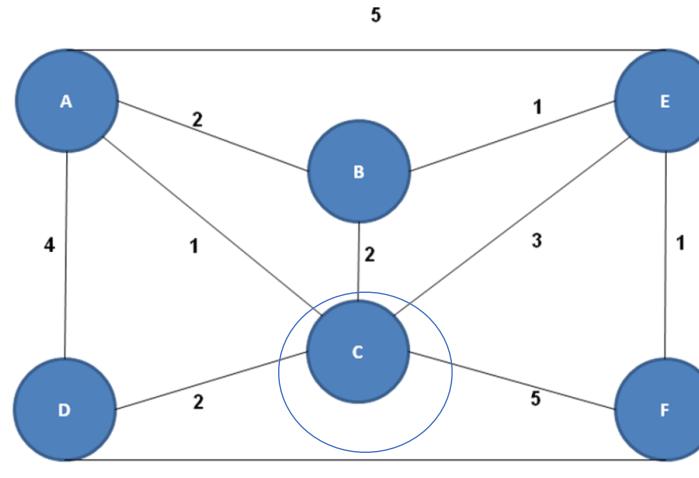
(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1						
2						
3						
4						
5						



# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



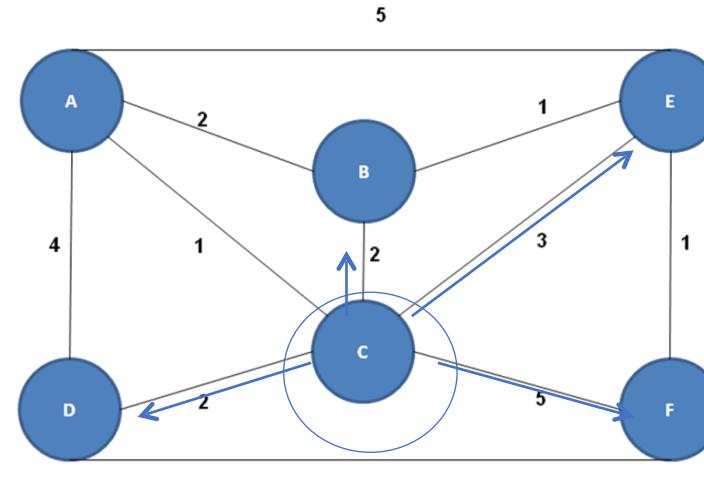
## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2						
3						
4						
5						

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



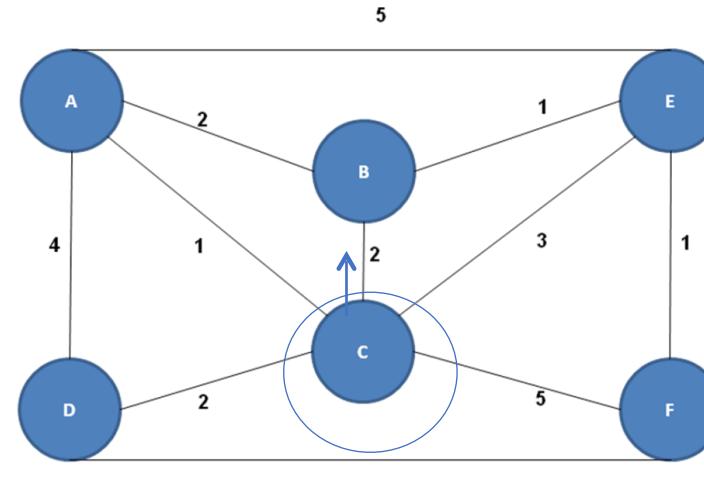
## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2						
3						
4						
5						

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



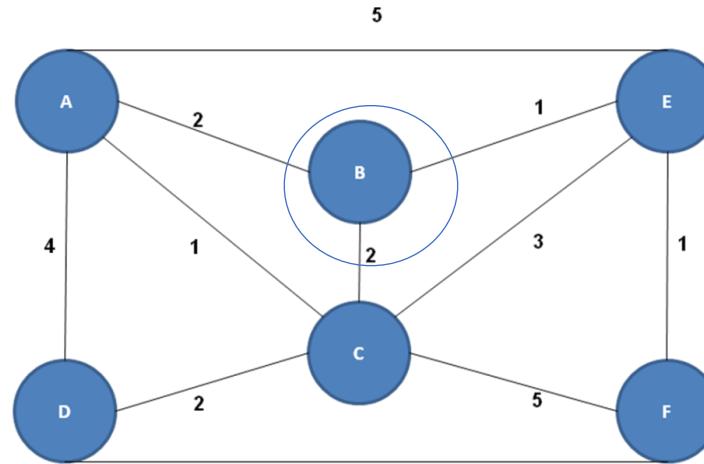
## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2						
3						
4						
5						

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



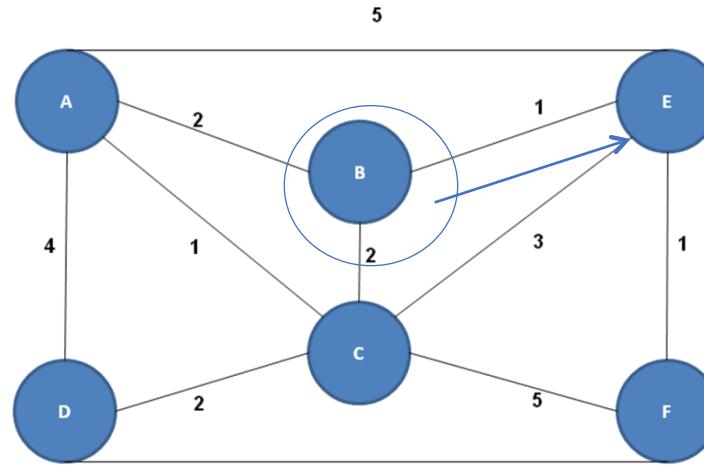
## Calculating on Node A

*(each node does its own calculation)*

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3						
4						
5						

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



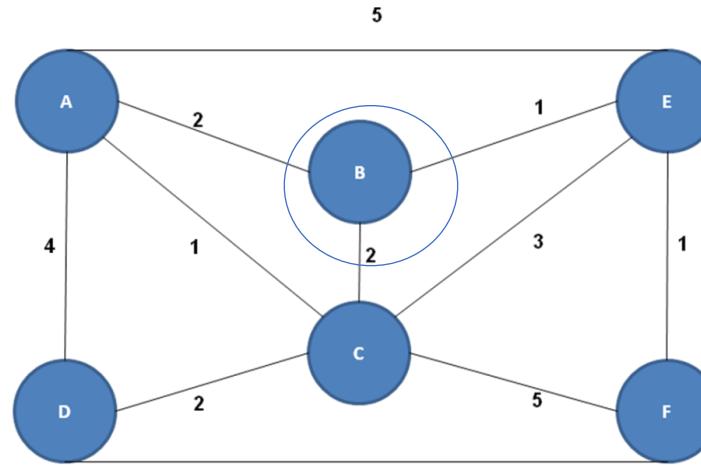
## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3						
4						
5						

# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3						
4						
5						

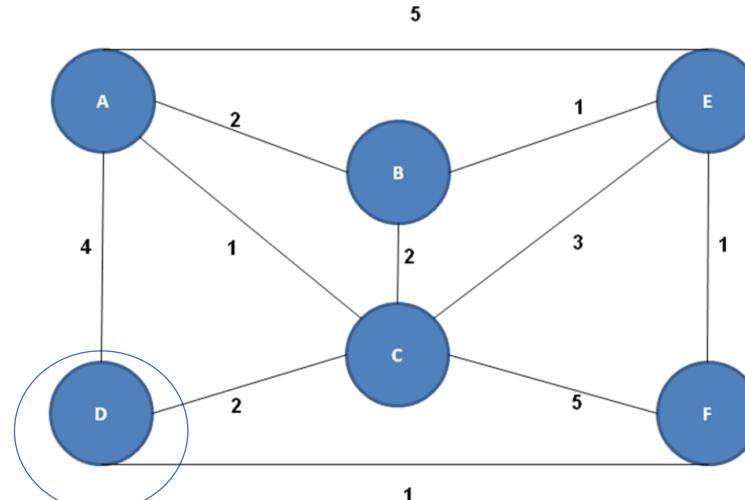
# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node

## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3	A,C,B,D	-	-	3/ACD	3/ABE	4/ACDF
4						
5						



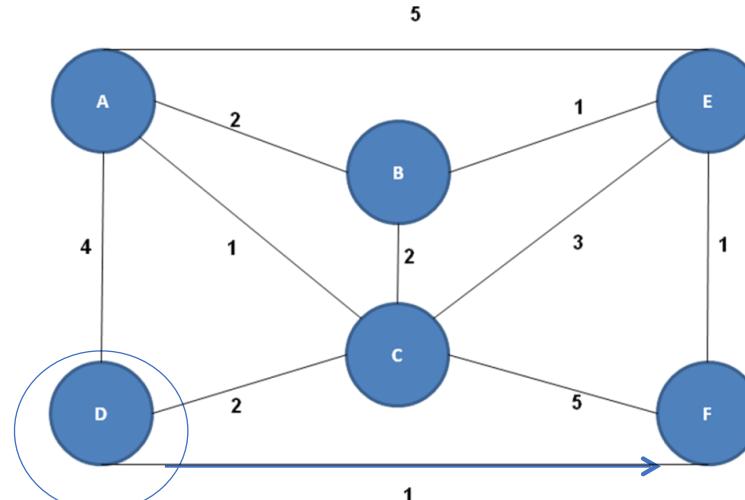
# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node

## Calculating on Node A

(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3	A,C,B,D	-	-	3/ACD	3/ABE	4/ACDF
4						
5						



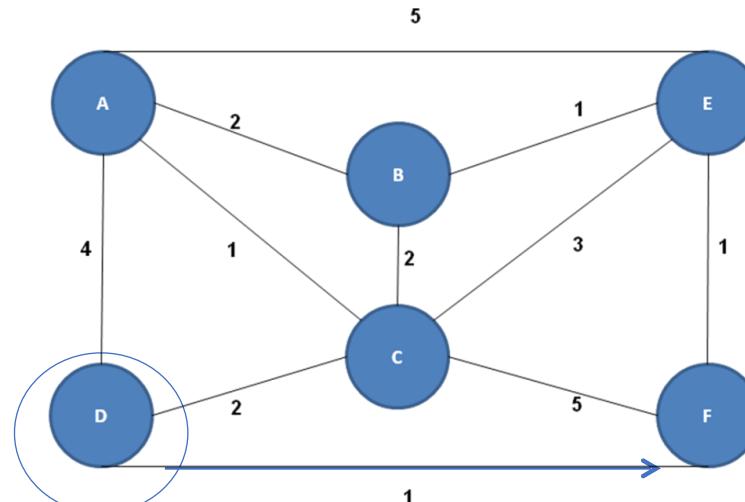
# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node

## Calculating on Node A

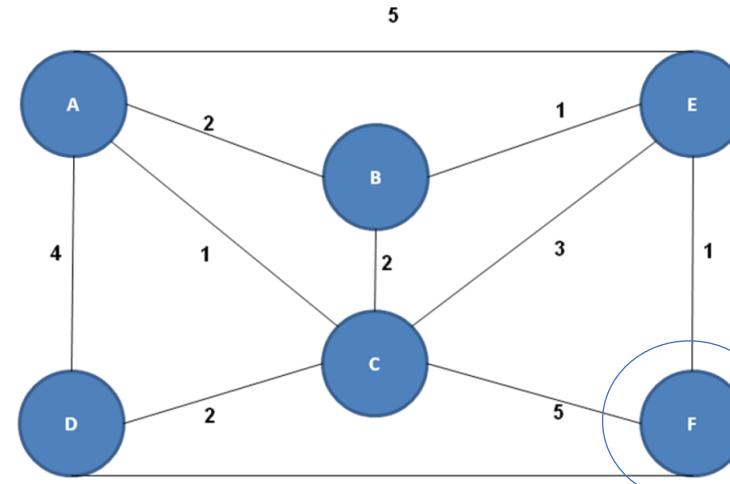
(each node does its own calculation)

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3	A,C,B,D	-	-	3/ACD	3/ABE	4/ACDF
4						
5						



# Link state in action

- Dijkstra's shortest path algorithm
- Uses global info, local algorithm (i.e. every node has this graph)
- $n-1$  iterations
- One least-cost route to a destination
- Runs simultaneously on each node



## Calculating on Node A

*(each node does its own calculation)*

Iteration Count	Nodes to which least-cost routes known	B Cost/route	C Cost/route	D cost/route	E Cost/route	F Cost/route
Init	A	2/AB	1/AC	4/AD	5/AE	$\infty$
1	A,C	2/AB	1/AC	3/ACD	4/ACE	6/ACF
2	A,C,B	2/AB	-	3/ACD	3/ABE	6/ACF
3	A,C,B,D	-	-	3/ACD	3/ABE	4/ACDF
4	A,C,B,D,E	-	-	-	3/ABE	4/ACDF
5	A,C,B,D,E,F	-	-	-	-	4/ACDF

# Hierarchical routing algorithms

---

# Hierarchical routing algorithms

---

- Network of networks

# Hierarchical routing algorithms

---

- Network of networks
- Very large scale, frequent state changes

# Hierarchical routing algorithms

---

- Network of networks
- Very large scale, frequent state changes
- Compartmentalize so routing information protocol issues don't spread and cause widespread outages

# Hierarchical routing algorithms

---

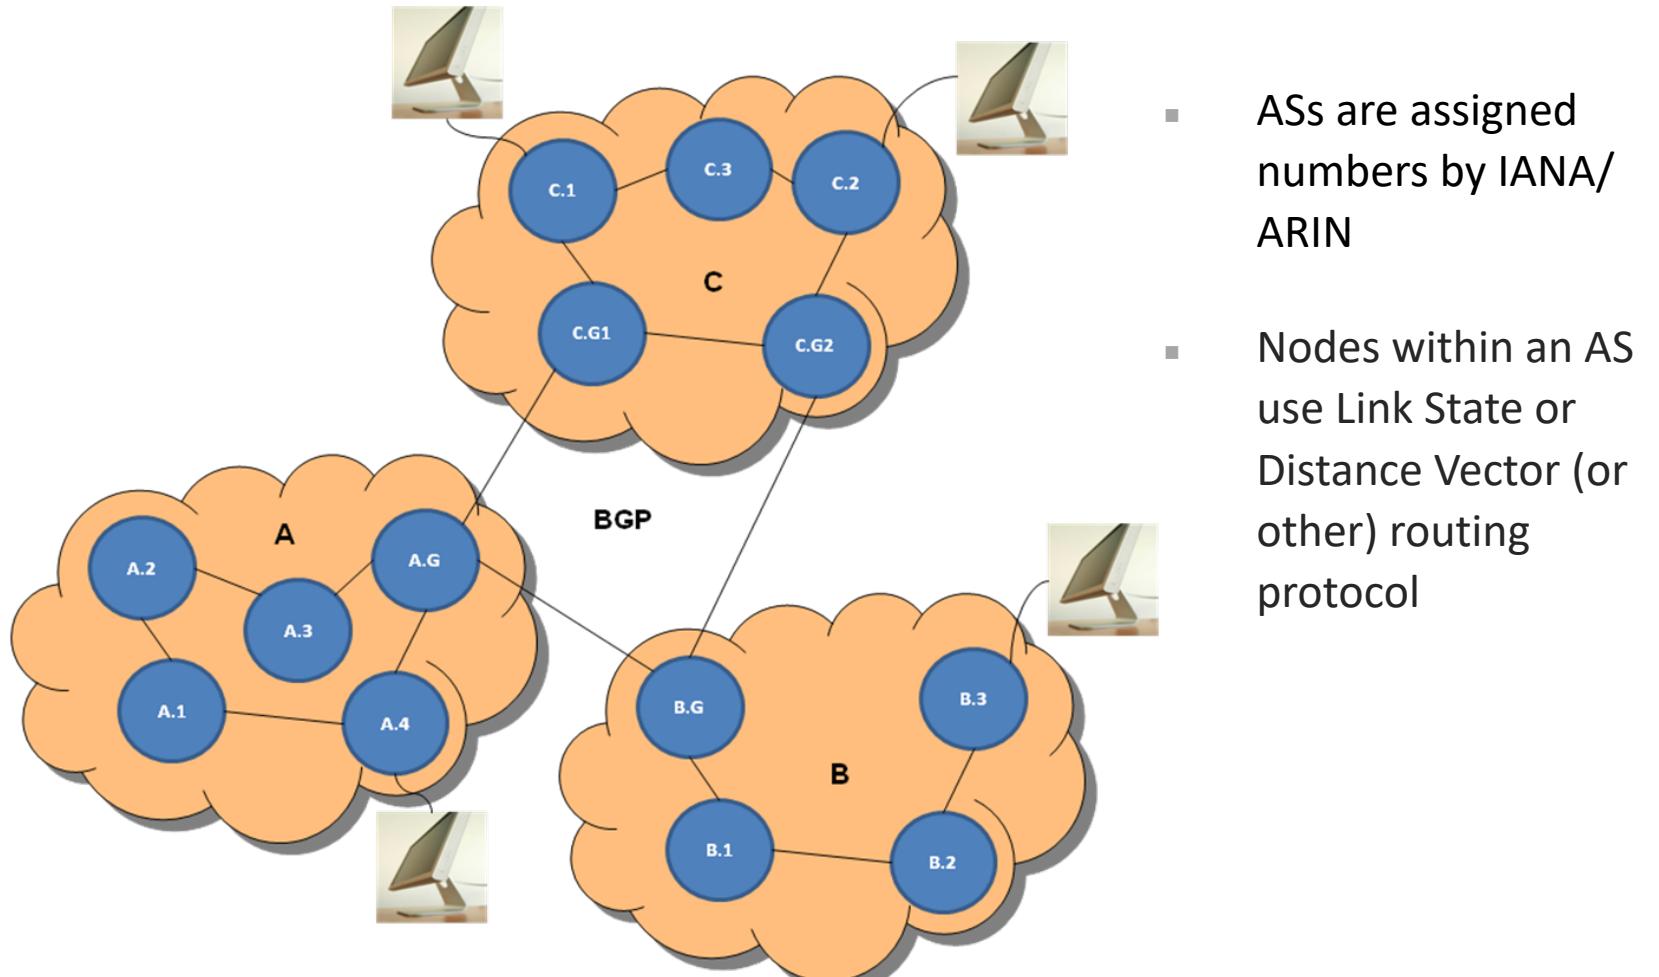
- Network of networks
- Very large scale, frequent state changes
- Compartmentalize so routing information protocol issues don't spread and cause widespread outages
- What better example than the Internet itself!

# Hierarchical routing algorithms

---

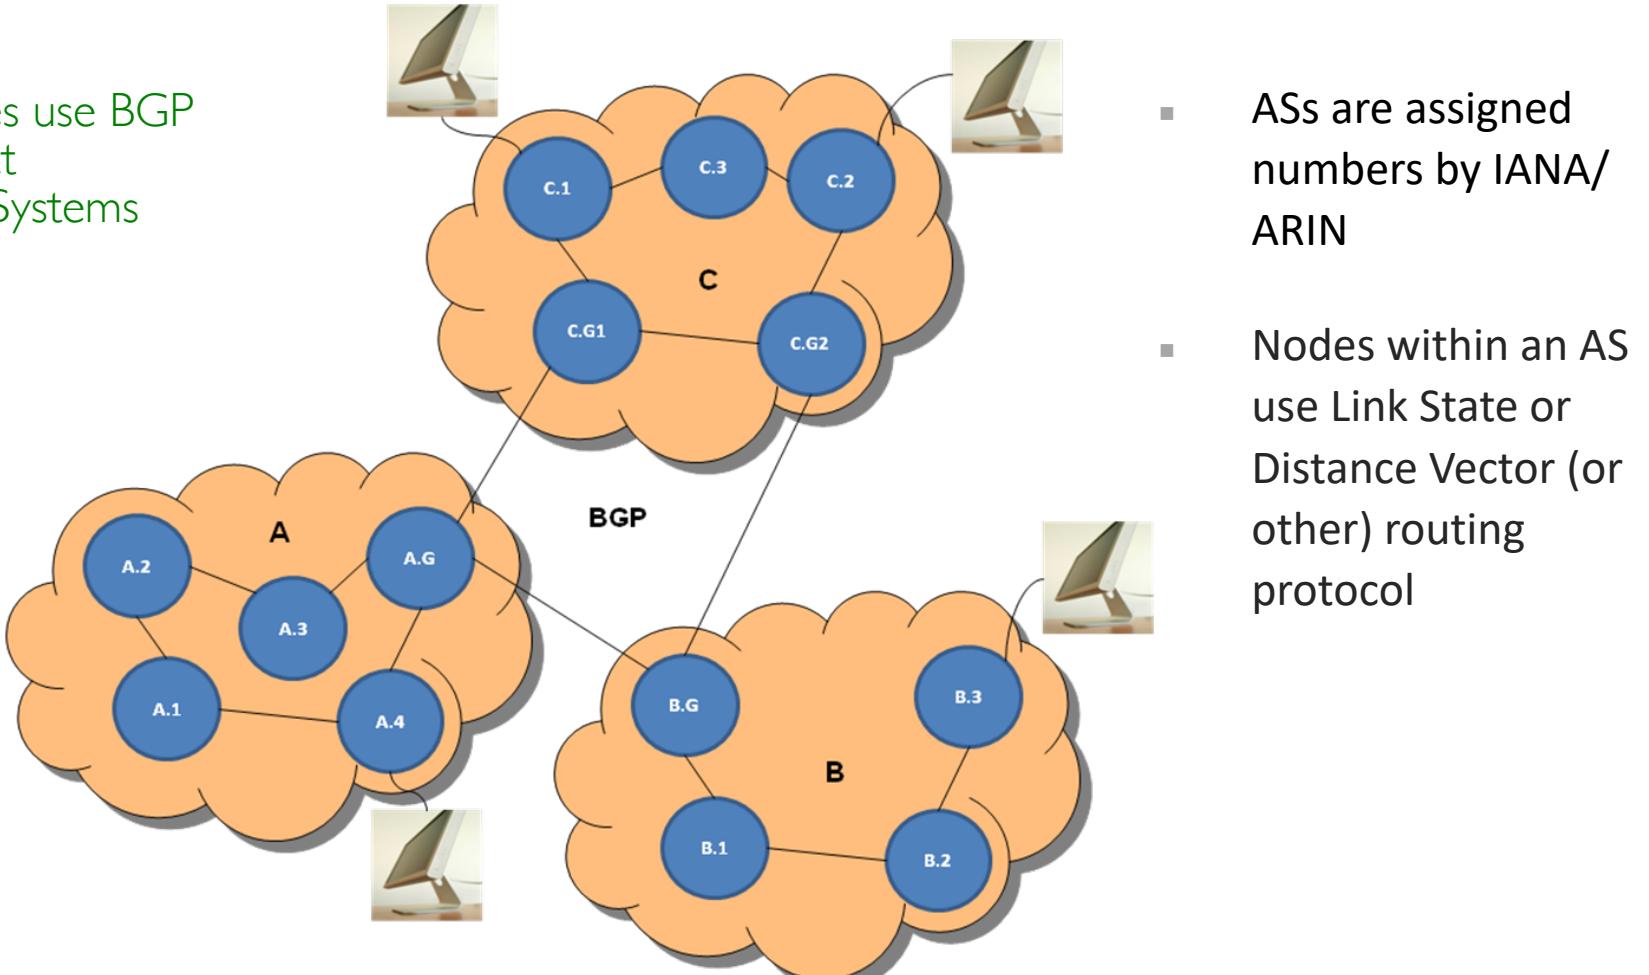
- Network of networks
- Very large scale, frequent state changes
- Compartmentalize so routing information protocol issues don't spread and cause widespread outages
- What better example than the Internet itself!
- Autonomous Systems (AS)
  - It is a collection of IP routing prefixes under the control of common administrative policy that presents a common, clearly defined routing policy to the internet.
  - Gateway nodes connect to other AS gateway nodes for inter-AS routing

# Hierarchical routing algorithms



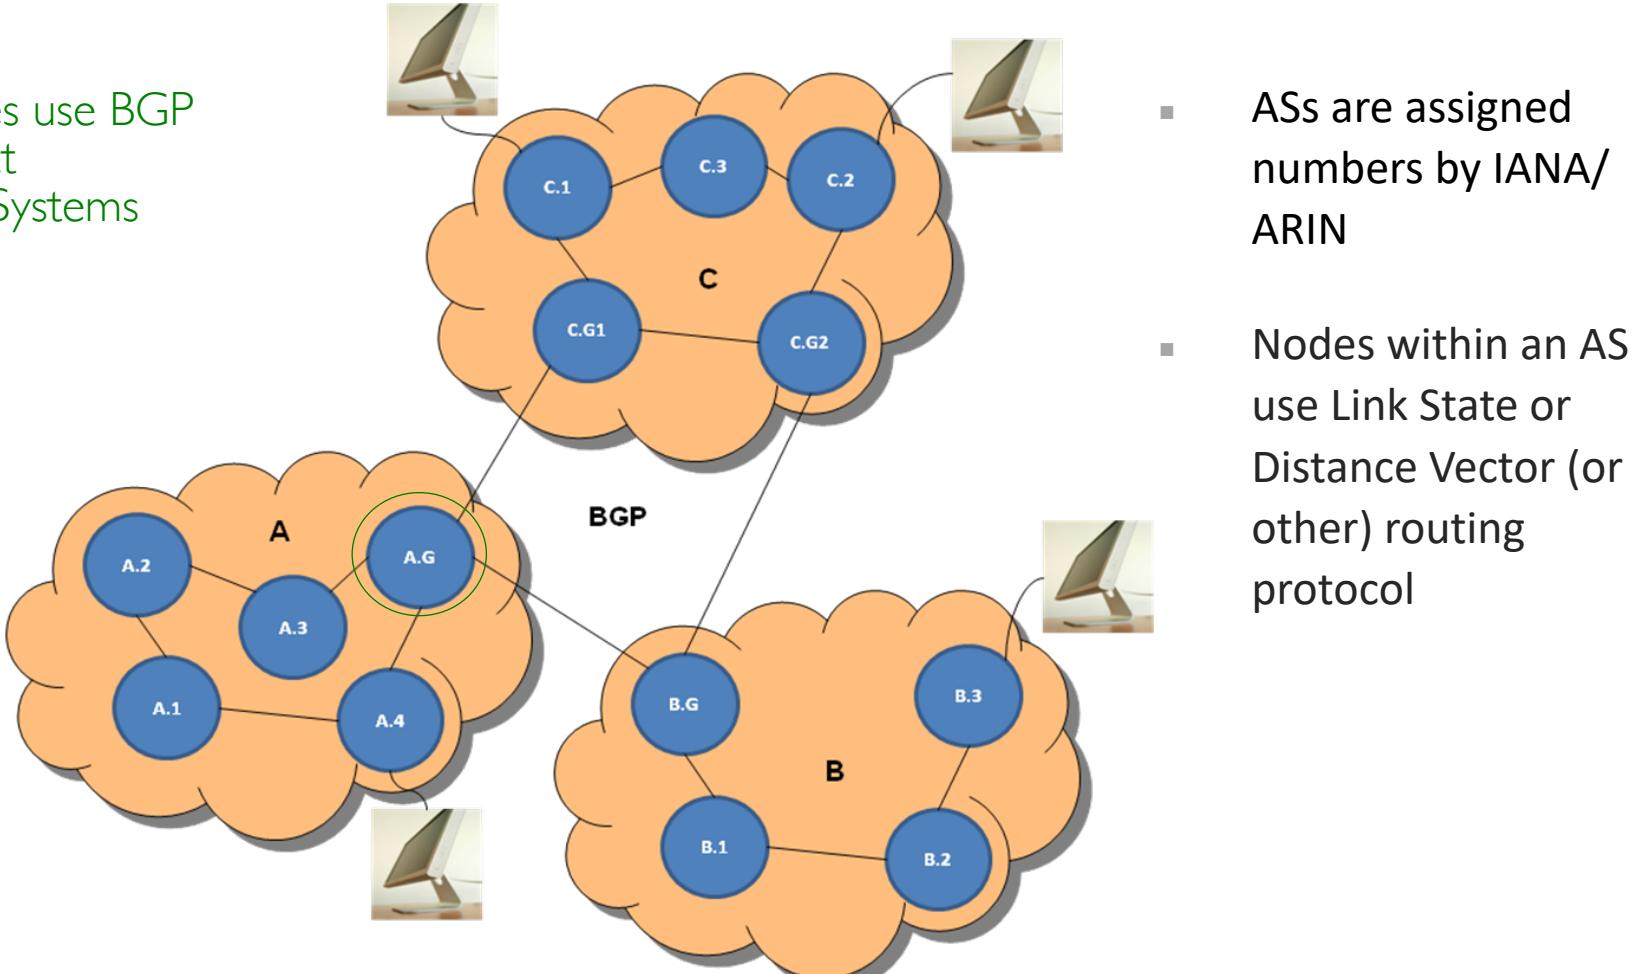
# Hierarchical routing algorithms

- Gateway nodes use BGP to interconnect Autonomous Systems



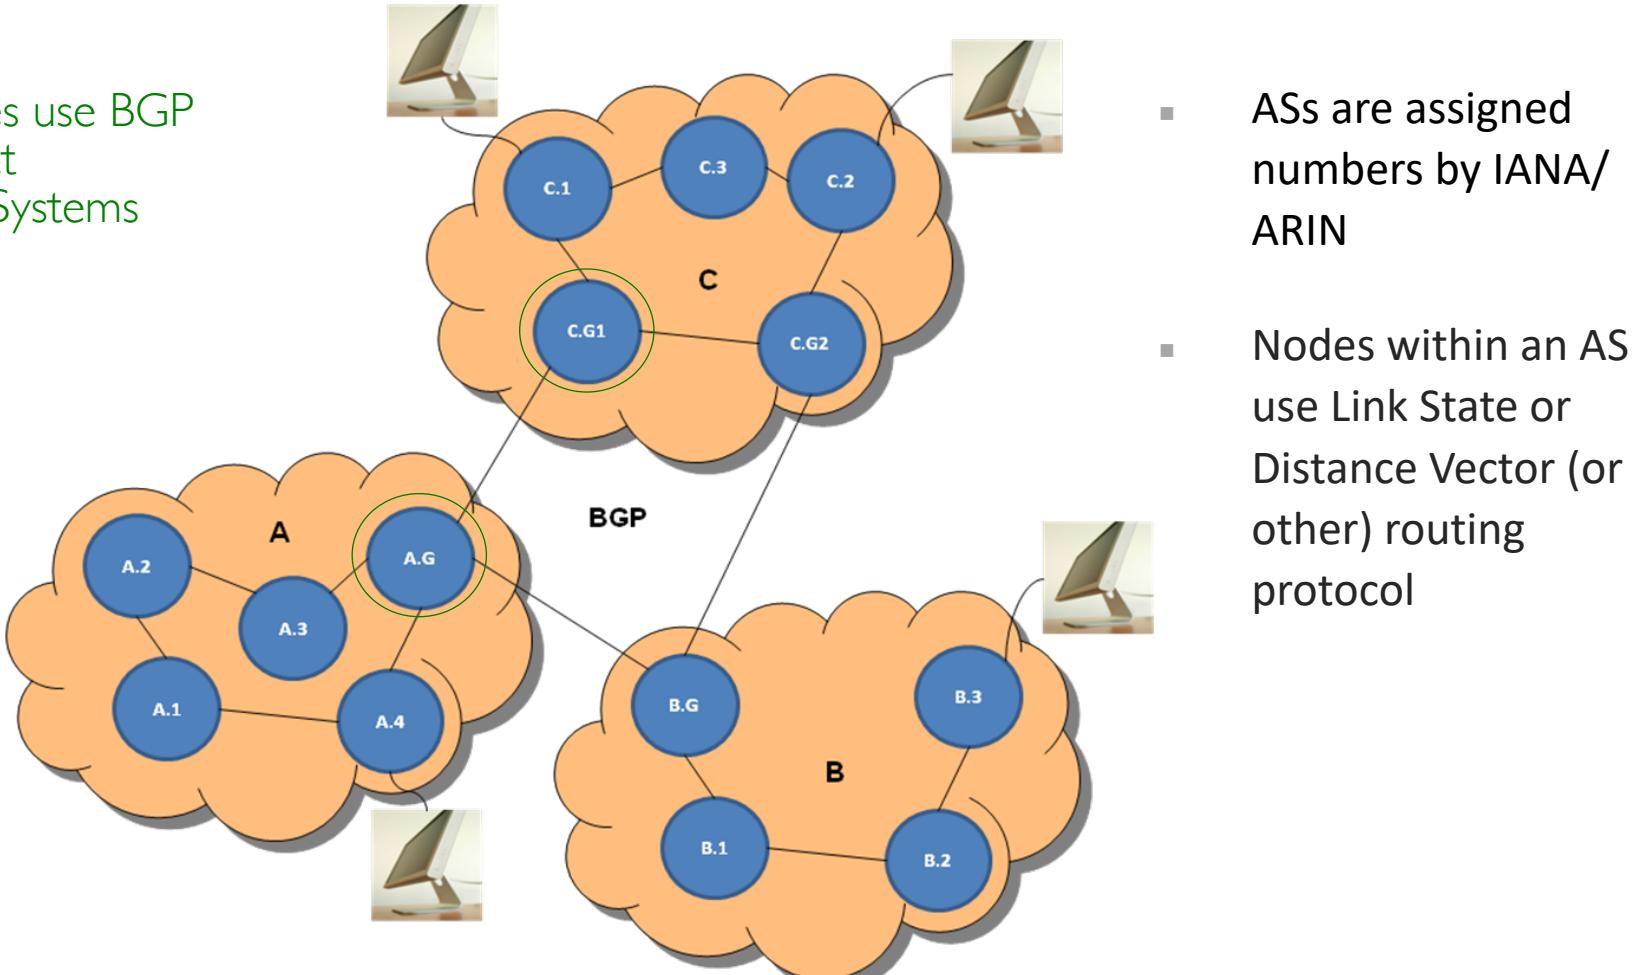
# Hierarchical routing algorithms

- Gateway nodes use BGP to interconnect Autonomous Systems



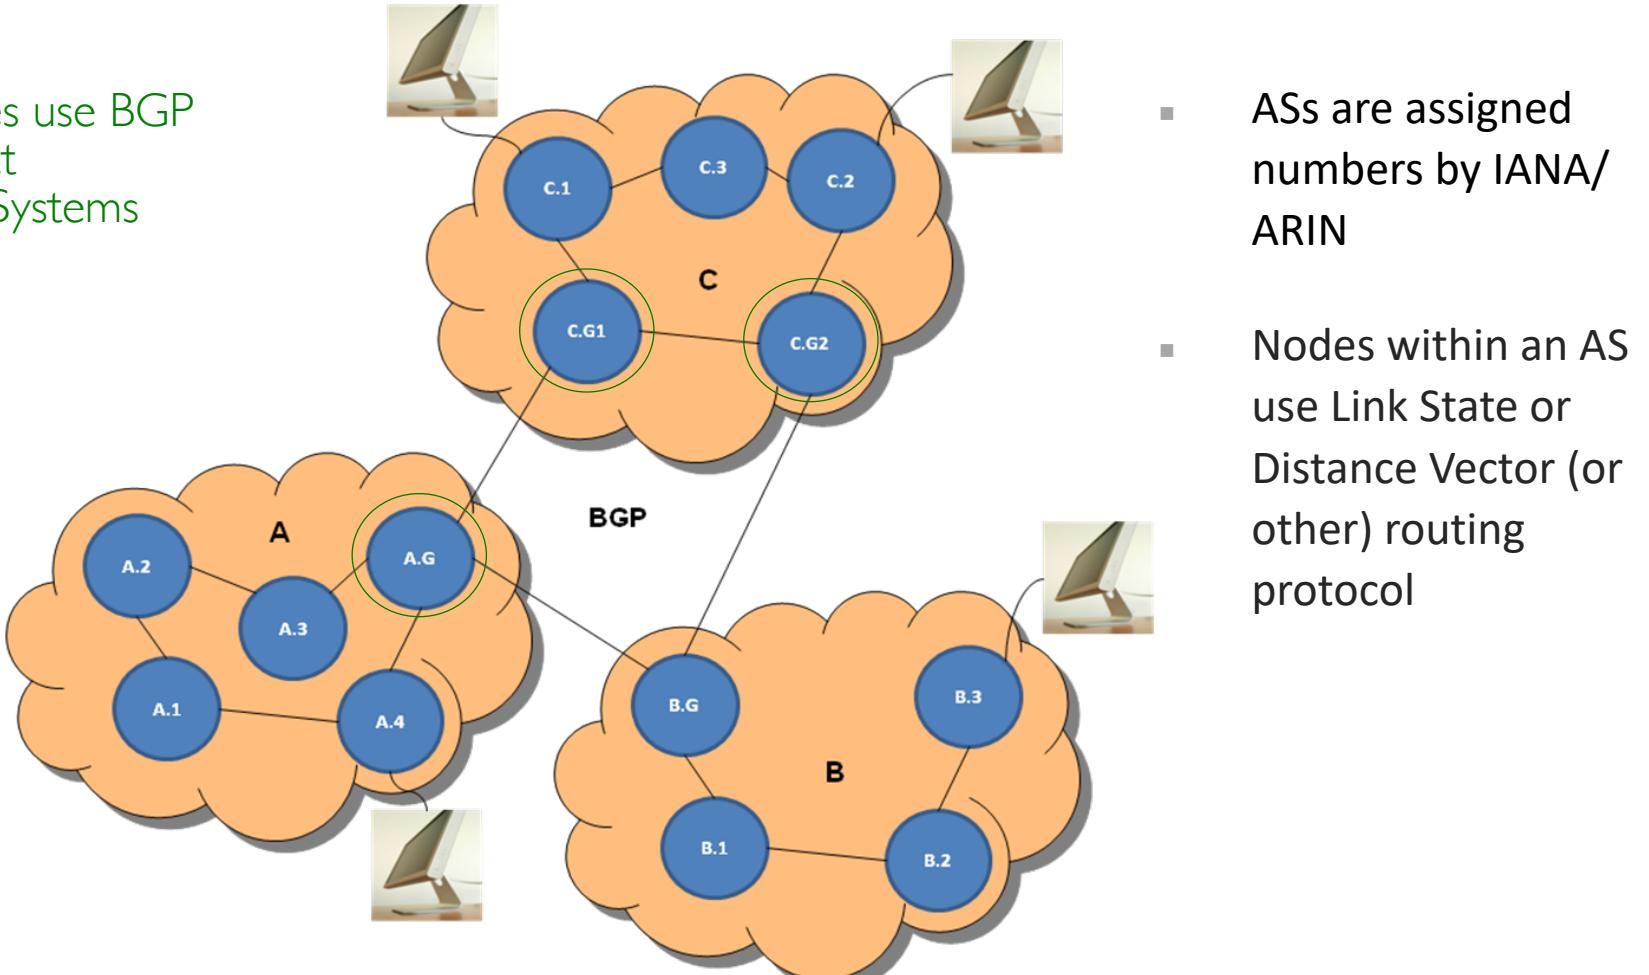
# Hierarchical routing algorithms

- Gateway nodes use BGP to interconnect Autonomous Systems



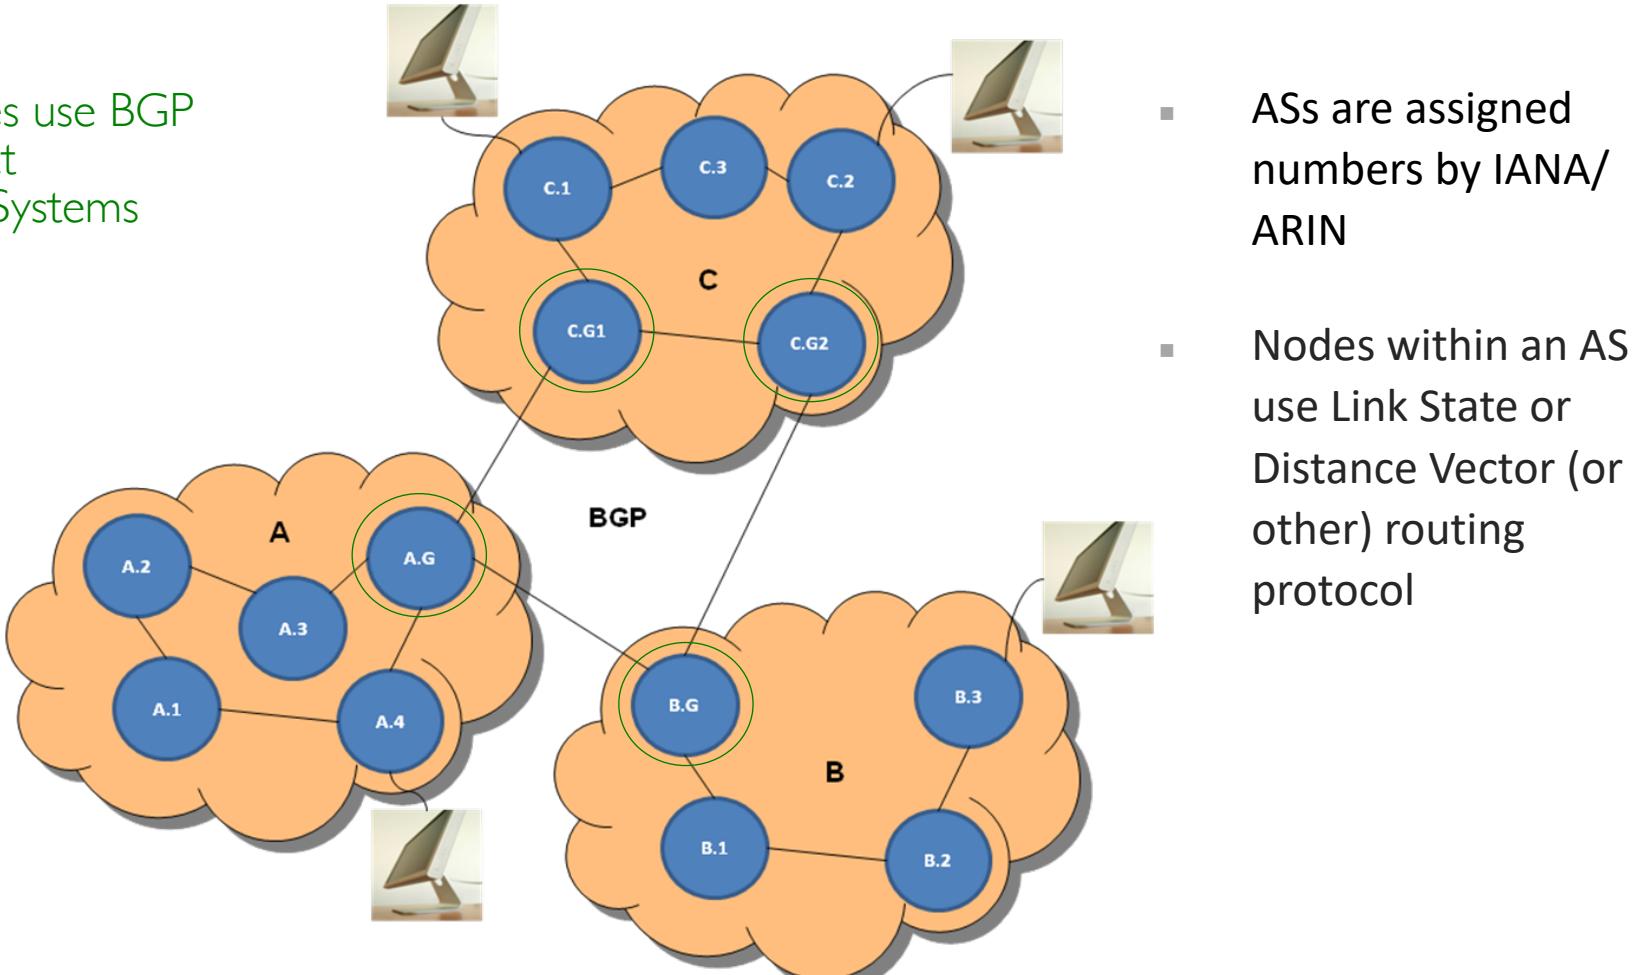
# Hierarchical routing algorithms

- Gateway nodes use BGP to interconnect Autonomous Systems



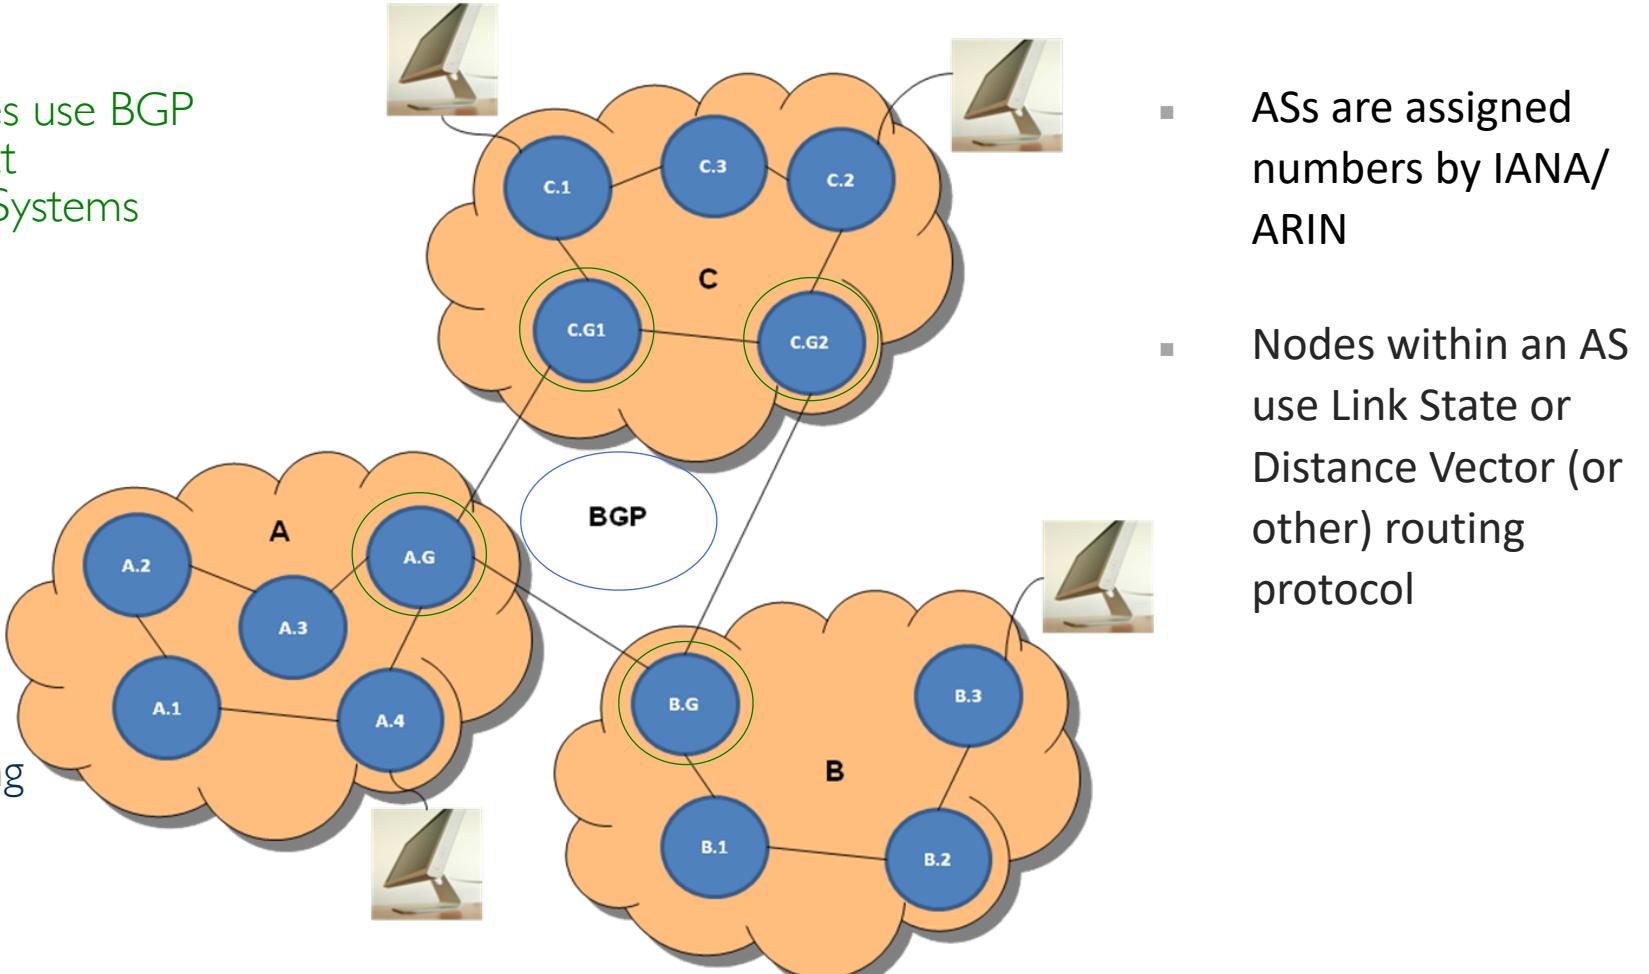
# Hierarchical routing algorithms

- Gateway nodes use BGP to interconnect Autonomous Systems



# Hierarchical routing algorithms

- Gateway nodes use BGP to interconnect Autonomous Systems



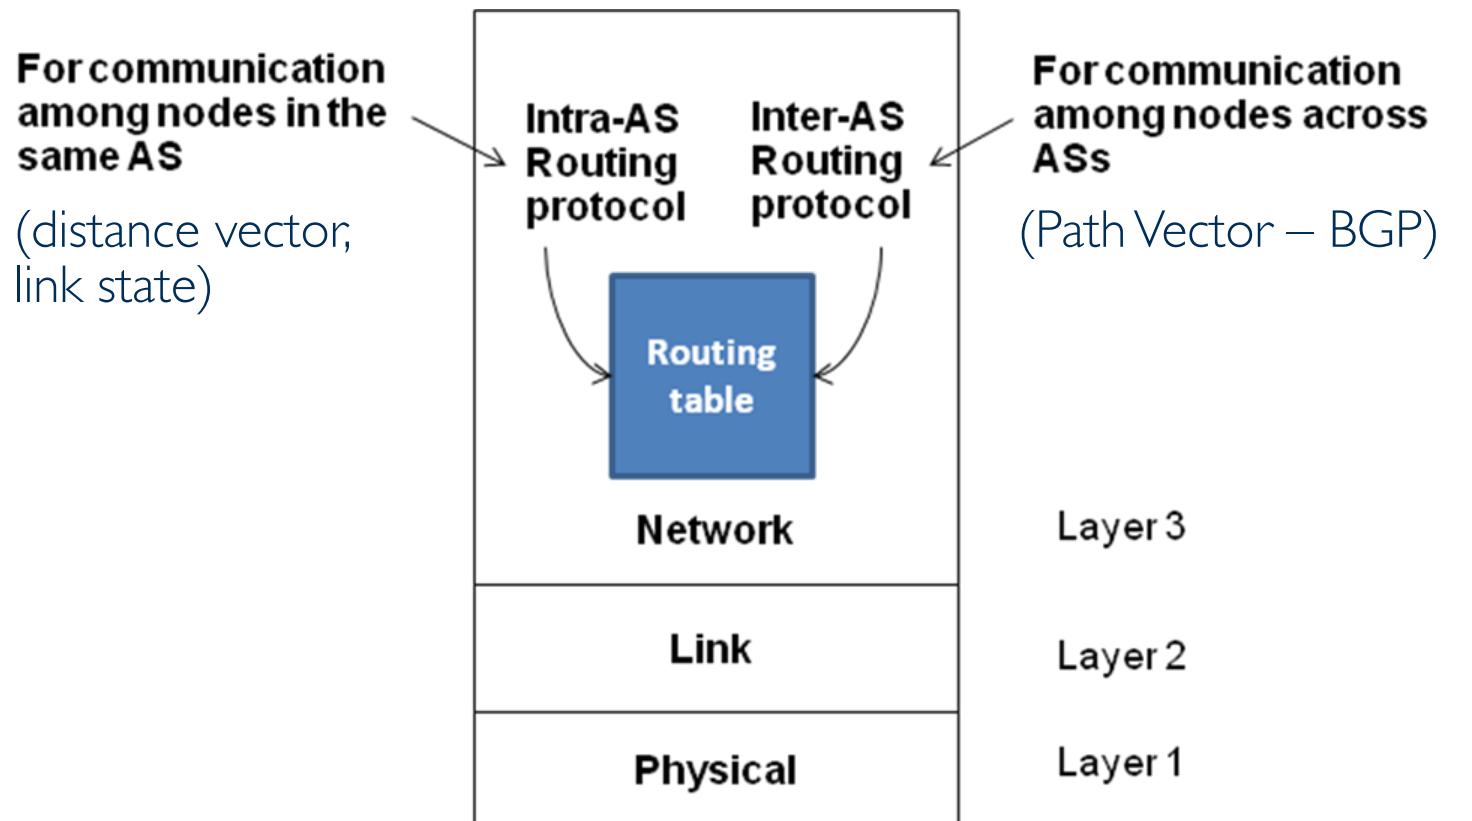
- We use BGP as the internet inter-AS routing protocol

# Border Gateway Protocol (BGP)

---

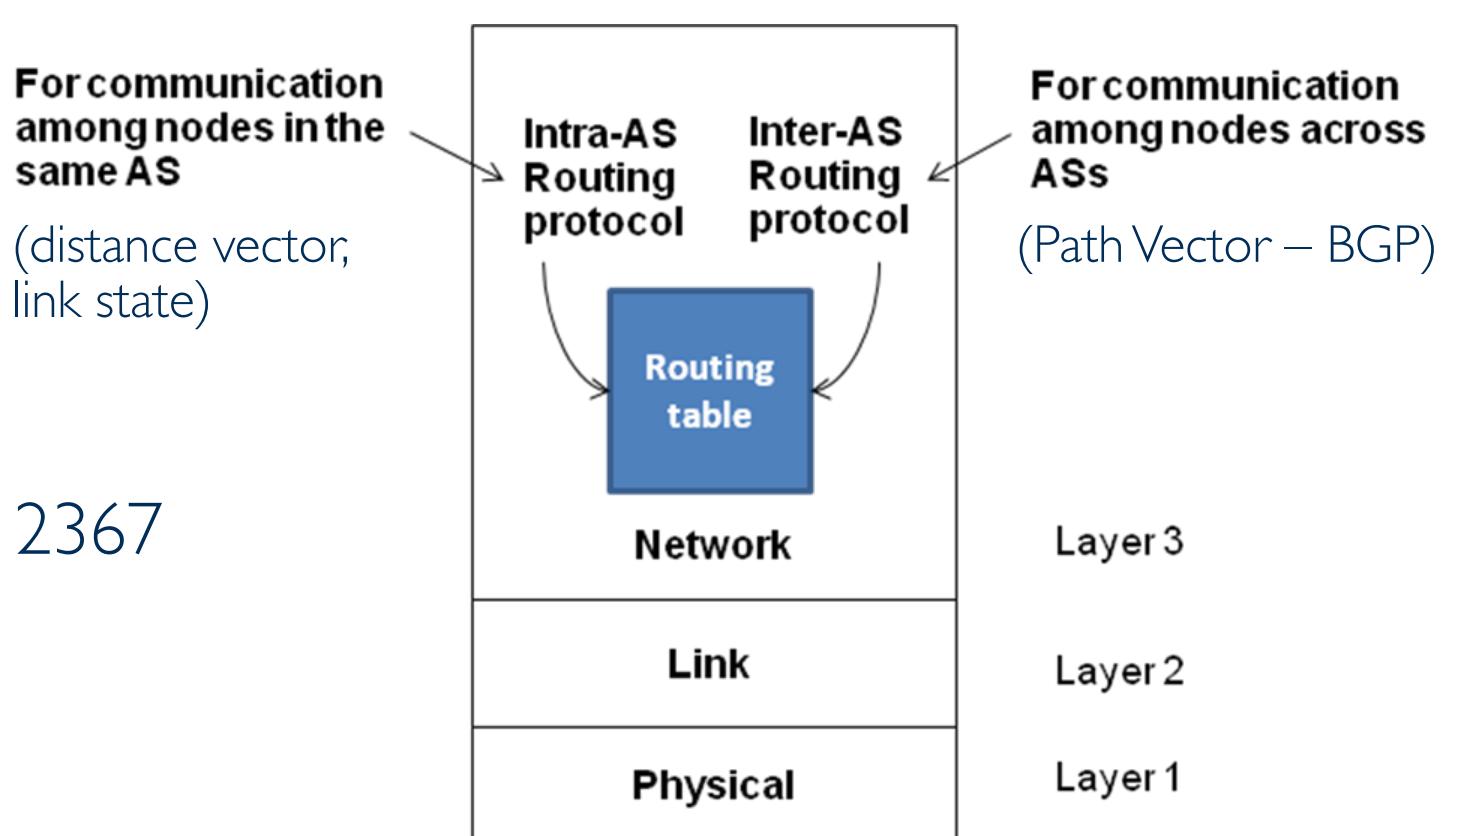
- BGP is a path-vector protocol. (We're not going into detail about that.)
- The gateway routers send path-vector messages to advertise the reachability of networks.
- Each router that receives a path vector message must verify the advertised path according to its policy.
- If the message is compliant, the router modifies its routing table and the message before sending the message to the next neighbor.
  - It modifies the routing table to maintain the autonomous systems that are traversed in order to reach the destination system.
  - It modifies the message to add its AS number and to replace the next router entry with its identification.

# Routing at GT



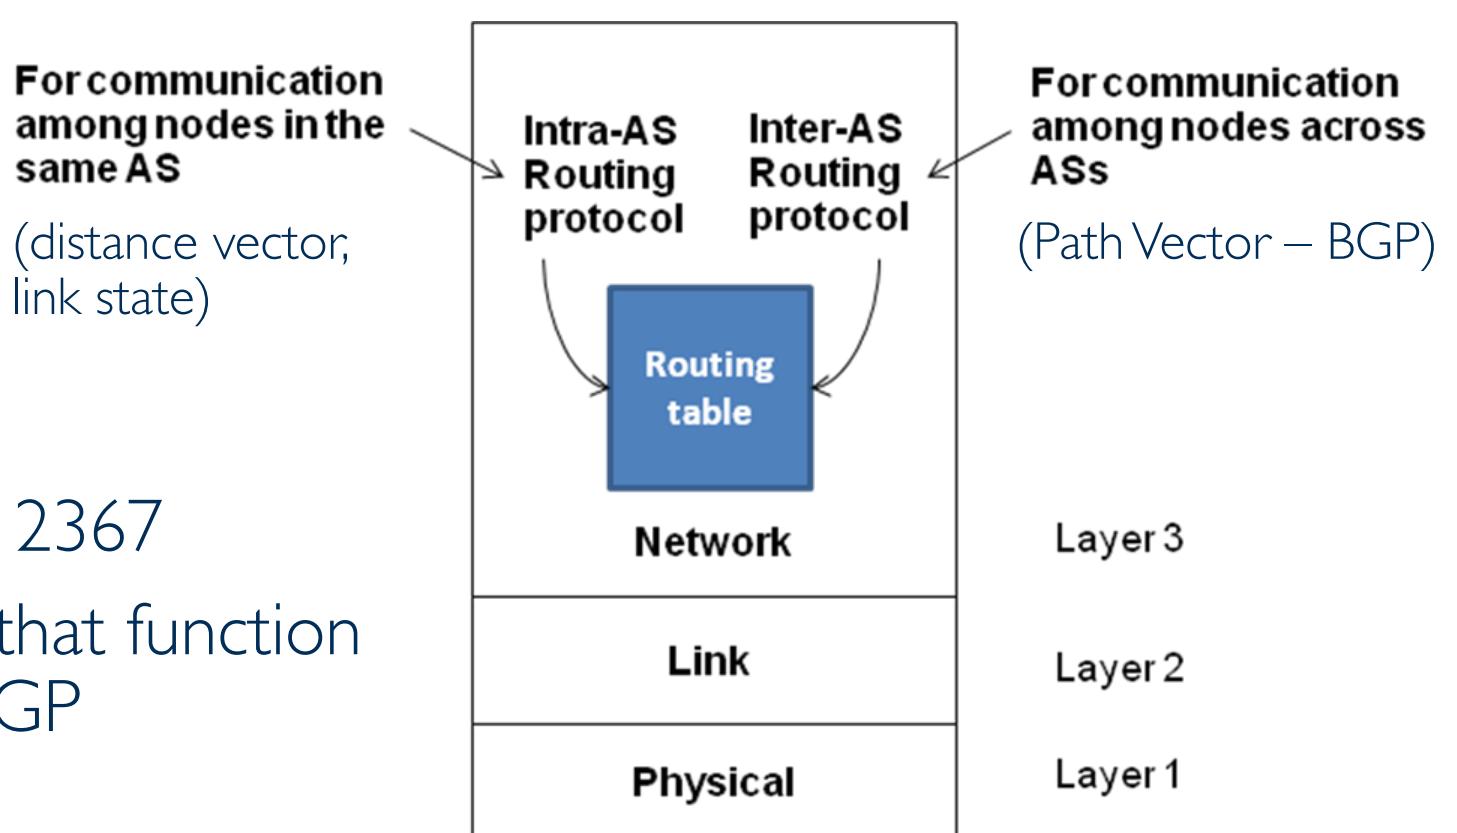
# Routing at GT

- The GT campus network is AS 2367



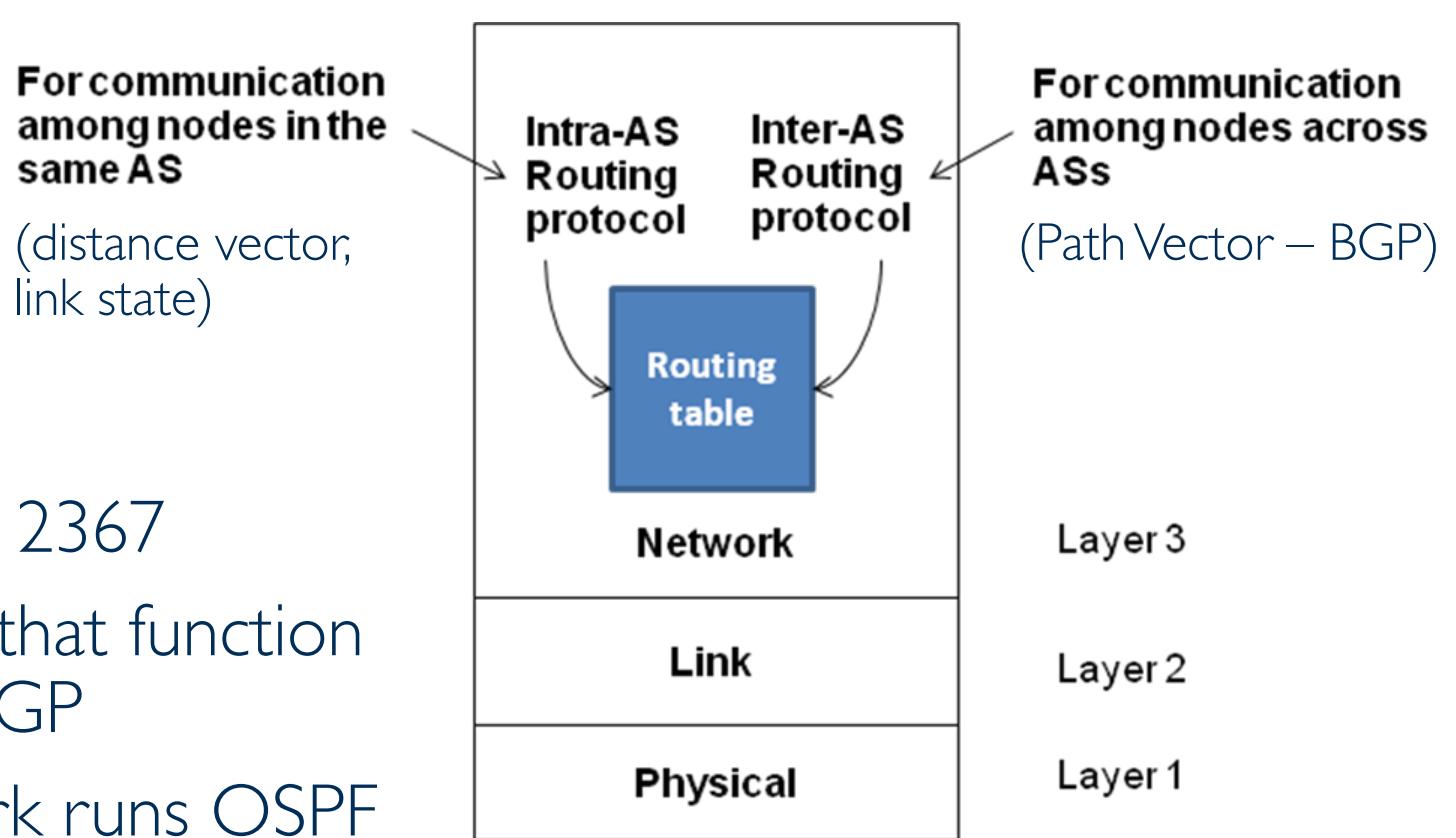
# Routing at GT

- The GT campus network is AS 2367
- There are two Juniper routers that function as Gateway Routers and run BGP



# Routing at GT

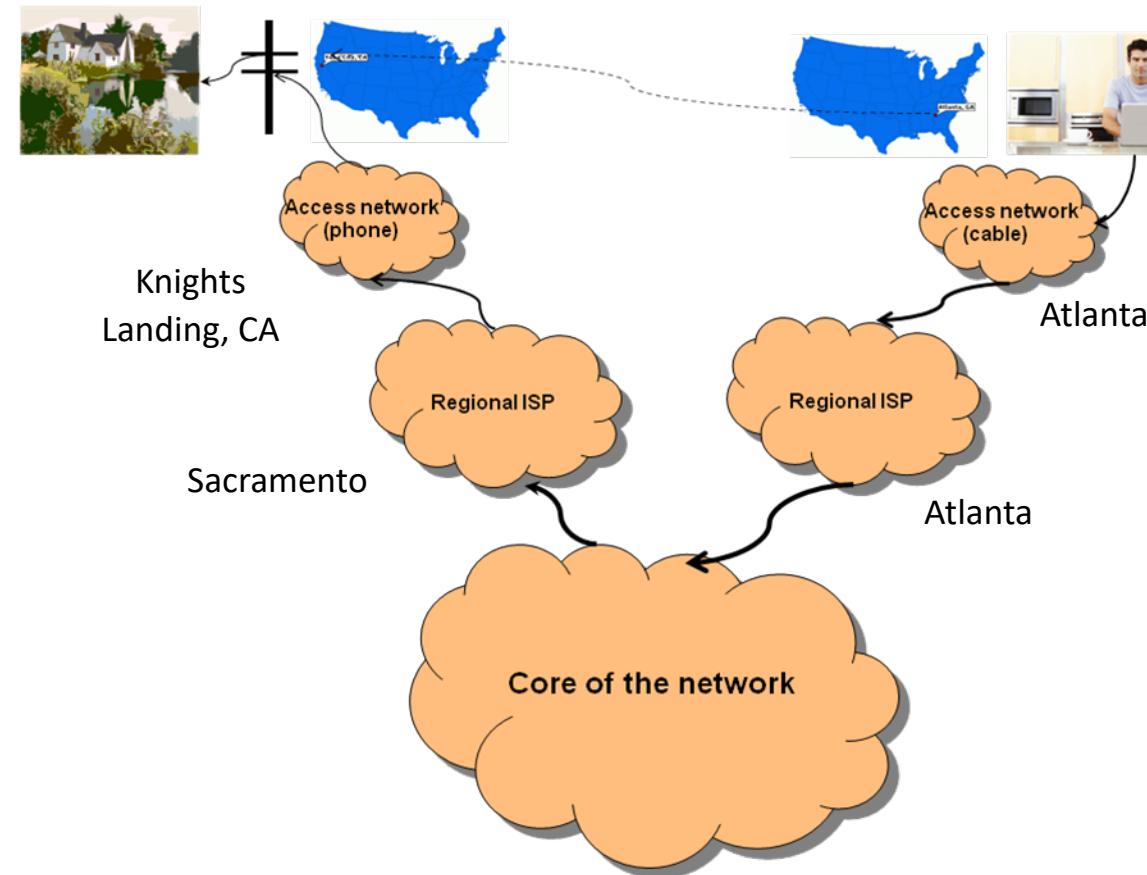
- The GT campus network is AS 2367
- There are two Juniper routers that function as Gateway Routers and run BGP
- The rest of the campus network runs OSPF (a link-state routing protocol)



# Remember this slide?

---

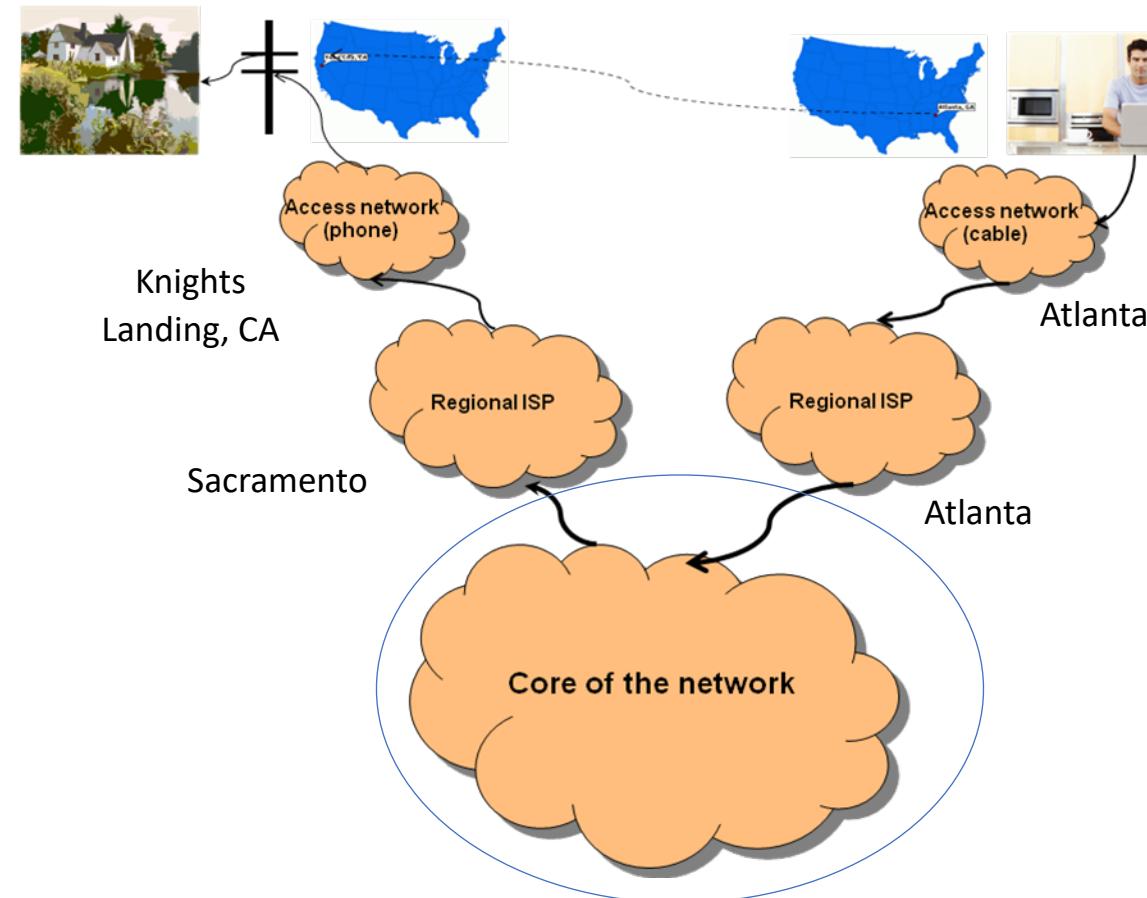
- Now consider an email from Joe to his grandmother



# Remember this slide?

---

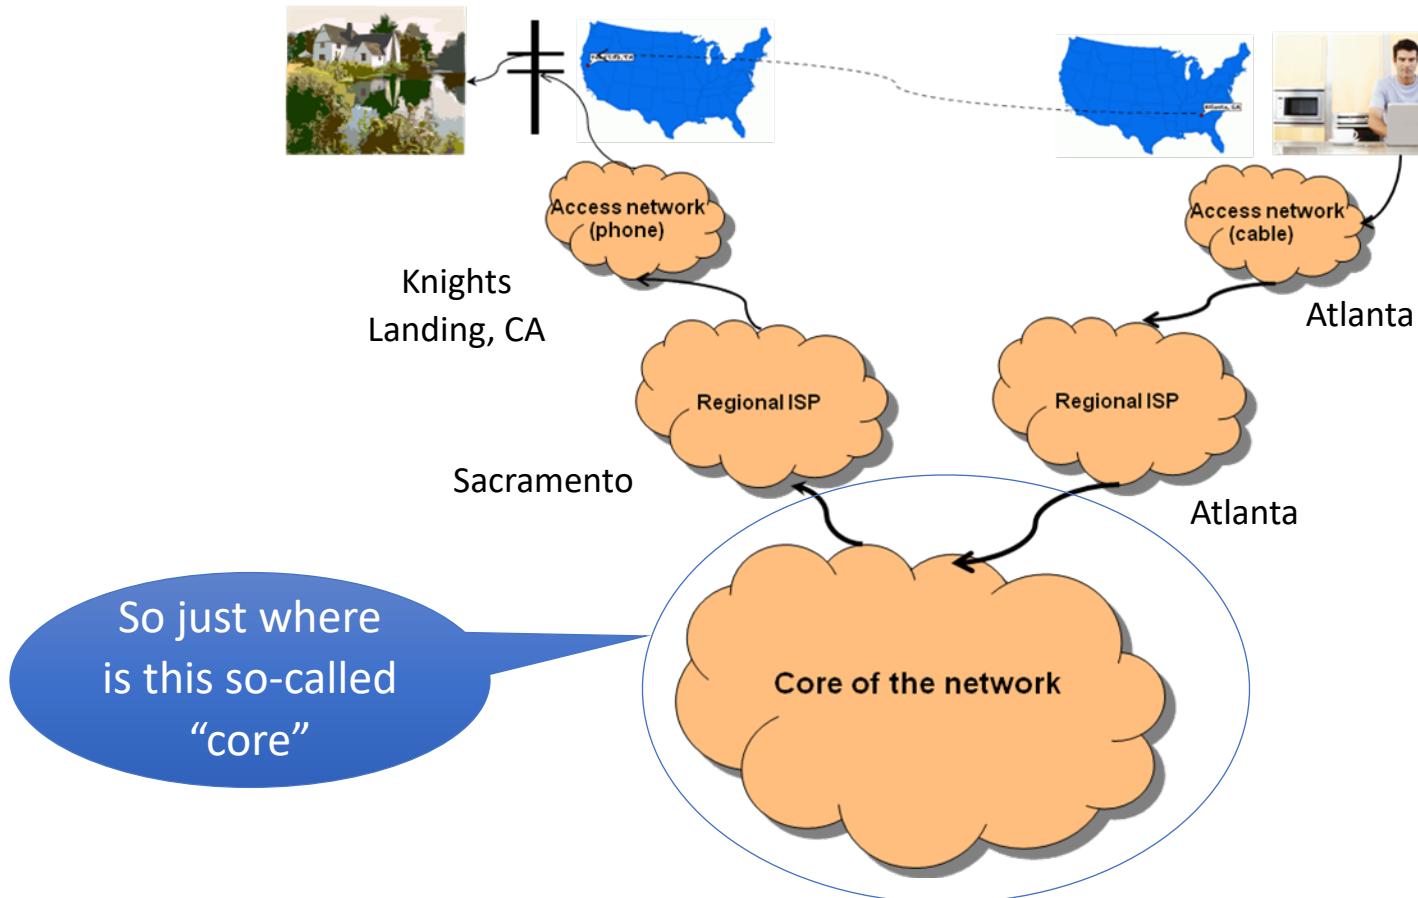
- Now consider an email from Joe to his grandmother



# Remember this slide?

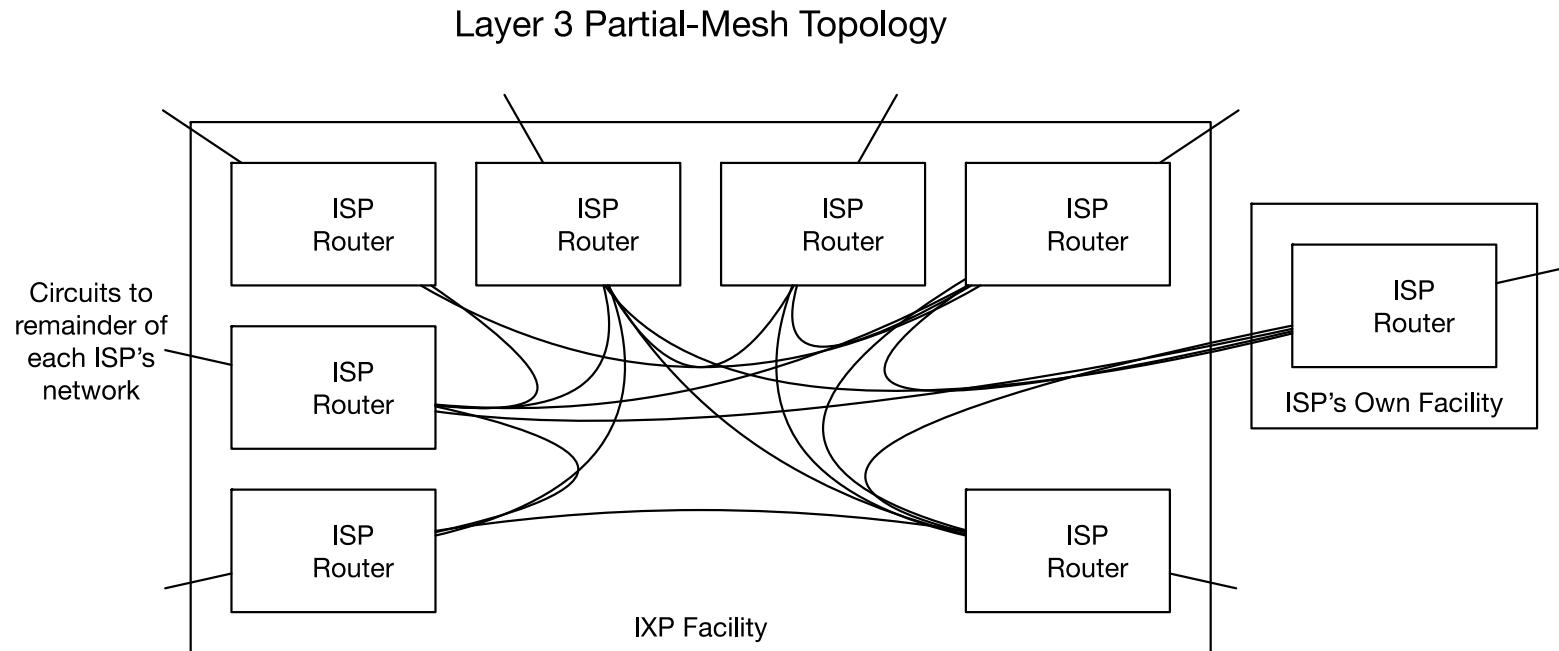
---

- Now consider an email from Joe to his grandmother



# There really is no internet core ...

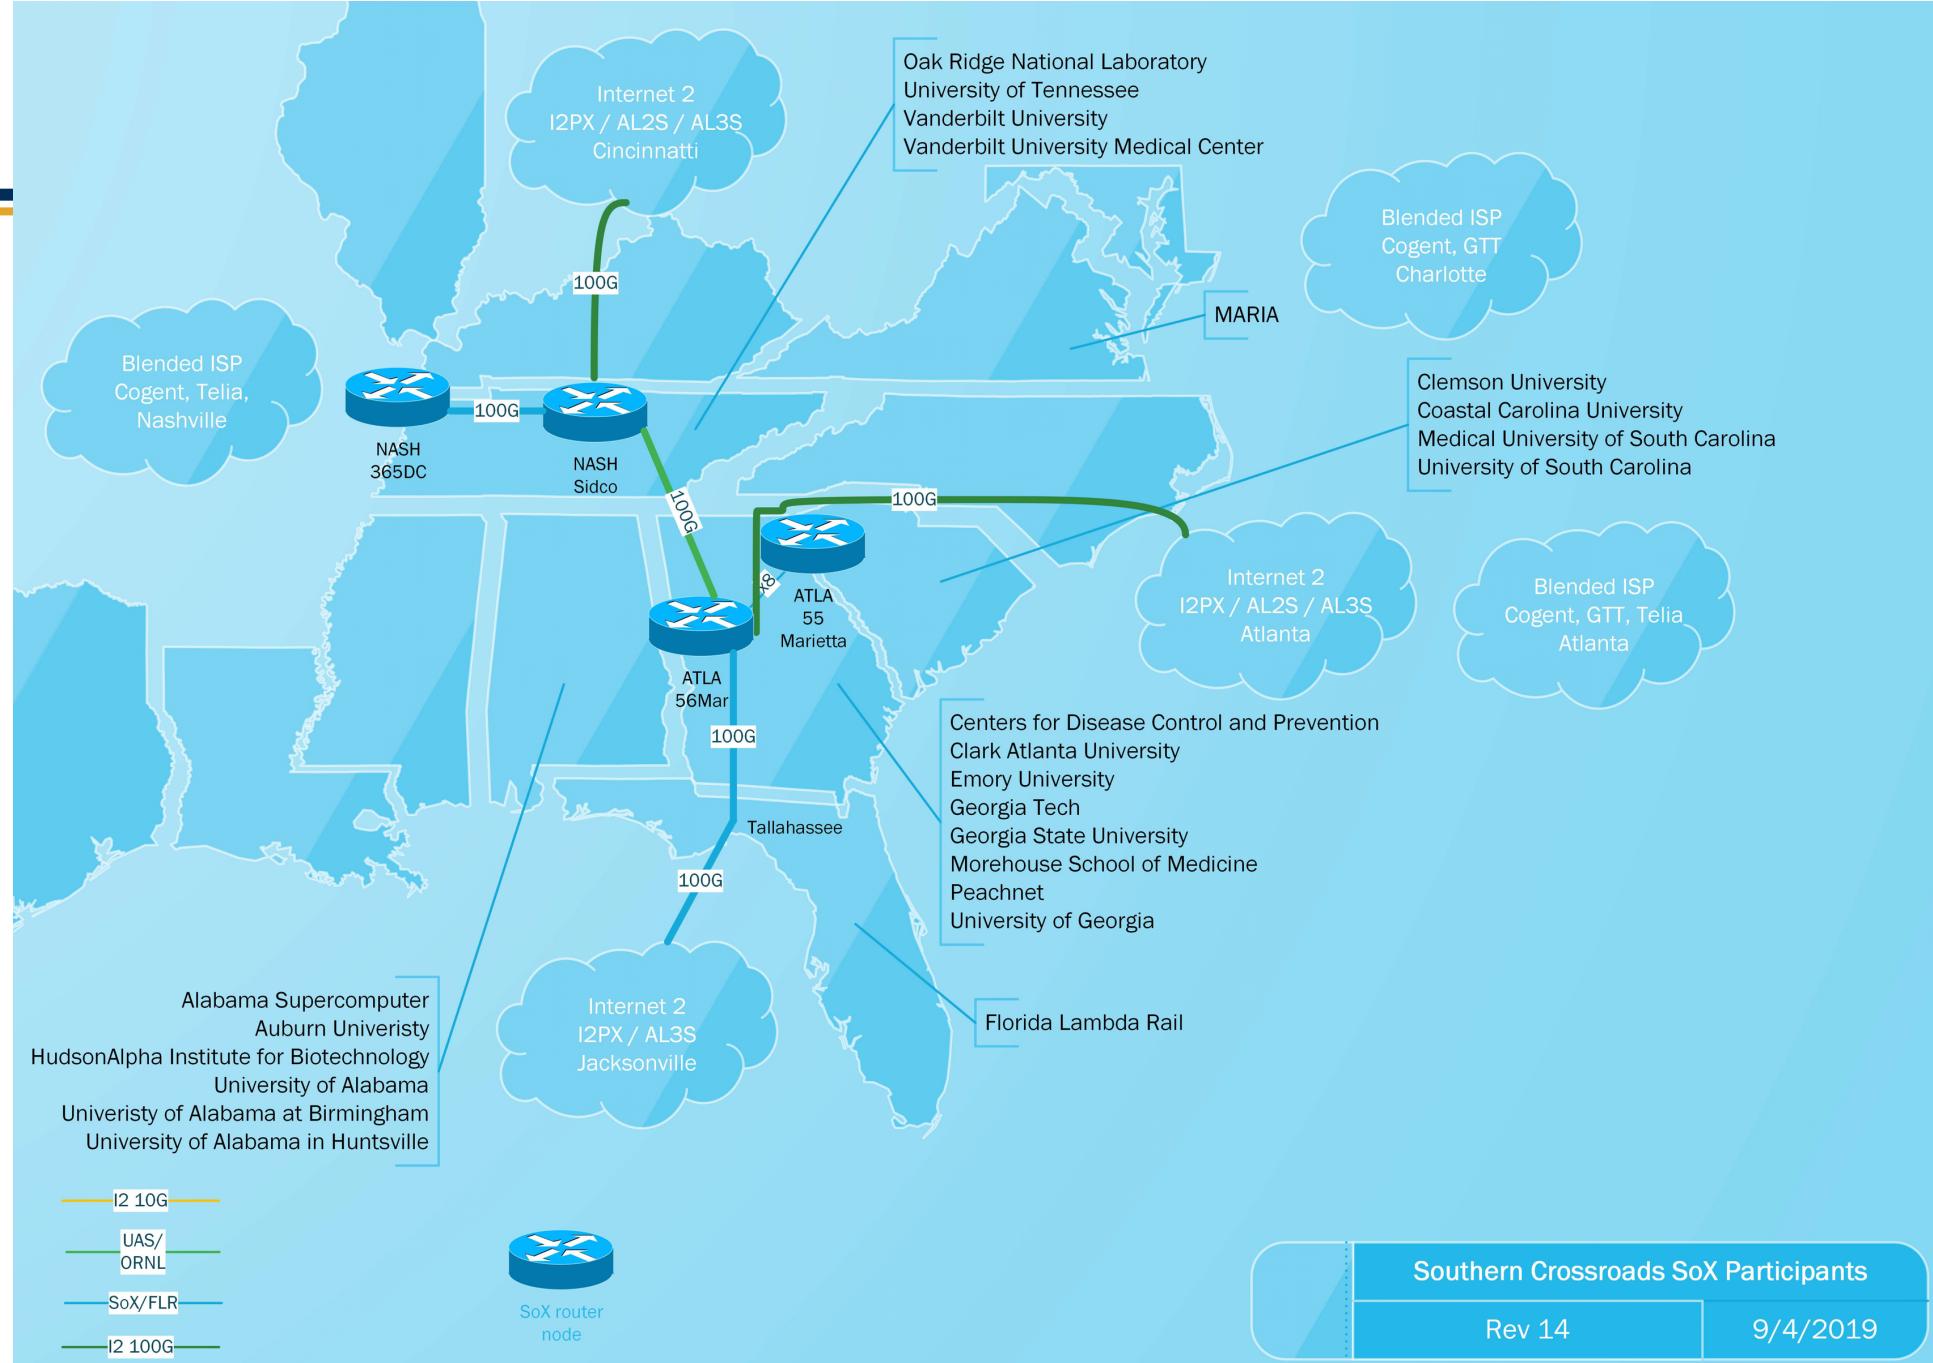
- Instead there are dozens of Internet Exchange Points (IXPs) where groups of Internet Service Providers (ISPs) interconnect their networks



# There really is no internet core ...

---

- The ISPs peer using BGP such that all the gateway routers know multiple routes to every regional ISP
- You can begin to see that the Internet has a highly distributed "core" created by multiple IXPs
- There are at least four IXPs in Atlanta; GT participates in and operates an ISP, SoX, which provides network access for about 3 dozen southeastern universities and research centers
- The diagram on the next slide shows just the IXPs in which SoX participates. Note that traffic is allowed to traverse SoX networks to get from one ISP to another.
- The diagram doesn't show the SoX peering with other entities such as Google, AT&T, NASA, Dept of Energy, and Comcast



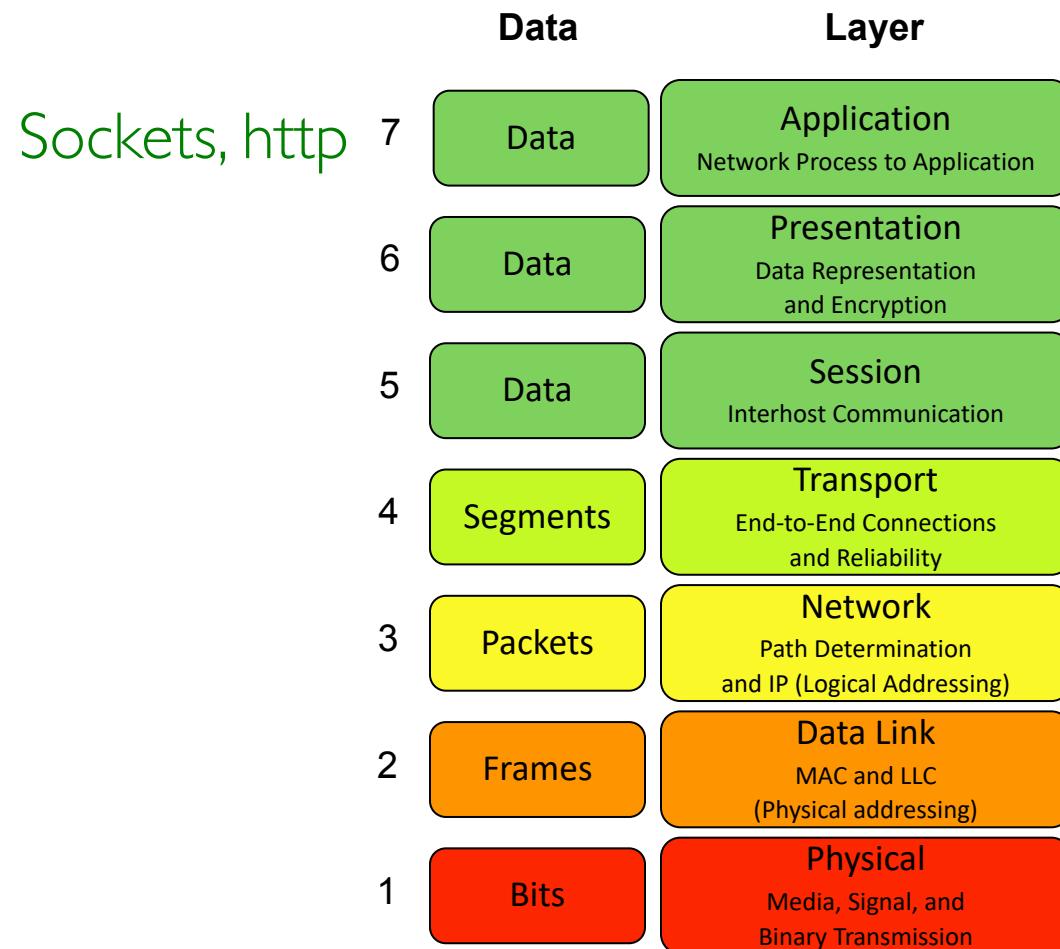
# Summary

---

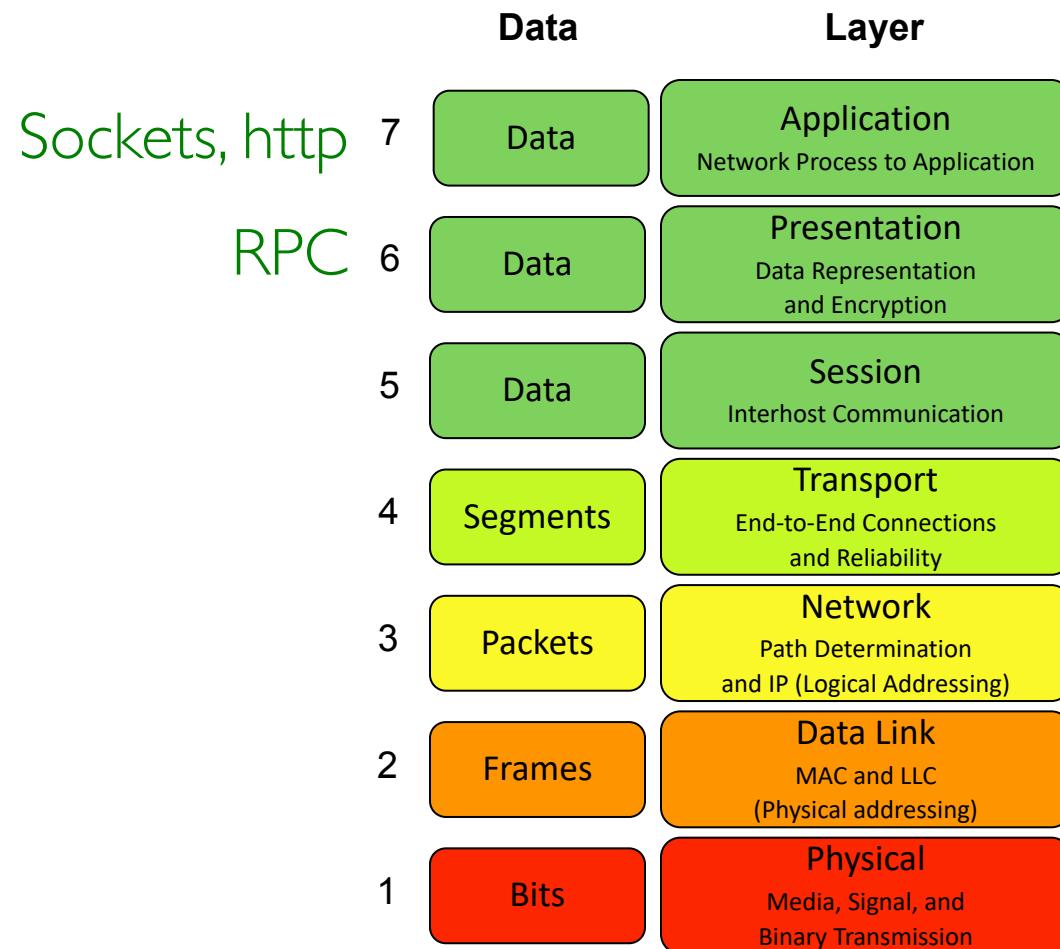
	Data	Layer
7	Data	Application Network Process to Application
6	Data	Presentation Data Representation and Encryption
5	Data	Session Interhost Communication
4	Segments	Transport End-to-End Connections and Reliability
3	Packets	Network Path Determination and IP (Logical Addressing)
2	Frames	Data Link MAC and LLC (Physical addressing)
1	Bits	Physical Media, Signal, and Binary Transmission

# Summary

---



# Summary



# Summary

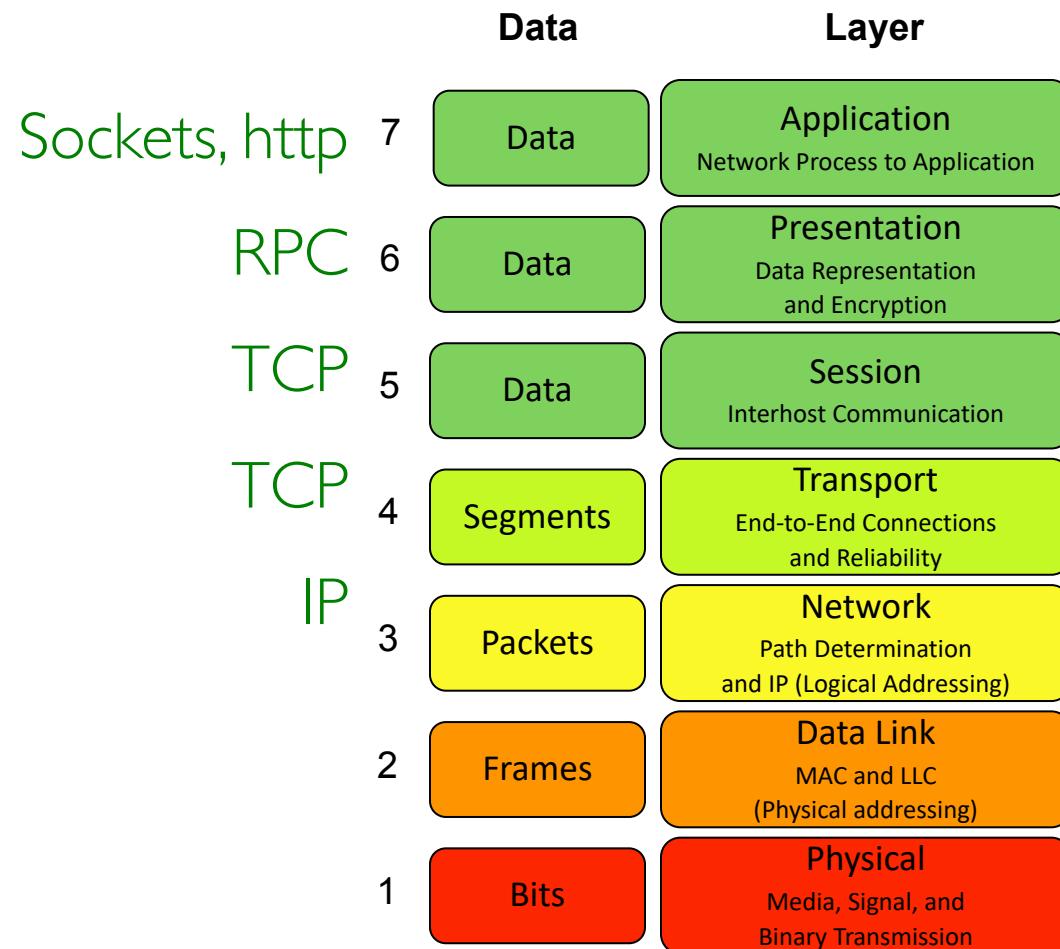
---

	Data	Layer
Sockets, http	7 Data	Application Network Process to Application
RPC	6 Data	Presentation Data Representation and Encryption
TCP	5 Data	Session Interhost Communication
	4 Segments	Transport End-to-End Connections and Reliability
	3 Packets	Network Path Determination and IP (Logical Addressing)
	2 Frames	Data Link MAC and LLC (Physical addressing)
	1 Bits	Physical Media, Signal, and Binary Transmission

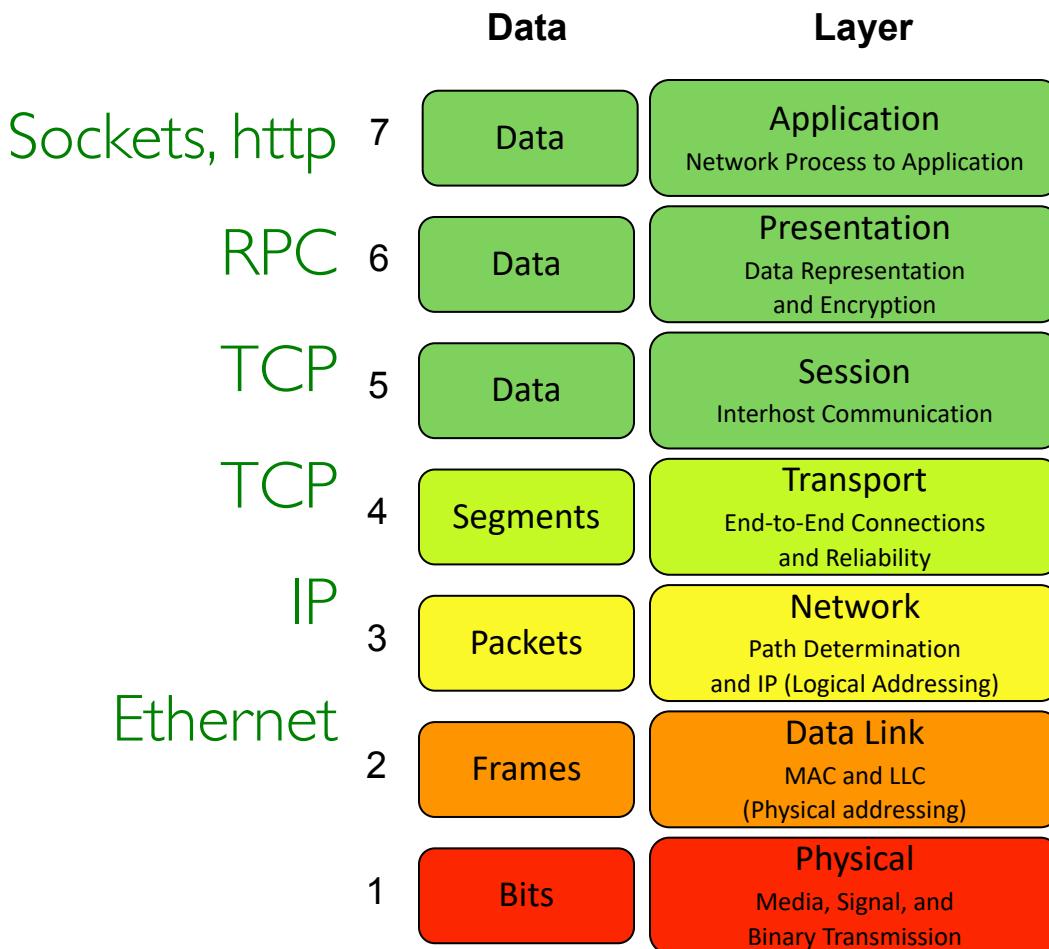
# Summary

		Data	Layer
Sockets, http	7	Data	Application Network Process to Application
RPC	6	Data	Presentation Data Representation and Encryption
TCP	5	Data	Session Interhost Communication
TCP	4	Segments	Transport End-to-End Connections and Reliability
	3	Packets	Network Path Determination and IP (Logical Addressing)
	2	Frames	Data Link MAC and LLC (Physical addressing)
	1	Bits	Physical Media, Signal, and Binary Transmission

# Summary



# Summary



# Summary

